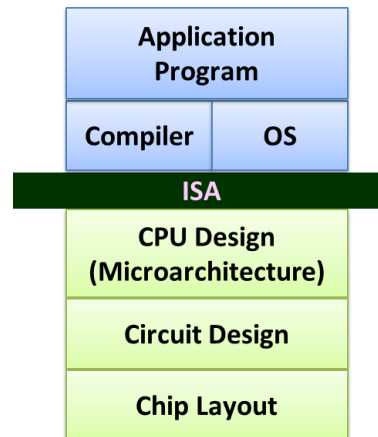


The HW/SW Interface

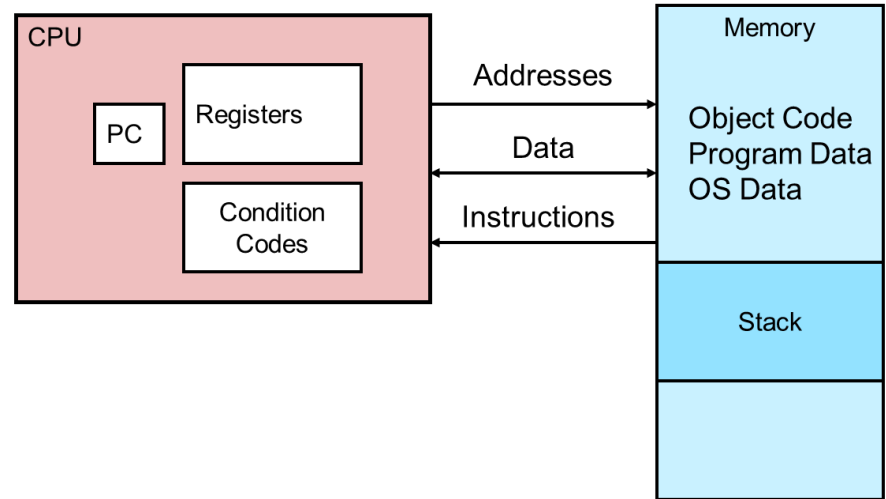
RISC-V

Machine-Level Programming Basics



Module Outline

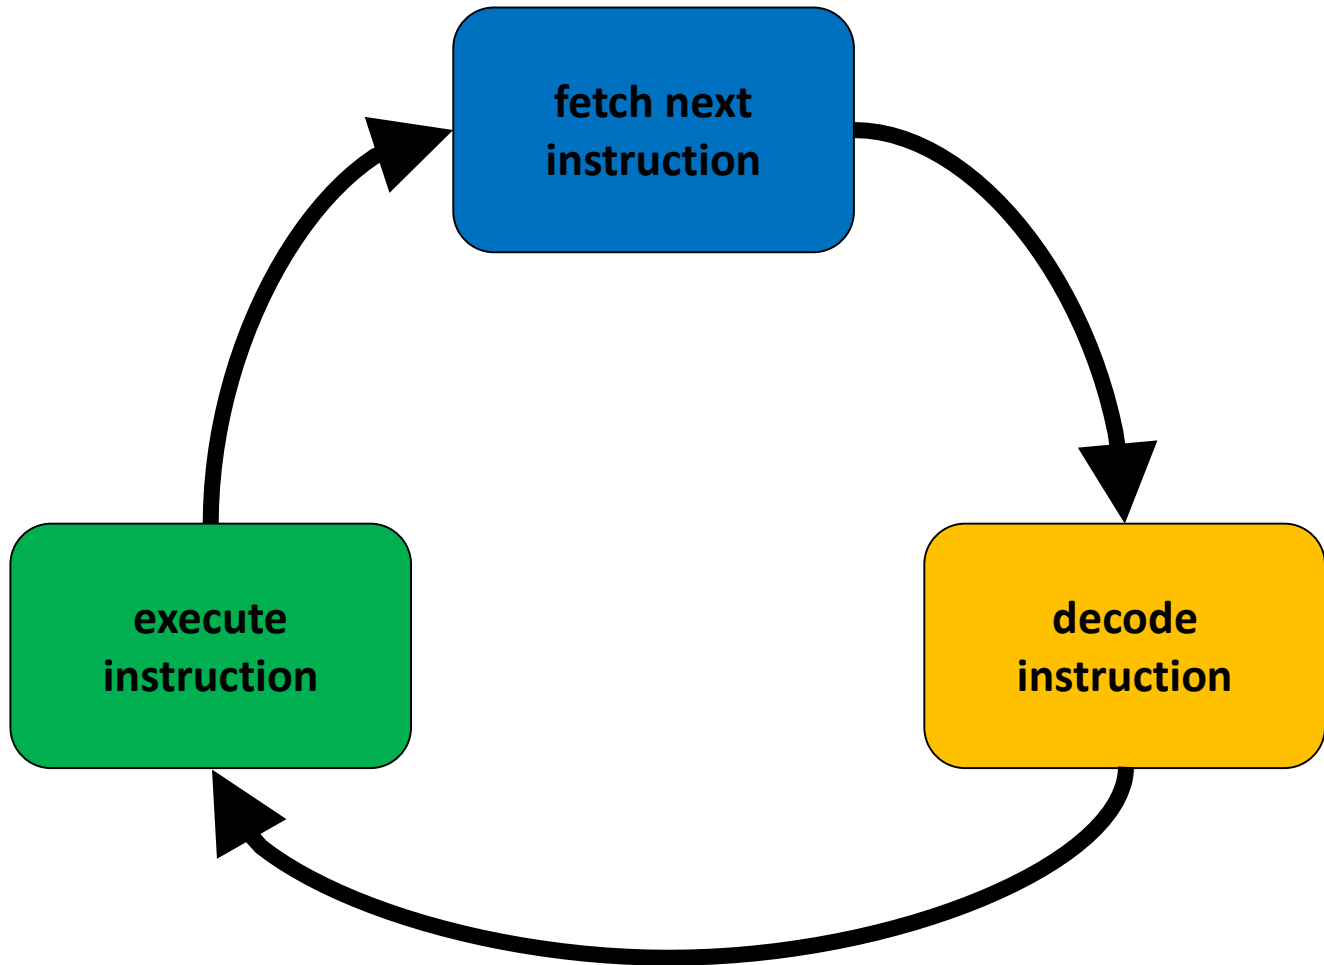
- Overview
- History of the Intel Processor Architecture
- The x86 ISA
 - The Programmer's View
 - Data Types and Alignment
 - First Steps
 - Accessing Information
 - Arithmetic and Logical Operations
 - Control Operations
 - Procedures
- Module Summary



The Hardware-Software Interface

Life as a Processor

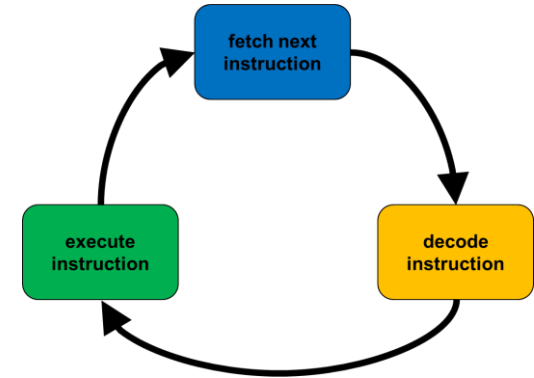
- As long as there is power

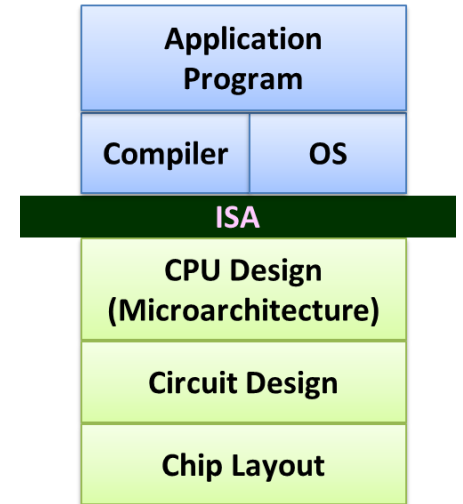


Life as a Processor

■ Lots of interesting problems to solve

- fetch next instruction
 - ▶ “next instruction”
 - ▶ fetch from where ?
- decode instruction
 - ▶ instruction format
 - ▶ supported operations
 - ▶ supported operands
 - ▶ where do operands come from/go to?
- execute instruction
 - ▶ implementation





Instruction Set Architecture: The Programmer's View

Instruction Set Architecture (ISA)

“the attributes of a system as seen by the programmer, i.e., the conceptual structure and functional behavior, as distinct from the organization of the data flow and controls, the logical design, and the physical implementation”

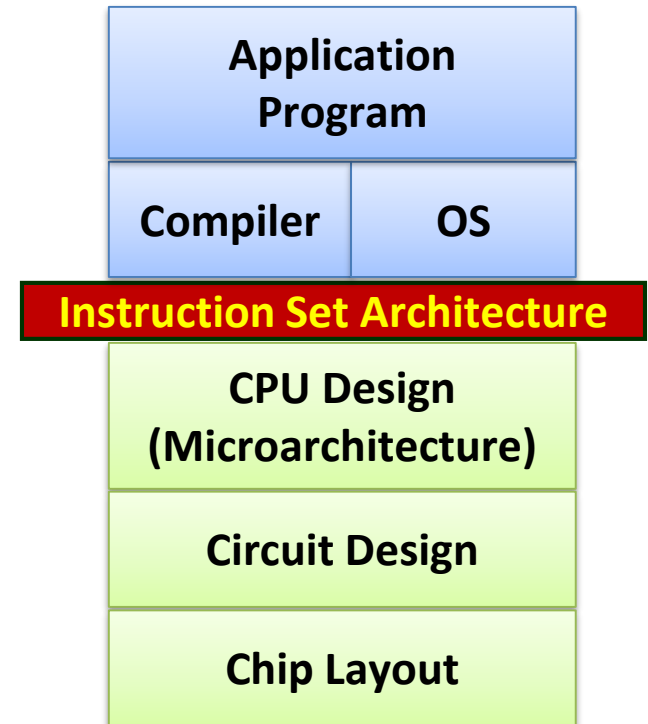
*-- Amdahl, Blaauw, and Brooks, Architecture of the IBM System/360,
IBM Journal of Research and Development, April 1964.*

- The visible interface between software and hardware
- What the user (OS, compiler, programmer, ...) needs to know in order to reason about how the machine behaves
- Abstracted from the details of how it may accomplish its task

Instruction Set Architecture (ISA)

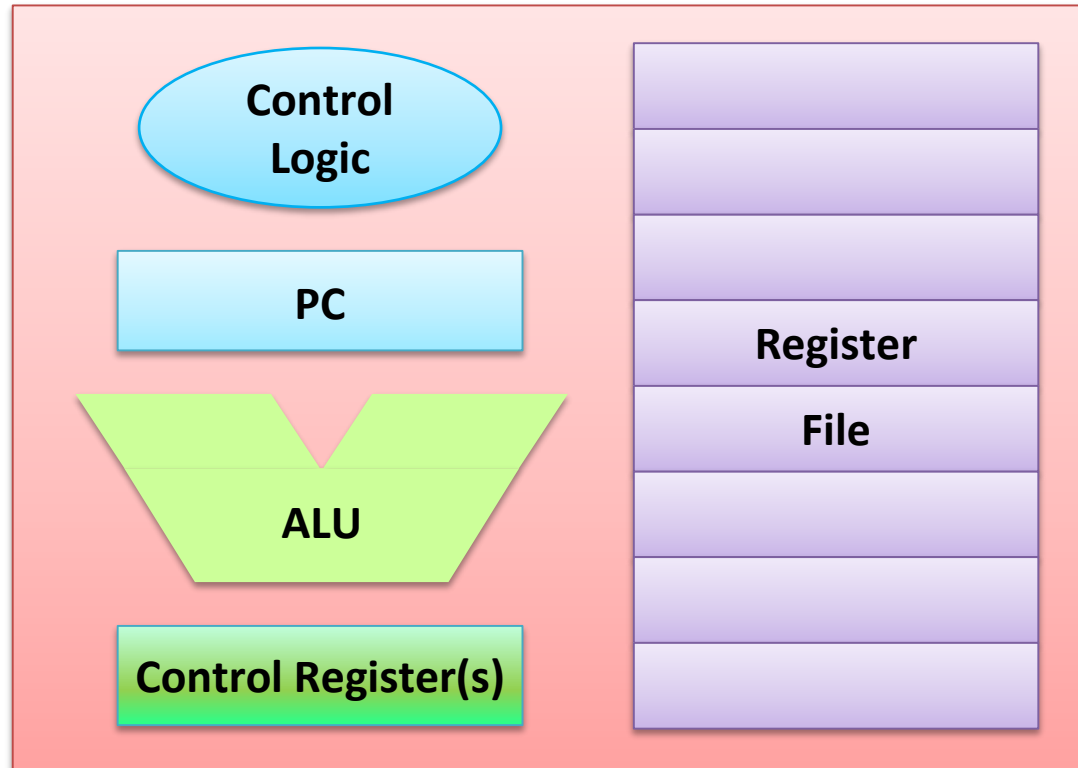
■ The ISA abstracts the details of the processor implementation

- Above: how to program machine
 - ▶ Processors execute instructions in sequence
- ISA: hardware/software interface
- Below: what needs to be built
 - ▶ Use variety of tricks to make it run fast



Instruction Set Architecture (ISA)

- The ISA defines the programmer-visible state of a processor architecture



- the ISA tells us what we can expect the processor to do when we execute instructions (but not *how* the processor does it)

Instruction Set Architecture (ISA)

- The ISA defines the programmer-visible state of a processor architecture
 - Memory addressing
 - ▶ interpretation of addresses
 - bytes, half words, words, double words, ...?
 - ▶ byte ordering in memory
 - little-endian vs. big-endian
 - ▶ addressing modes
 - how to access objects in memory

Instruction Set Architecture (ISA)

■ The ISA defines the programmer-visible state of a processor architecture

● Type and size of operands

- ▶ specifying types
 - byte, half word, word, float, double...

● Operations

- ▶ arithmetic and logical
- ▶ data transfer
- ▶ control
- ▶ system
- ▶ floating point
- ▶ string
- ▶ multimedia

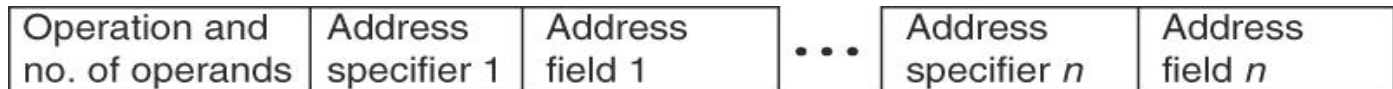
Rank	x86 operation	%
1	load	22
2	conditional branch	20
3	compare	16
4	store	12
5	add	8
6	and	6
7	sub	5
8	move reg-reg	4
9	call	1
10	return	1
	total	96

Instruction Set Architecture (ISA)

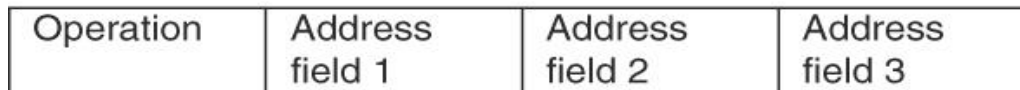
■ The ISA defines the programmer-visible state of a processor architecture

● Instruction encoding

- ▶ defines binary representation for each operation and operand
- ▶ fixed size vs. variable size



(a) Variable (e.g., Intel 80x86, VAX)



(b) Fixed (e.g., Alpha, ARM, MIPS, PowerPC, SPARC, SuperH)

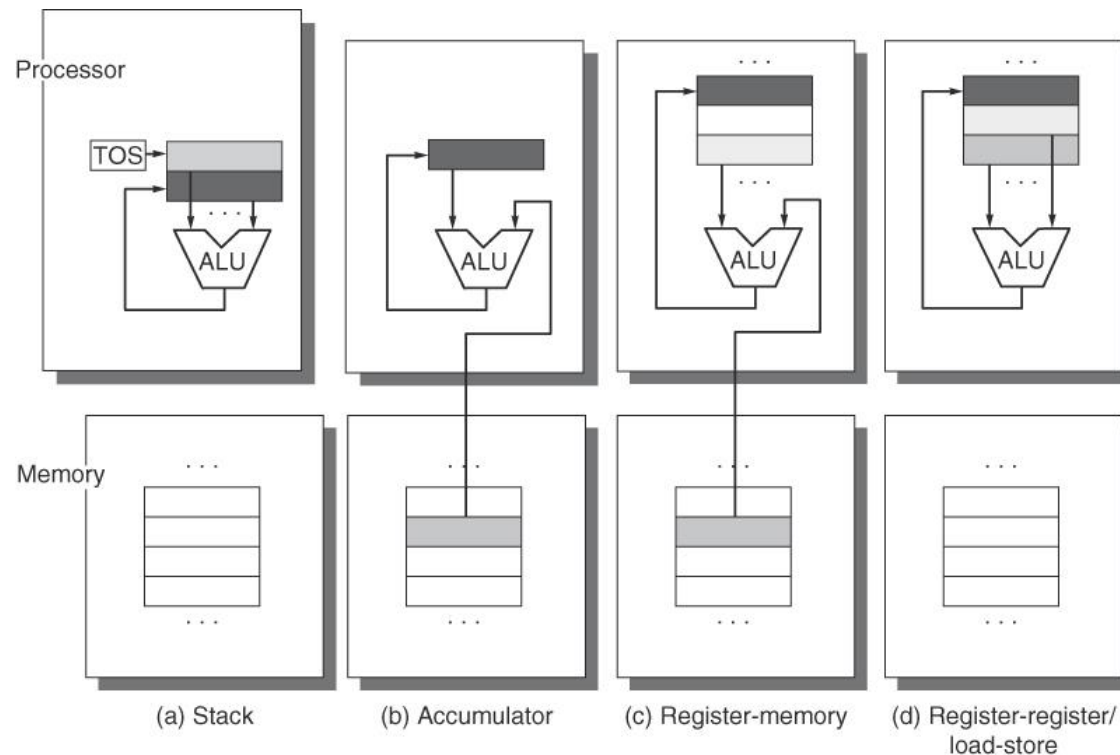
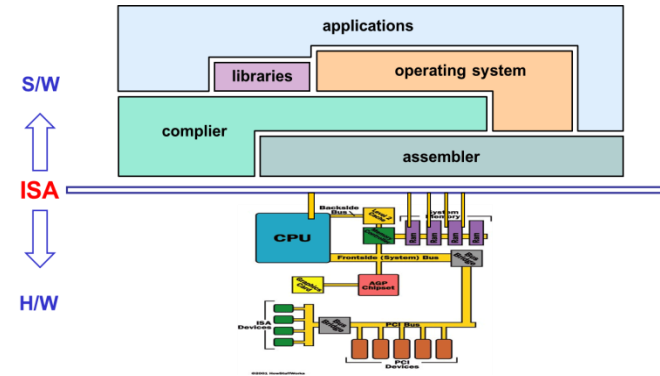
Instruction Set Architecture (ISA)

ISA classification

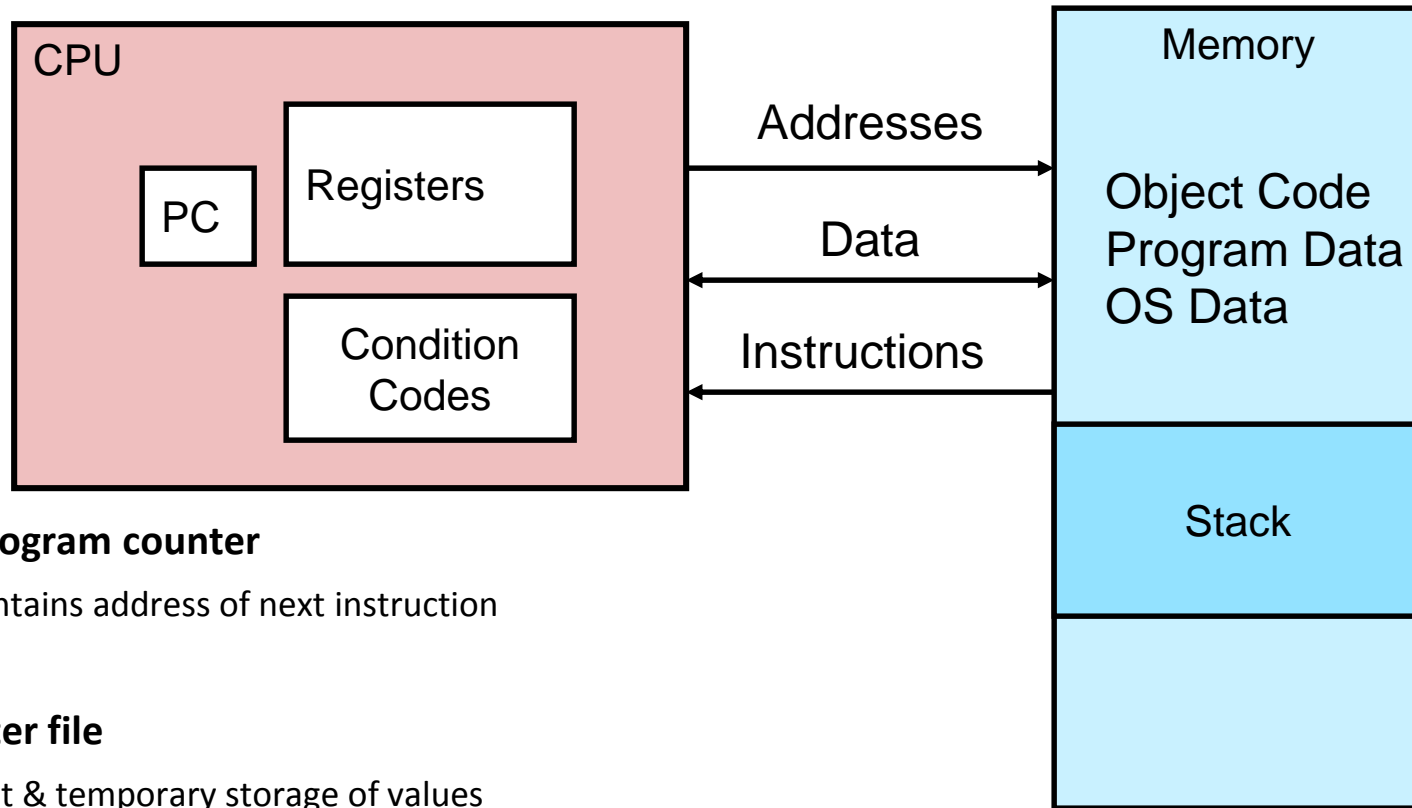
- RISC vs CISC — *Complex instruction set*
Reduce instruction set

- general-purpose
 - ▶ two or three operands?
 - ▶ # of memory operands?
 - ▶ classification
 - load-store
 - register-memory
 - memory-memory

- stack
- accumulator



Assembly Programmer's View



■ PC: Program counter

- Contains address of next instruction

■ Register file

- Fast & temporary storage of values

■ Condition codes

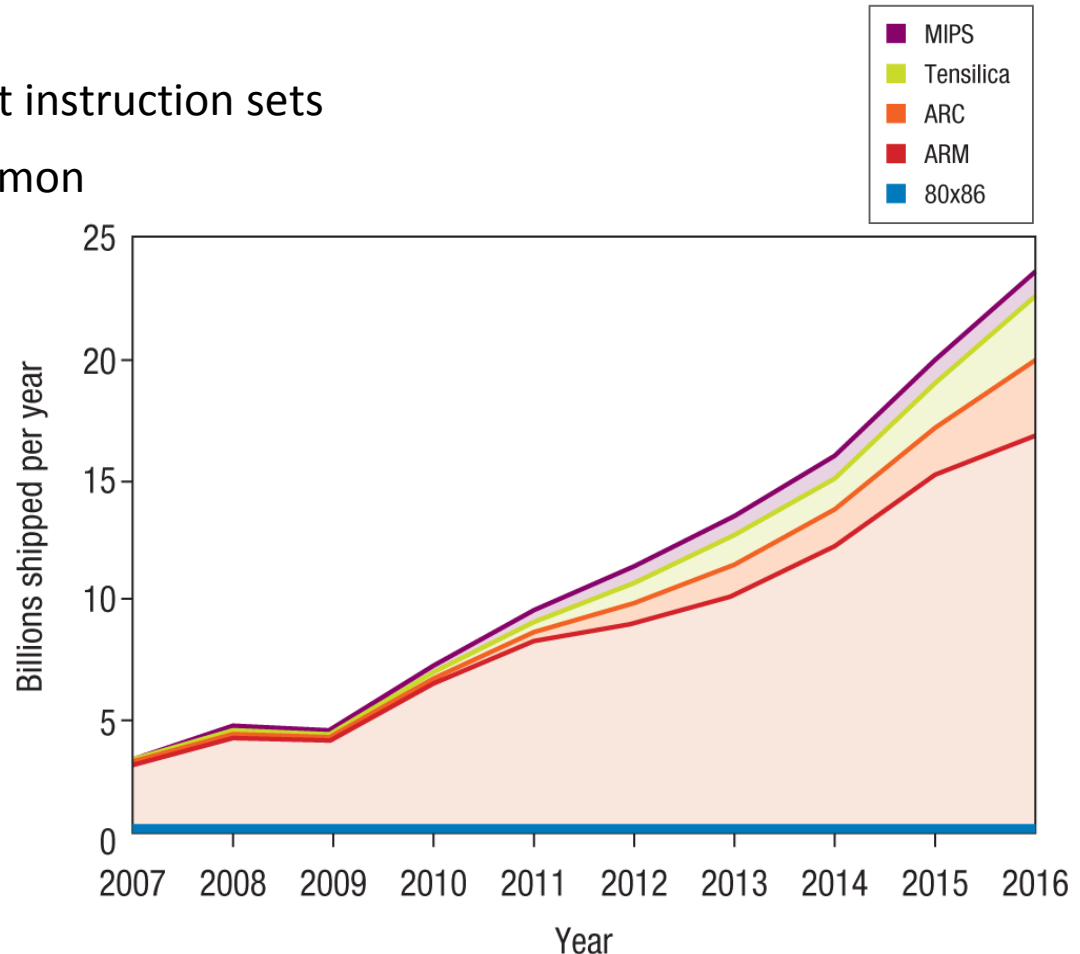
- Status information about most recent arithmetic operation
- Used for conditional branching

■ Memory

- Byte addressable array
- Code, user data, (some) OS data
- Includes stack to support procedures

The Instruction Set

- The repertoire of instructions of a computer
- Different computers have different instruction sets
 - But with many aspects in common
- Early computers had very simple instruction sets
 - Simplified implementation
- Many modern computers also have simple instruction sets
 - RISC (Reduced Instruction Set Computer)
- The Intel Architecture is an entirely different beast



Source: David Patterson, "Reduced Instruction Set Computers Then and Now," IEEE Computer, 2017.

The Instruction Set

- ARM processors (ARMv8-A A64)
 - basic instructions: ~300
 - with bells and whistles: ~1300
 - [ARMv8-A A64 instruction set](#)
- RISC-V
 - base integer instruction set: 47 instructions
 - with bells and whistles: <200
 - [RISC-V ISA specification](#)

Source: David Patterson, "Reduced Instruction Set Computers Then and Now," IEEE Computer, 2017.

The Instruction Set

■ Intel x86_64

- basic instructions: ~300
- with bells and whistles: ~4000
(includes different operand types)
 - ▶ Multimedia Extensions (MMX): 47 instructions
 - ▶ Streaming SIMD Extensions (SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2)
About $70 + 145 + 15 + 30 + 50 = 310$ new instructions
 - ▶ Advanced Vector Extensions (AVX)
 - AVX: ~100 instructions
 - AVX2: ~250 instructions
 - AVX-512: ~2,000 new instructions
- [Intel Architecture Reference Manual](#)

Source: David Patterson, "Reduced Instruction Set Computers Then and Now," IEEE Computer, 2017.



The RISC-V ISA

The RISC-V Instruction Set

- A completely open ISA that is freely available to academia and industry
- Fifth RISC ISA design developed at UC Berkeley
 - RISC-I (1981), RISC-II (1983), SOAR (1984), SPUR (1989), and RISC-V (2010)
- Now managed by the RISC-V Foundation (<http://riscv.org>)
- Typical of many modern ISAs
 - See RISC-V Reference Card (or Green Card)
- Similar ISAs have a large share of embedded core market
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...

Why Open the ISA at no Cost?

- Greater innovation via free-market competition
 - From many core designers, closed-source and open-source
- Shared open core designs
 - Shorter time to market, lower cost from reuse, fewer errors given more eyeballs, transparency makes it difficult for government agencies to add secret trap doors
- Processors becoming affordable for more devices
 - Help expand the Internet of Things (IoT), which could cost as little as \$1
- Software stack survive for long time
- Make architectural research and education more real
 - Fully open hardware and software stacks

Source: Krste Asanovic, RISC-V Tutorial at HPCA, 2015.

RISC-V ISAs

- Three base integer ISAs, one per address width
 - RV32I, RV64I, RV128I
 - RV32I: Only 40 instructions defined
 - RV32E: Reduced version of RV32I with 16 registers for embedded systems
- Standard extensions
 - Standard RISC encoding in a fixed 32-bit instruction format
 - C extension offers shorter 16-bit versions of common 32-bit RISC-V instructions (can be intermixed with 32-bit instructions)

Name	Extension
M	Integer Multiply/Divide
A	Atomic Instructions
F	Single-precision FP
D	Double-precision FP
G	General-purpose (= IMAFD)
Q	Quad-precision FP
C	Compressed Instructions

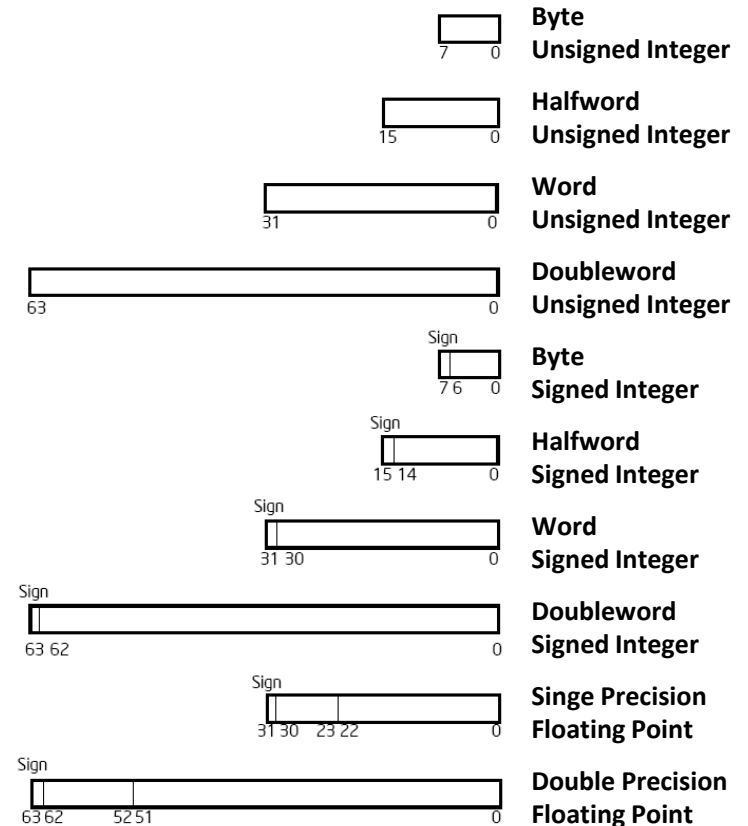
RISC-V Registers

#	Name	Usage
x0	zero	Hard-wired zero
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5	t0	Temporaries (Caller-save registers)
x6	t1	
x7	t2	
x8	s0/fp	Saved register / Frame pointer
x9	s1	Saved register
x10	a0	Function arguments / Return values
x11	a1	
x12	a2	Function arguments
x13	a3	
x14	a4	
x15	a5	

#	Name	Usage
x16	a6	Function arguments
x17	a7	
x18	s2	Saved registers (Callee-save registers)
x19	s3	
x20	s4	
x21	s5	
x22	s6	
x23	s7	
x24	s8	
x25	s9	
x26	s10	
x27	s11	
x28	t3	Temporaries (Caller-save registers)
x29	t4	
x30	t5	
x31	t6	
	pc	Program counter

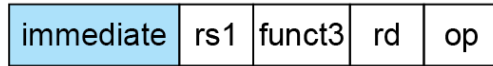
Data Types

- Integer data of 1, 2, 4, or 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4 or 8 bytes (with F or D extension)
 - Just contiguously allocated bytes in memory
- No aggregated types such as arrays or structures
 - Just contiguously allocated bytes in memory

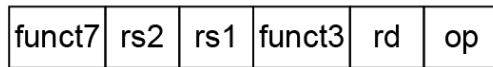


Operand Types: RISC-V Addressing

1. Immediate addressing



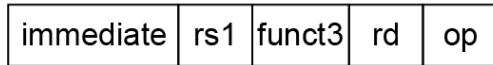
2. Register addressing



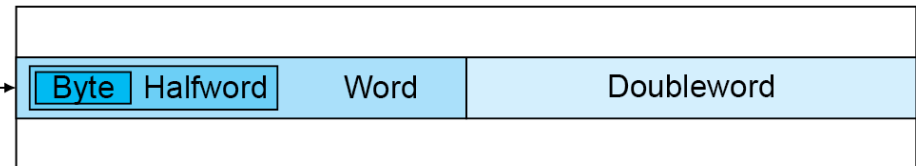
Registers

Register

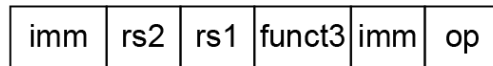
3. Base addressing



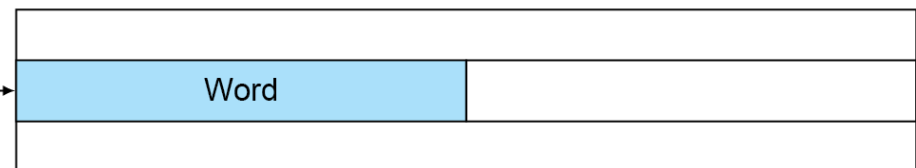
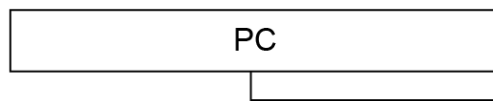
Memory



4. PC-relative addressing



Memory



Operations

- Perform an arithmetic or logical function on register data
- Transfer data between memory and register
 - Load data from memory into register
 - Store register data into memory
- Transfer control
 - Unconditional jump
 - Conditional branch
 - Procedure call and return

Comparison

- ARMv8 (AArch64)
 - 32 general purpose registers
 - 1 reserved, 1 zero

- Intel IA32 (32-bit)
 - 8 general purpose registers
 - 2 of which have predetermined role

- Intel x86_64 (64-bit)
 - 16 general purpose registers
 - 1 of which is reserved

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```



```
sum:
    add    a0,a0,a1
    ret
```

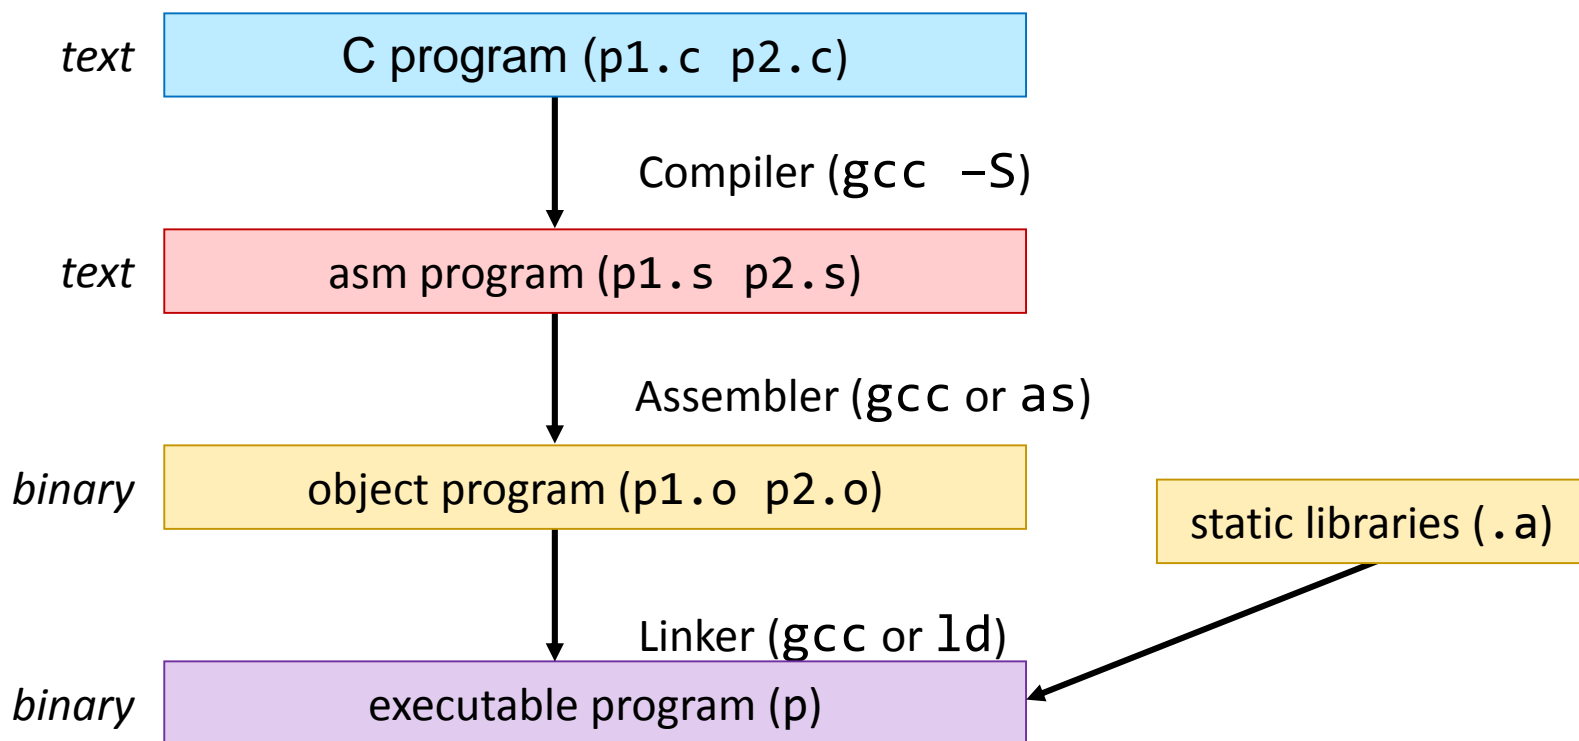
The RISC-V Instruction Set Architecture:

First Steps

From C to Machine Code

Prepend all commands with
`risc64-unknown-elf-`
in the VM (cross-compiler)!

- Code in files **p1.c** **p2.c**
- Compile with command: **gcc -O p1.c p2.c -o p**
 - Use optimizations (**-O**)
 - Put resulting binary in file **p**



Compiling to 32-bit RISC-V

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

`$ risc64-unknown-elf-gcc -march=rv32i -mabi=ilp32 -O -S sum.c`
produces the file `sum.s`:

```
sum:
    add    a0,a0,a1
    ret
```

→ code generated by `gcc -march=rv32i -mabi=ilp32 -O0 -S sum.c` ?

Compiling to 64-bit RISC-V

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

`$ risc64-unknown-elf-gcc -march=rv64i -mabi=lp64 -O -S sum.c`
produces the file `sum.s`:

```
sum:
    addw    a0,a0,a1
    ret
```

Compiling to 64-bit Intel (x86_64)

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

\$ gcc -march=x86-64 -O -S sum.c

produces the file sum.s:

```
sum:
.LFB0:
    .cfi_startproc
    leal    (%rdi,%rsi), %eax
    ret
    .cfi_endproc
```

→ code generated by gcc -march=x86-64 -O0 -S sum.c ?

Compiling to 32-bit Intel (i386)

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

\$ gcc -m32 -march=i386 -O -S sum.c

produces the file sum.s:

```
sum:
.LFB0:
    .cfi_startproc
    movl    8(%esp), %eax
    addl    4(%esp), %eax
    ret
    .cfi_endproc
```


Object Code

Machine code for sum (32-bit RISC-V)

0x1014c <sum>:

0x00

0xb5

0x05

0x33

0x00

0x00

0x80

0x67

- 8 bytes total
- Each instruction 4 bytes
- Starts at address 0x1014c

■ Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

■ Linker

- Resolves references between files
- Combines with static run-time libraries
 - ▶ E.g., code for malloc, printf
- Some libraries are *dynamically linked*
 - ▶ Linking occurs when program begins execution

Object Code

64-bit RISC-V

0x10154 <sum>:

0x00

0xb5

0x05

0x3b

0x00

0x00

0x80

0x67

- 8 bytes total
- Each instruction 4 bytes
- Starts at address 0x1014c

64-bit Intel

0x113b <sum>:

0x8d

0x04

0x37

0xc3

0x90

- 5 bytes total
- Three instructions with 3, 1, 1 bytes
- Starts at address 0x113b

Machine Instruction Example – RISC-V

```
int t = x+y;
```

■ C Code

- add two signed integers

```
add    a0, a0, a1
```

■ Assembly

- add computes the sum of two registers and places the result in another register
- Operands:
 - x: register **a0**
 - y: register **a1**
 - t: register **a0**
(return value also in **a0**)

```
0x1014c:  00b50533
```

■ Object Code

- 4-byte instruction
- Stored at address **0x1014c**

Machine Instruction Example – IA32

```
int t = x+y;
```

■ C Code

- add two signed integers

```
addl 8(%ebp),%eax
```

■ Assembly

- `addl` adds the 4-byte value located at memory address `%ebp+8` to register `%eax`
- Operands:

x:	register	%eax
y:	memory	mem[%ebp+8]
t:	register	%eax

(return value also in **%eax**)

Intel code similar to expression:

`x += y`

More precisely:

```
int eax;
```

```
int *ebp;
```

```
eax += mem[ebp + 8]
```

```
0x0804809a:  03 45 08
```

■ Object Code

- 3-byte instruction
- Stored at address **0x0804809a**

Disassembling Object Code

```
$ riscv64-unknown-elf-gcc -march=rv32i -mabi=ilp32 -O -o p main.c sum.c  
$ riscv64-unknown-elf-objdump -d p32
```

```
0001014c <sum>:  
    1014c:    00b50533      add    a0,a0,a1  
    10150:    00008067      ret
```

■ Disassembler (objdump)

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces *approximate* rendition of assembly code
- Can be run on either a.out (complete executable) or .o file

Disassembling Object Code – IA32

```
$ gcc -m32 -c -O sum.c  
$ objdump -d sum.o
```

```
00000000 <sum>:
```

0:	55	push	%ebp
1:	89 e5	mov	%esp,%ebp
3:	8b 45 0c	mov	0xc(%ebp),%eax
6:	03 45 08	add	0x8(%ebp),%eax
9:	5d	pop	%ebp
a:	c3	ret	

Alternate Disassembly using gdb

- Within **gdb** Debugger

```
$ gcc -m32 -O -o p main.c sum.c
```

```
$ riscv64-unknown-elf-gdb p32
```

```
GNU gdb (GDB) 8.3.0.20190516-git
...
Reading symbols from p32...
(No debugging symbols found in p32)
(gdb) disassemble sum
Dump of assembler code for function sum:
    0x0001014c <+0>:      add      a0,a0,a1
    0x00010150 <+4>:      ret
End of assembler dump.
(gdb) x/8xb sum
0x1014c <sum>:  0x33 0x05 0xb5 0x00 0x67 0x80 0x00 0x00
(gdb)
```

What Can be Disassembled?

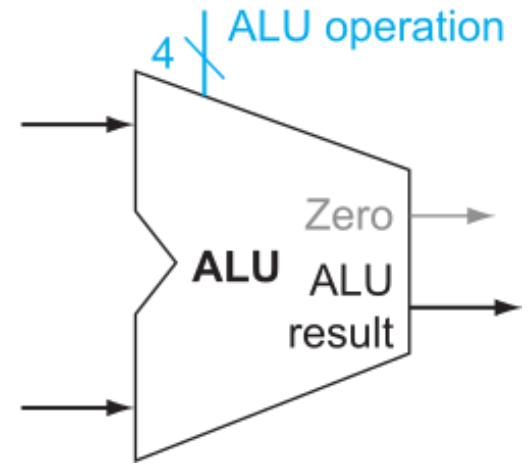
```
% objdump -d WINWORD.EXE

WINWORD.EXE:  file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:  55                push    %ebp
30001001:  8b ec            mov     %esp,%ebp
30001003:  6a ff            push    $0xffffffff
30001005:  68 90 10 00 30    push    $0x30001090
3000100a:  68 91 dc 4c 30    push    $0x304cdc91
```

- (Almost) anything that can be interpreted as executable code (not everything makes sense, of course)
- Disassembler examines bytes and reconstructs assembly source



RISC-V

Arithmetic and Logical Operations

Arithmetic Operations

- Add and subtract: three operands
 - Two sources and one destination

add a, b, c	// a ← b + c
sub a, b, c	// a ← b - c

- All arithmetic operations have this form

- **Design Principle 1: Simplicity favors regularity**

- Regularity makes implementation simpler
- Simplicity enables higher performance at lower cost

Register Operands

- Arithmetic instructions use register operands
- RISC-V has a 32 x 64-bit register file: $x0 \sim x31$
 - Use for frequently accessed data
 - 64-bit data is called a “doubleword”
 - 32-bit data is called a “word”
- **Design Principle 2: Smaller is faster**
 - cf. Main memory: millions of locations

Register Operand Example

C code:

```
// f in x19
// g in x20
// h in x21
// i in x22
// j in x23

f = (g + h) - (i + j);
```

Compiled RISC-V code:

```
add  x5, x20, x21
add  x6, x22, x23
sub  x19, x5, x6
```

Registers vs. Memory

- Registers are faster to access than memory
- In RISC-V, data in memory cannot be directly addressed by ALU instructions
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler tries to use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Immediate Operands

- Constant data specified in an instruction

```
addi x22, x22, 4
```

- Make the common case fast
 - Small constants are common (limited to 12 bits, signed)
 - Immediate operand avoids a load instruction

Arithmetic Operations

Instruction	Type	Example	Meaning
Add	R	add rd, rs1, rs2	$R[rd] = R[rs1] + R[rs2]$
Subtract	R	sub rd, rs1, rs2	$R[rd] = R[rs1] - R[rs2]$
Add immediate	I	addi rd, rs1, imm12	$R[rd] = R[rs1] + \text{SignExt}(\text{imm12})$
Set less than	R	slt rd, rs1, rs2	$R[rd] = (R[rs1] < R[rs2])? 1 : 0$
Set less than immediate	I	slti rd, rs1, imm12	$R[rd] = (R[rs1] < \text{SignExt}(\text{imm12}))? 1 : 0$
Set less than unsigned	R	sltu rd, rs1, rs2	$R[rd] = (R[rs1] <_u R[rs2])? 1 : 0$
Set less than immediate unsigned	I	sltiu rd, rs1, imm12	$R[rd] = (R[rs1] <_u \text{SignExt}(\text{imm12}))? 1 : 0$
Load upper immediate	U	lui rd, imm20	$R[rd] = \text{SignExt}(\text{imm20} \ll 12)$
Add upper immediate to PC	U	auipc rd, imm20	$R[rd] = \text{PC} + \text{SignExt}(\text{imm20} \ll 12)$

Example: arith

x in a0
y in a1
z in a2

```
long arith (long x,  
            long y,  
            long z) {  
    long t1 = x + y;  
    long t2 = z + t1;  
    long t3 = x + 4;  
    long t4 = y * 48;  
    long t5 = t3 + t4;  
    long rv = t2 - t5;  
    return rv;  
}
```

```
arith:  
    add    a5, a0, a1    # a5 = x + y (t1)  
    add    a2, a5, a2    # a2 = t1 + z (t2)  
    addi   a0, a0, 4     # a0 = x + 4 (t3)  
    slli   a5, a1, 1     # a5 = y * 2  
    add    a1, a5, a1    # a1 = a5 + y  
    slli   a5, a1, 4     # a5 = a1 * 16 (t4)  
    add    a0, a0, a5    # a0 = t3 + t4 (t5)  
    sub    a0, a2, a0    # a0 = t2 - t5 (rval)  
    ret
```


Logical Operations

■ Instructions for bitwise manipulation

Operation	C	Java	RISC-V
Shift left	<<	<<	sll, slli
Shift right (arithmetic)	>>	>>	sra, srai
Shift right (logical)	>>	>>>	srl, srli
Bit-by-bit AND	&	&	and, andi
Bit-by-bit OR			or, ori
Bit-by-bit XOR	^	^	xor, xori
Bit-by-bit NOT	~	~	

■ Useful for extracting and inserting groups of bits in a word

AND Operations

- Useful to mask bits in a word
 - Select some bits, clears others to 0

and x9, x10, x11

x10 00000000 00000000 00000000 00000000 00000000 00000000 00001101 11000000

x11 00000000 00000000 00000000 00000000 00000000 00000000 00111100 00000000

x9 00000000 00000000 00000000 00000000 00000000 00000000 00001100 00000000

OR Operation

- Useful to include bits in a word
 - Set some bits to 1, leave others unchanged

```
or    x9, x10, x11
```

x10	00000000 00000000 00000000 00000000 00000000 00000000 00001101 11000000
-----	---

x11	00000000 00000000 00000000 00000000 00000000 00000000 00111100 00000000
-----	---

x9	00000000 00000000 00000000 00000000 00000000 00000000 00111101 11000000
----	---

XOR Operation

- Differencing operation
 - Clear when bits are the same, set if they are different

```
xor x9, x10, x12
```

x10	00000000	00000000	00000000	00000000	00000000	00000000	00001101	11000000
-----	----------	----------	----------	----------	----------	----------	----------	----------

x12	11111111	11111111	11111111	11111111	11111111	11111111	11111111	11111111
-----	----------	----------	----------	----------	----------	----------	----------	----------

x9	11111111	11111111	11111111	11111111	11111111	11111111	11110010	00111111
----	----------	----------	----------	----------	----------	----------	----------	----------

32-bit Constants

- Most constants are small
 - 12-bit immediate is sufficient
- For the occasional 32-bit constant:
 - Copies 20-bit constant to bits [31:12] of rd
 - Extends bit 31 to bits [63:32]
 - Clears bits [11:0] of rd to 0

```
lui    rd, constant
```

- Example: `x19 <- 0x003D0500`

<code>lui x19, 0x003D0</code>	0000 0000 0000 0000	0000 0000 0000 0000	0000 0000 0011 1101 0000	0000 0000 0000
<code>addi x19,x19,0x500</code>	0000 0000 0000 0000	0000 0000 0000 0000	0000 0000 0011 1101 0000	0101 0000 0000

Logical Operations

Instruction	Type	Example	Meaning
AND	R	and rd, rs1, rs2	$R[rd] = R[rs1] \& R[rs2]$
OR	R	or rd, rs1, rs2	$R[rd] = R[rs1] R[rs2]$
XOR	R	xor rd, rs1, rs2	$R[rd] = R[rs1] \wedge R[rs2]$
AND immediate	I	andi rd, rs1, imm12	$R[rd] = R[rs1] \& \text{SignExt}(\text{imm12})$
OR immediate	I	ori rd, rs1, imm12	$R[rd] = R[rs1] \text{SignExt}(\text{imm12})$
XOR immediate	I	xori rd, rs1, imm12	$R[rd] = R[rs1] \wedge \text{SignExt}(\text{imm12})$
Shift left logical	R	sll rd, rs1, rs2	$R[rd] = R[rs1] \ll R[rs2]$
Shift right logical	R	srl rd, rs1, rs2	$R[rd] = R[rs1] \gg R[rs2]$ (logical)
Shift right arithmetic	R	sra rd, rs1, rs2	$R[rd] = R[rs1] \gg R[rs2]$ (arithmetic)
Shift left logical immediate	I	slli rd, rs1, shamt	$R[rd] = R[rs1] \ll \text{shamt}$
Shift right logical imm.	I	srli rd, rs1, shamt	$R[rd] = R[rs1] \gg \text{shamt}$ (logical)
Shift right arithmetic immediate	I	srai rd, rs1, shamt	$R[rd] = R[rs1] \gg \text{shamt}$ (arithmetic)

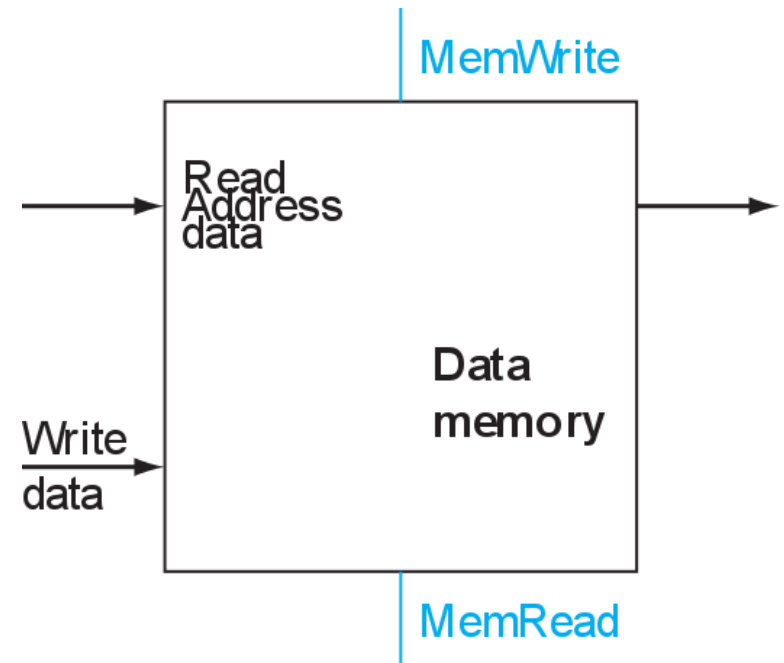
Example: logical

x in a0
y in a1

```
long logical (long x, long y)
{
    long t1 = x ^ y;
    long t2 = t1 >> 17;
    long mask = (1 << 8) - 7;
    long rval = t2 & mask;
    return rval;
}
```

logical:

```
xor    a0, a0, a1 # a0 = x ^ y (t1)
srai   a0, a0, 17 # a0 = t1 >> 17 (t2)
andi   a0, a0, 249 # a0 = t2 & ((1 << 8) - 7)
ret
```



RISC-V

Data Transfer Operations

Memory Operands

- Main memory used for composite data
 - Arrays, structures, dynamic data
- To apply arithmetic operations
 - Load values from memory to registers
 - Store result from register to memory
- Memory is byte addressed: each address identifies an 8-bit byte
- RISC-V is Little Endian
 - Least-significant byte at least address of a word
- RISC-V does not require words to be aligned in memory
 - Unlike some other ISAs

Memory Operand Example

C code:

```
// h in x21
// base address of A in x22

A[12] = h + A[8]
```

Compiled RISC-V code:

```
// 8 bytes per doubleword
// &A[8] = A + 64

ld    x9, 64(x22)
add   x9, x21, x9
sd    x9, 96(x22)
```

Byte/Halfword/Word Operations

- Load byte/halfword/word:
Sign extend to 64 bits in rd

lb	rd, offset(rs1)
lh	rd, offset(rs1)
lw	rd, offset(rs1)

- Load byte/halfword/word:
Zero extend to 64 bits in rd

lbu	rd, offset(rs1)
lhu	rd, offset(rs1)
lwu	rd, offset(rs1)

- Store byte/halfword/word:
Store rightmost 8/16/32 bits

sb	rs2, offset(rs1)
sh	rs2, offset(rs1)
sw	rs2, offset(rs1)

Data Transfer Operations

Instruction	Type	Example	Meaning
Load doubleword	I	ld rd, imm12(rs1)	$R[rd] = \text{Mem}_8[R[rs1] + \text{SignExt}(\text{imm12})]$
Load word	I	lw rd, imm12(rs1)	$R[rd] = \text{SignExt}(\text{Mem}_4[R[rs1] + \text{SignExt}(\text{imm12})])$
Load halfword	I	lh rd, imm12(rs1)	$R[rd] = \text{SignExt}(\text{Mem}_2[R[rs1] + \text{SignExt}(\text{imm12})])$
Load byte	I	lb rd, imm12(rs1)	$R[rd] = \text{SignExt}(\text{Mem}_1[R[rs1] + \text{SignExt}(\text{imm12})])$
Load word unsigned	I	lwu rd, imm12(rs1)	$R[rd] = \text{ZeroExt}(\text{Mem}_4[R[rs1] + \text{SignExt}(\text{imm12})])$
Load halfword unsigned	I	lhu rd, imm12(rs1)	$R[rd] = \text{ZeroExt}(\text{Mem}_2[R[rs1] + \text{SignExt}(\text{imm12})])$
Load byte unsigned	I	lbu rd, imm12(rs1)	$R[rd] = \text{ZeroExt}(\text{Mem}_1[R[rs1] + \text{SignExt}(\text{imm12})])$
Store doubleword	S	sd rs2, imm12(rs1)	$\text{Mem}_8[R[rs1] + \text{SignExt}(\text{imm12})] = R[rs2]$
Store word	S	sw rs2, imm12(rs1)	$\text{Mem}_4[R[rs1] + \text{SignExt}(\text{imm12})] = R[rs2](31:0)$
Store halfword	S	sh rs2, imm12(rs1)	$\text{Mem}_2[R[rs1] + \text{SignExt}(\text{imm12})] = R[rs2](15:0)$
Store byte	S	sb rs2, imm12(rs1)	$\text{Mem}_1[R[rs1] + \text{SignExt}(\text{imm12})] = R[rs2](7:0)$

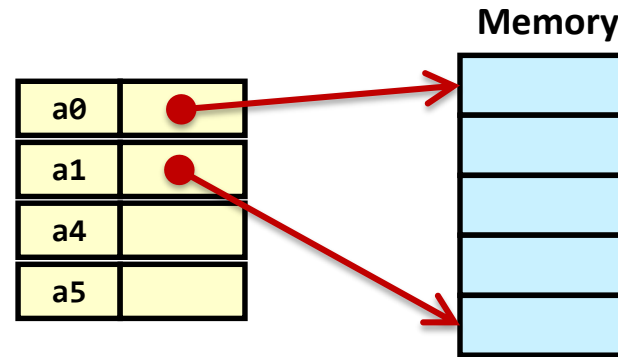
Swap Example

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    ld    a4, 0(a0)
    ld    a5, 0(a1)
    sd    a5, 0(a0)
    sd    a4, 0(a1)
    ret
```

Understanding Swap (1)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

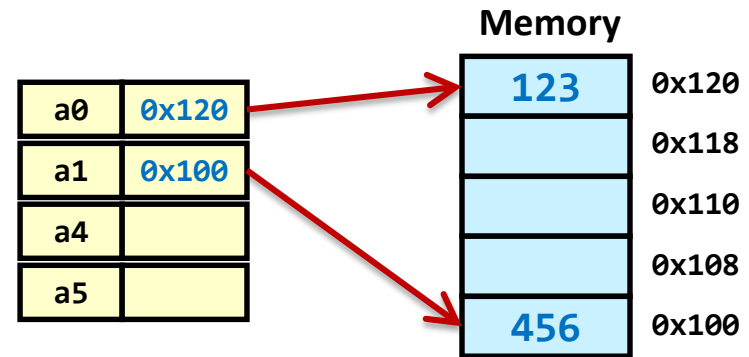
Register	Variable
<code>a0</code>	<code>xp</code>
<code>a1</code>	<code>yp</code>
<code>a4</code>	<code>t0</code>
<code>a5</code>	<code>t1</code>

swap:

<code>ld</code>	<code>a4, 0(a0)</code>	<code># t0 = *xp</code>
<code>ld</code>	<code>a5, 0(a1)</code>	<code># t1 = *yp</code>
<code>sd</code>	<code>a5, 0(a0)</code>	<code># *xp = t1</code>
<code>sd</code>	<code>a4, 0(a1)</code>	<code># *yp = t0</code>
<code>ret</code>		

Understanding Swap (2)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

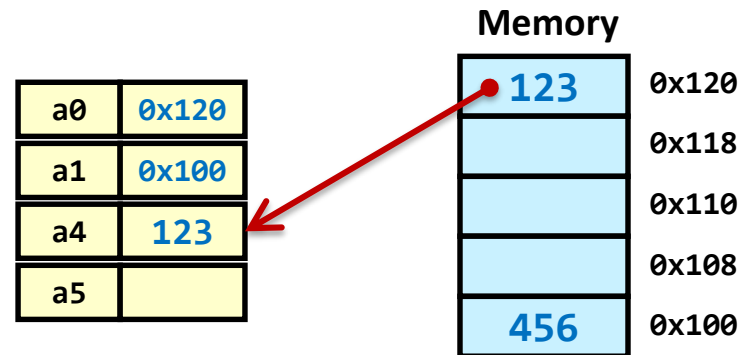
Register	Variable
a0	xp
a1	yp
a4	t0
a5	t1

swap:

ld	a4, 0(a0)	# t0 = *xp
ld	a5, 0(a1)	# t1 = *yp
sd	a5, 0(a0)	# *xp = t1
sd	a4, 0(a1)	# *yp = t0
ret		

Understanding Swap (3)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

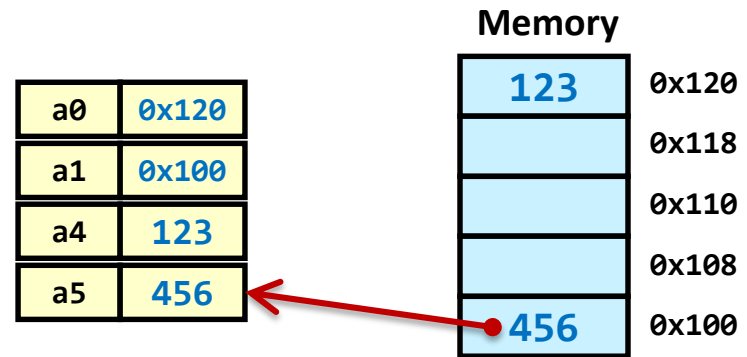
Register	Variable
a0	xp
a1	yp
a4	t0
a5	t1

swap:

ld	a4, 0(a0)	# t0 = *xp
ld	a5, 0(a1)	# t1 = *yp
sd	a5, 0(a0)	# *xp = t1
sd	a4, 0(a1)	# *yp = t0
ret		

Understanding Swap (4)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

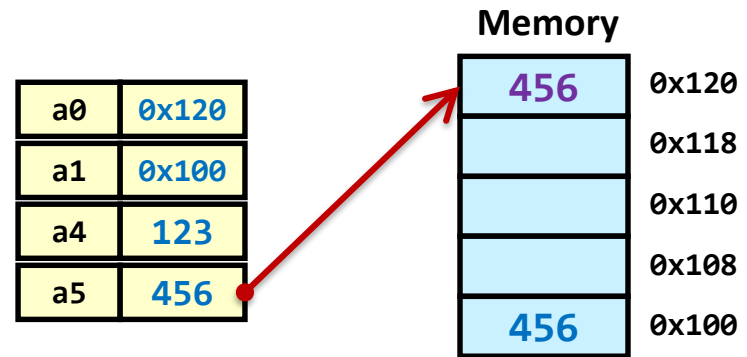
Register	Variable
a0	xp
a1	yp
a4	t0
a5	t1

swap:

ld	a4, 0(a0)	# t0 = *xp
ld	a5, 0(a1)	# t1 = *yp
sd	a5, 0(a0)	# *xp = t1
sd	a4, 0(a1)	# *yp = t0
ret		

Understanding Swap (5)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

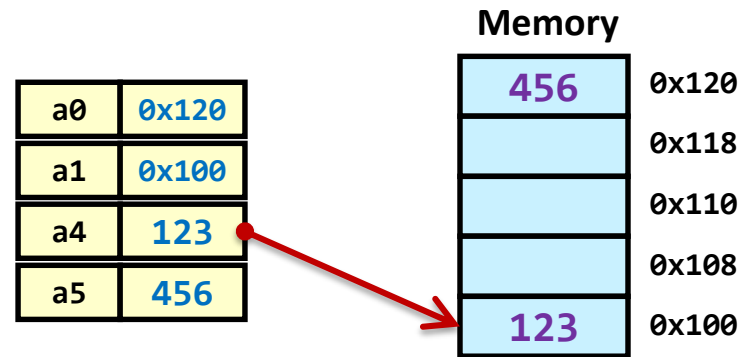
Register	Variable
a0	xp
a1	yp
a4	t0
a5	t1

swap:

ld	a4, 0(a0)	# t0 = *xp
ld	a5, 0(a1)	# t1 = *yp
sd	a5, 0(a0)	# *xp = t1
sd	a4, 0(a1)	# *yp = t0
ret		

Understanding Swap (6)

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Allocation
(by compiler)

Register	Variable
a0	xp
a1	yp
a4	t0
a5	t1

swap:

ld	a4, 0(a0)	# t0 = *xp
ld	a5, 0(a1)	# t1 = *yp
sd	a5, 0(a0)	# *xp = t1
sd	a4, 0(a1)	# *yp = t0
ret		

String Copy Example

```
void strcpy(char x[],
            char y[])
{
    size_t i;

    i = 0;
    while
        ((x[i]=y[i])!='\0')
        i += 1;
}
```

strcpy:

```
    ; x is in a0
    ; y is in a1
    ; i is allocated in t3

    add    t3,zero,zero    ; t3 <- 0
L1:  add    t0,t3,a1        ; t0 <- &y[i]
    lbu    t1,0(t0)        ; t1 <- y[i]
    add    t2,t3,a0        ; t2 <- &x[i]
    sb     t1,0(t2)        ; x[i] <- y[i]
    beq    t1,zero,L2      ; if (t1==0), goto L2
    addi   t3,t3,1         ; t3 += 1
    j      L1              ; goto L1
L2:  ret                  ; return
```