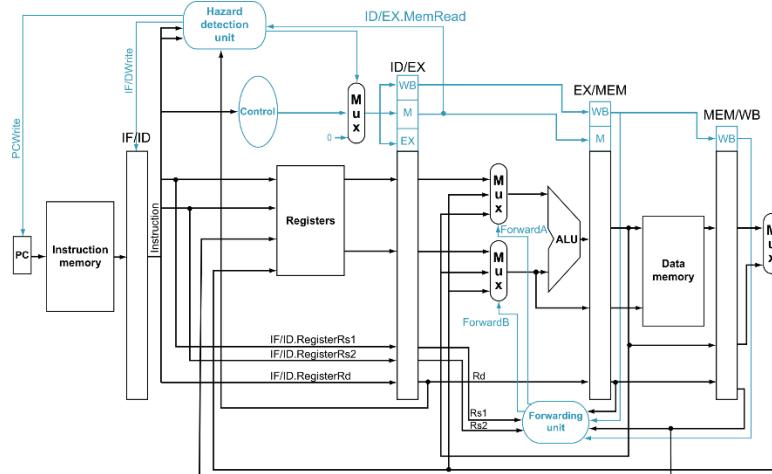


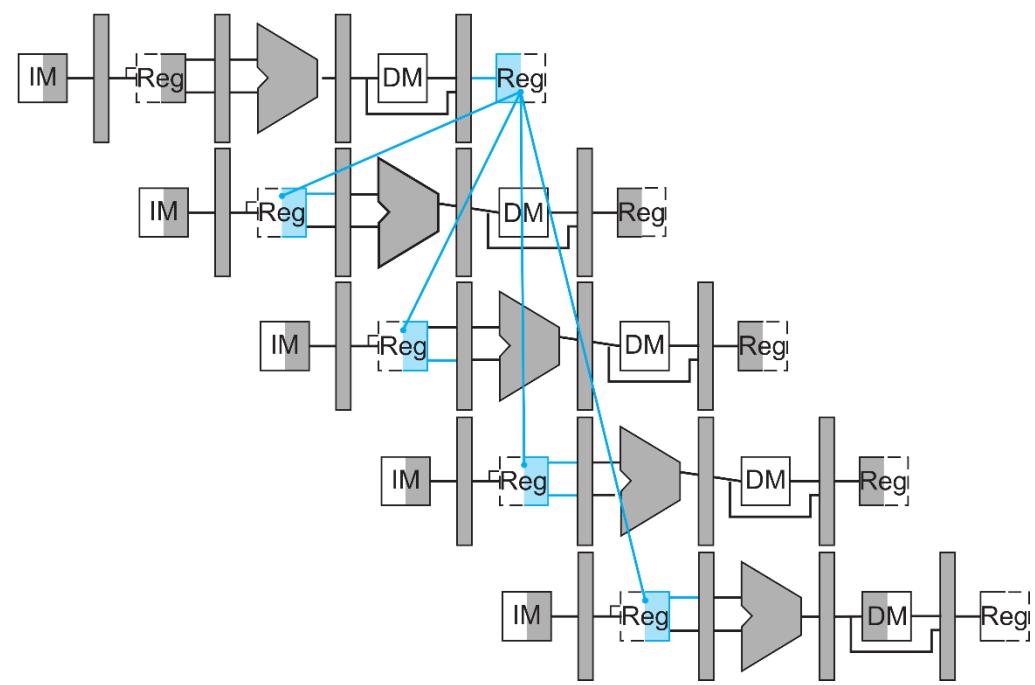
Processor Architecture

Pipelined RISC-V Implementation (Part II)



Module Outline

- Data Hazards
- Control Hazards
- Exceptions
- Module Summary

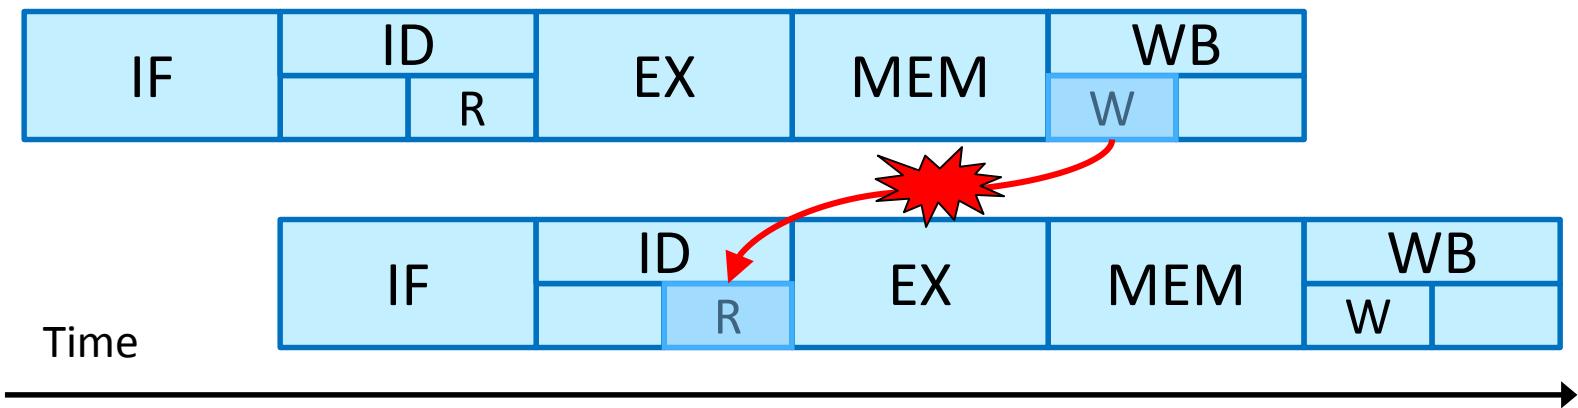


Data Hazards

Recap: Data Hazard

- Caused by data dependencies (arithmetic or load)

OP i add x_5, x_{11}, x_{12}
OP $i+1$ sub x_7, x_5, x_7



Information cannot flow
backwards in time!

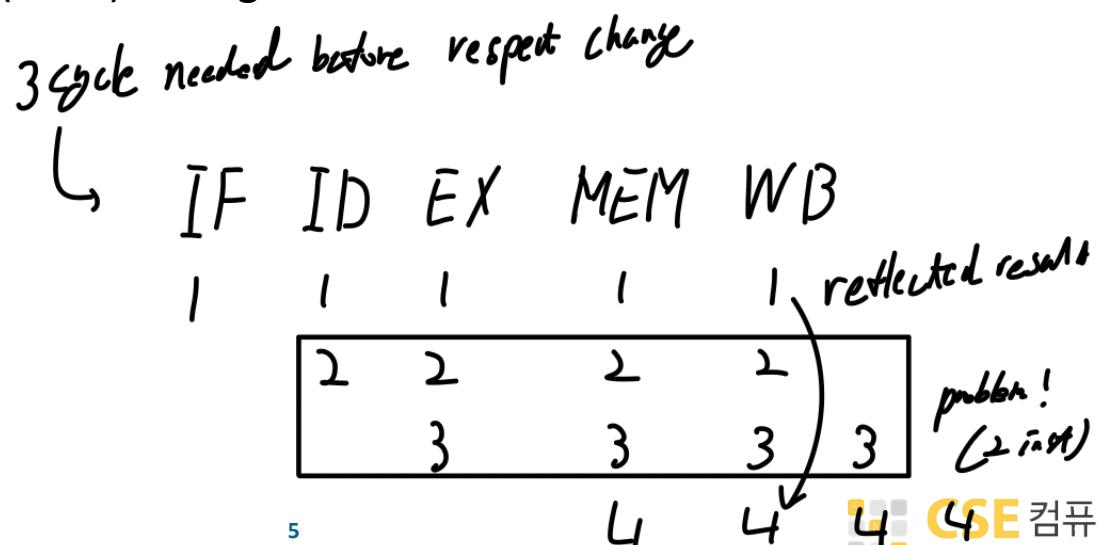
Data Hazards in ALU Instructions

- Consider this code sequence:

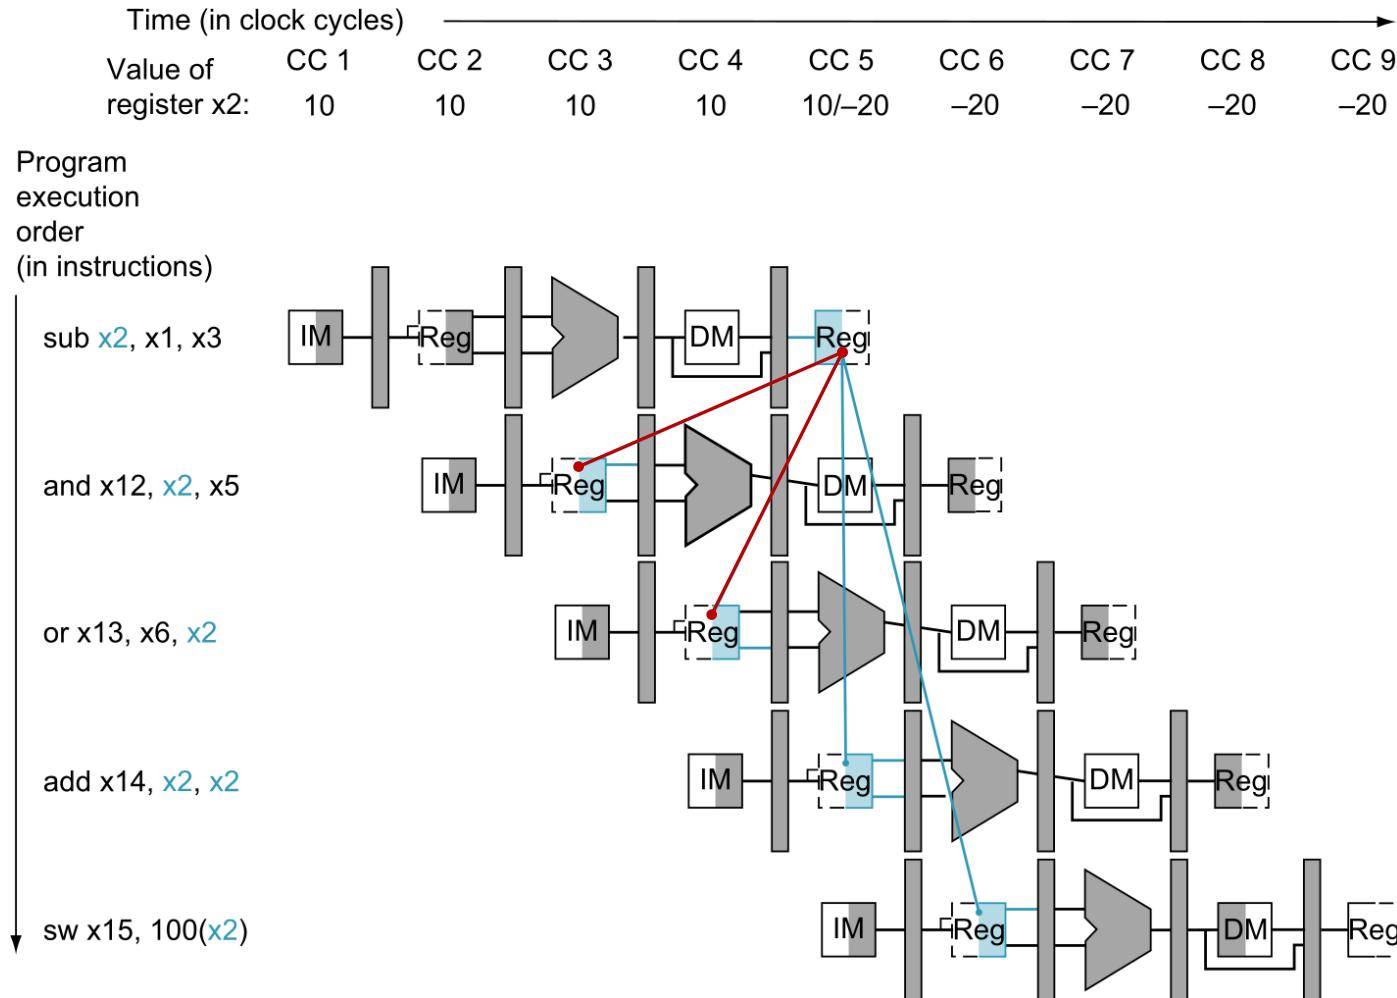
```
sub x2, x1, x3  
and x12, x2, x5  
or x13, x6, x2  
add x14, x2, x2  
sw x15, 100(x2)
```

- Several data dependencies (RAW) on register x2

```
sub x2, x1, x3  
and x12, x2, x5  
or x13, x6, x2  
add x14, x2, x2  
sw x15, 100(x2)
```



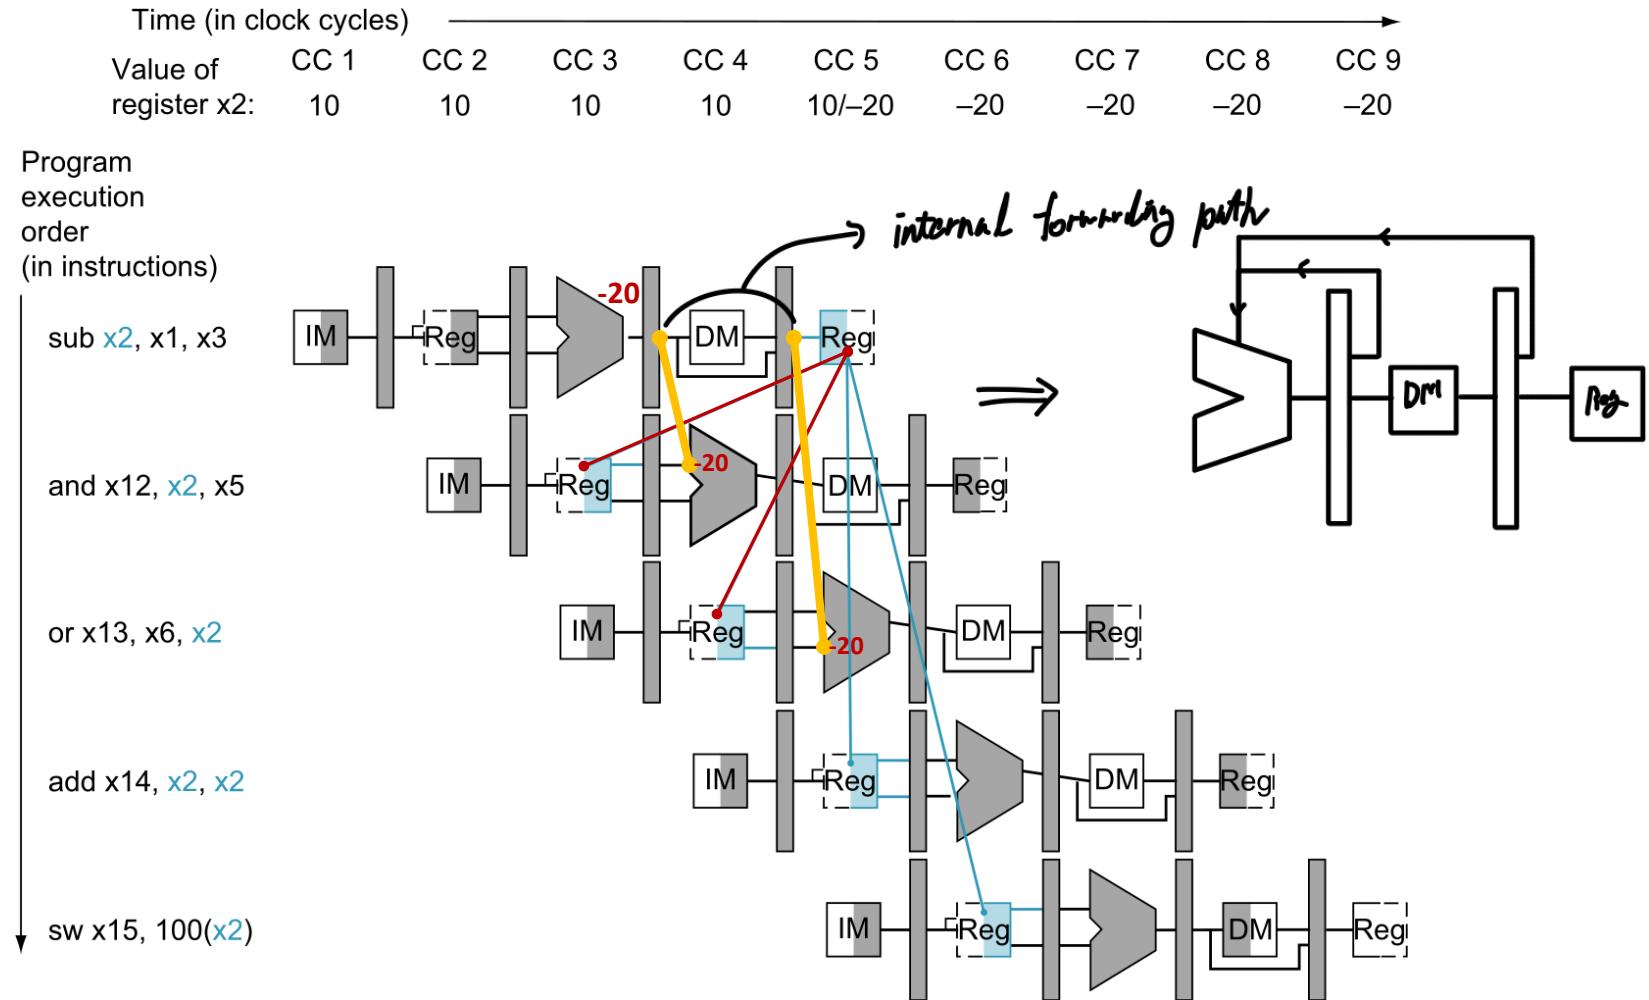
Data Hazards in ALU Instructions



- How do we detect data hazards in the pipeline and how do we fix them?

Data Hazards in ALU Instructions

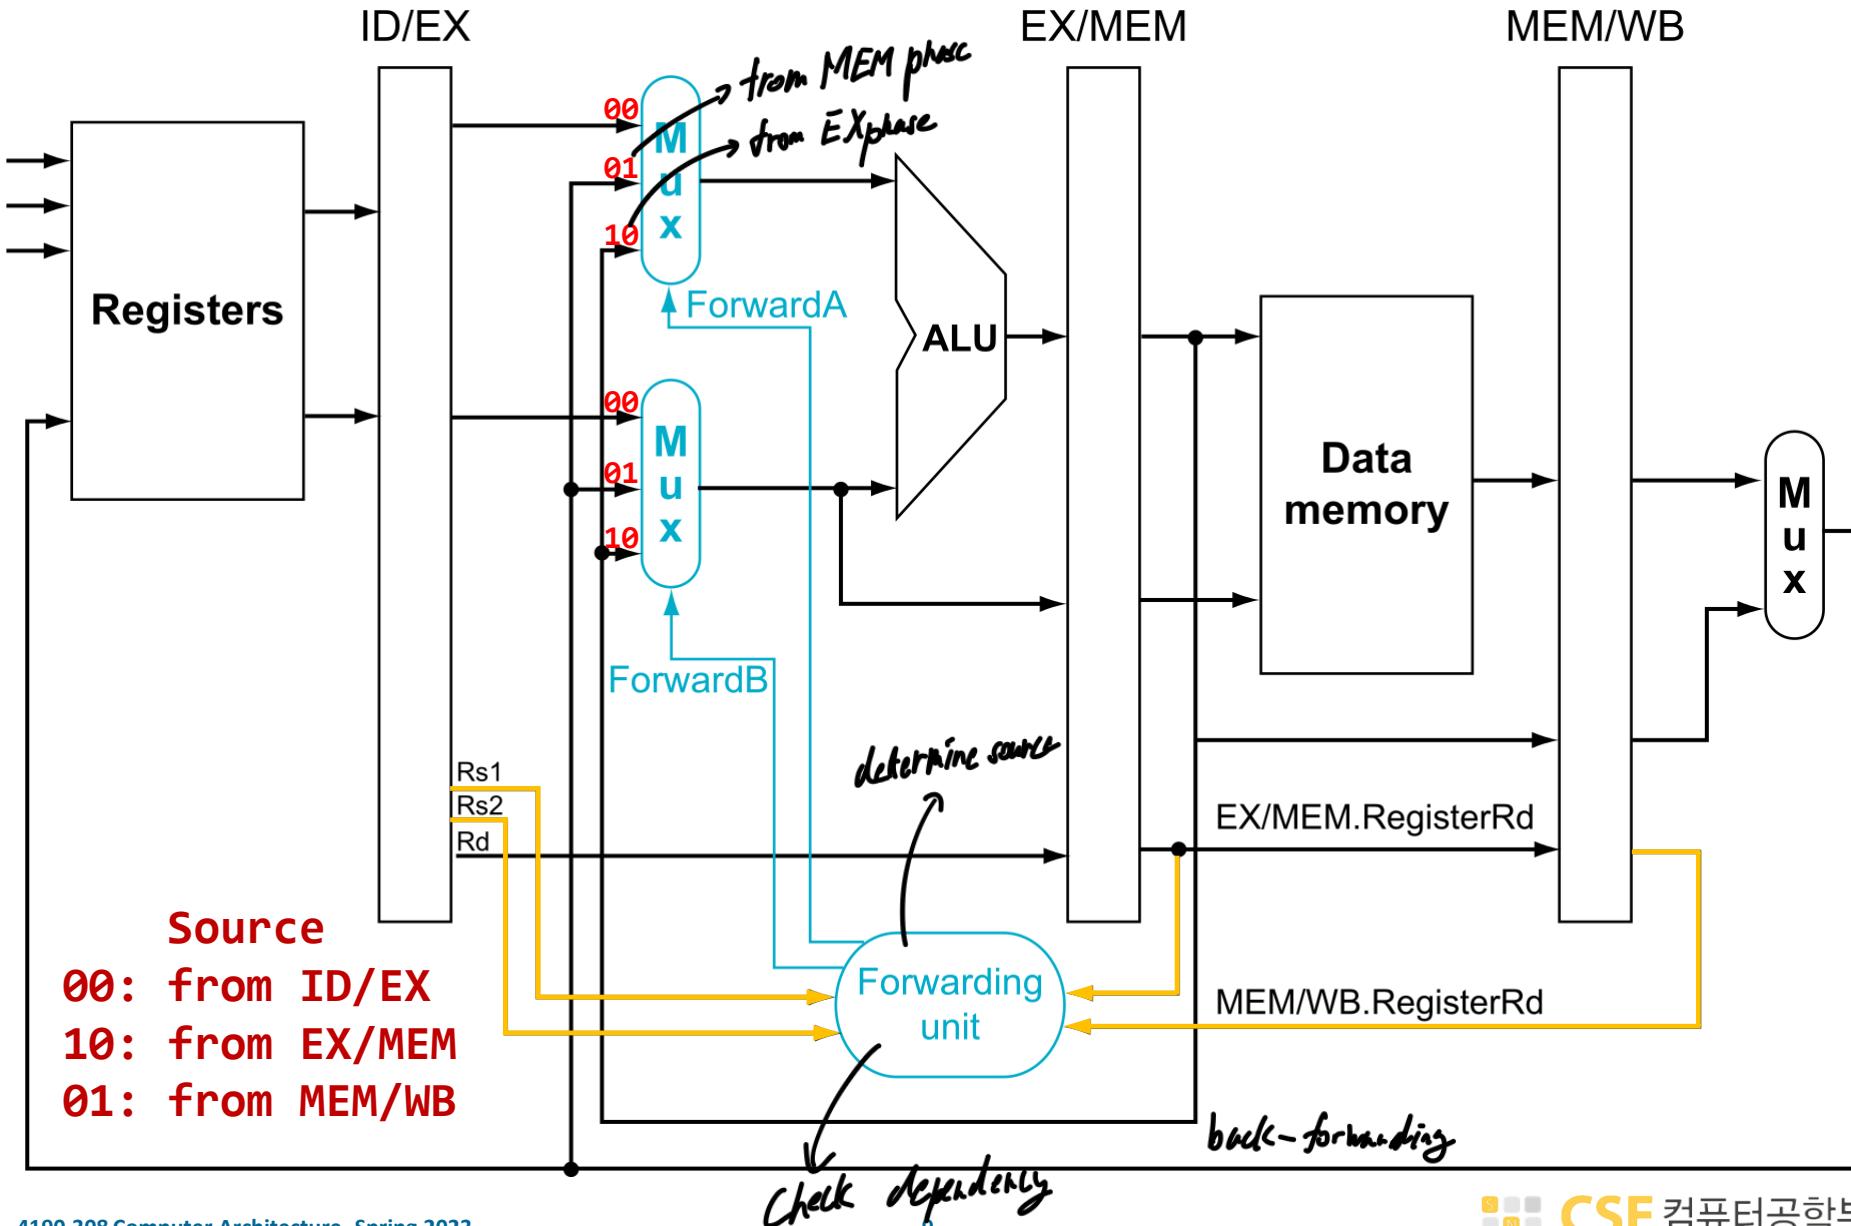
- The updated value of x_2 is available as soon as sub has passed the EX stage



Forwarding Paths and Control

- **Pass source register indices along in pipeline** (destination register already passed on)
 - Pass register source 1 and register source 2 indices from ID to EX stage
 - Naming
 - ▶ ID/EX.RegisterRs1 = register # for Rs1 sitting in ID/EX pipeline register
 - ▶ ID/EX.RegisterRs2 = register # for Rs2 sitting in ID/EX pipeline register

Forwarding Paths and Control



Forwarding Control

MUX control	Source	Example
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

Detecting the Need to Forward

Condition for data hazards

EX/MEM.RegisterRd = ID/EX.RegisterRs1

EX/MEM.RegisterRd = ID/EX.RegisterRs2

MEM/WB.RegisterRd = ID/EX.RegisterRs1

MEM/WB.RegisterRd = ID/EX.RegisterRs2

Arithmetic data hazard

Forward from EX/MEM
pipeline register

Forward from MEM/WB
pipeline register

Memory data hazard

Detecting the Need to Forward (cont'd)

- But only if forwarding instruction will write to a register

```
EX/MEM.RegWrite = 1,  
MEM/WB.RegWrite = 1
```

- And only if Rd for that instruction is not x0

```
EX/MEM.RegisterRd ≠ 0  
MEM/WB.RegisterRd ≠ 0
```

Forwarding Conditions : Computing inside Forwarding unit $\Rightarrow (rd == rs1/2) \& (\text{regwrite}) \& (!\text{regrd})$

■ EX hazard

```
if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
    and (EX/MEM.RegisterRd = ID/EX.RegisterRs1))  
    forwardA = 10  
if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)  
    and (EX/MEM.RegisterRd = ID/EX.RegisterRs2))  
    forwardB = 10
```

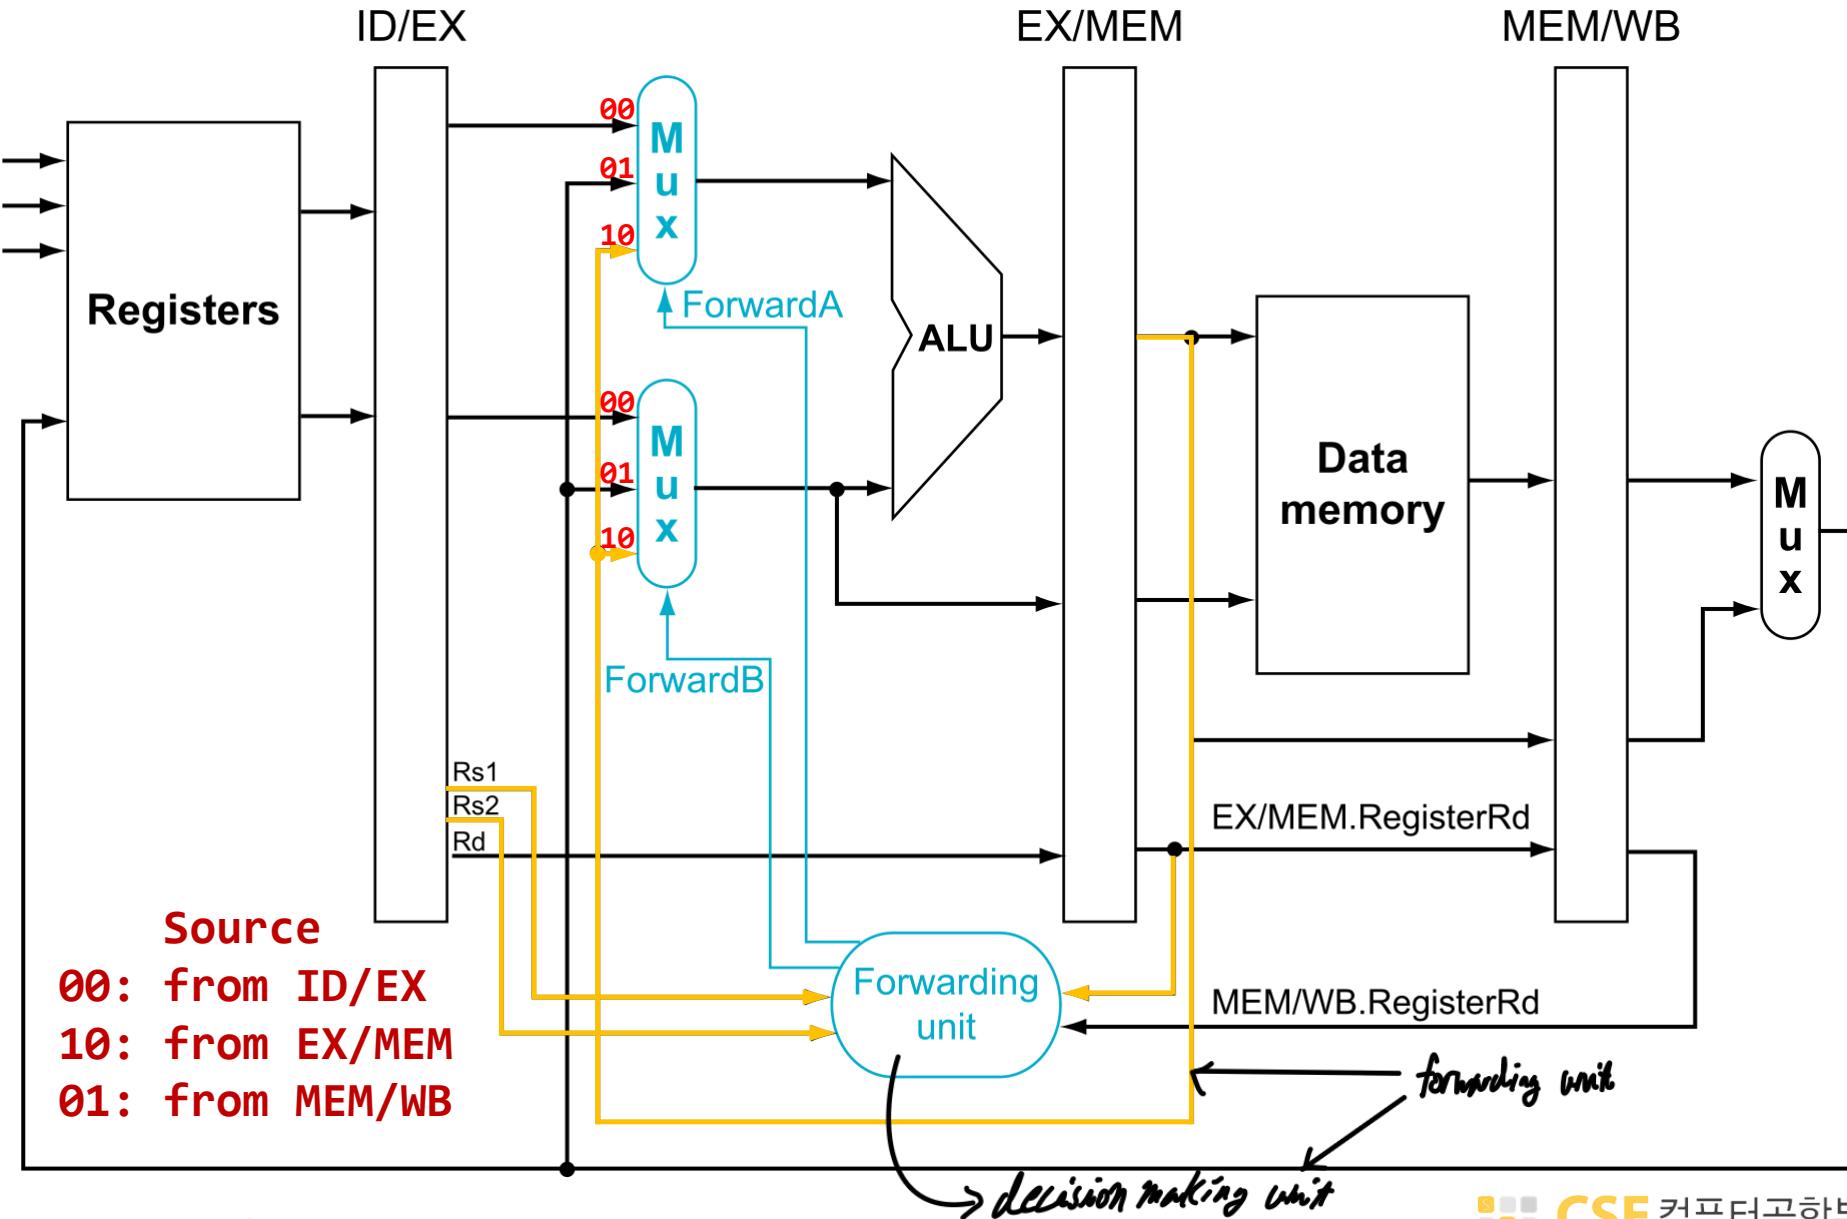
Completed

■ MEM hazard

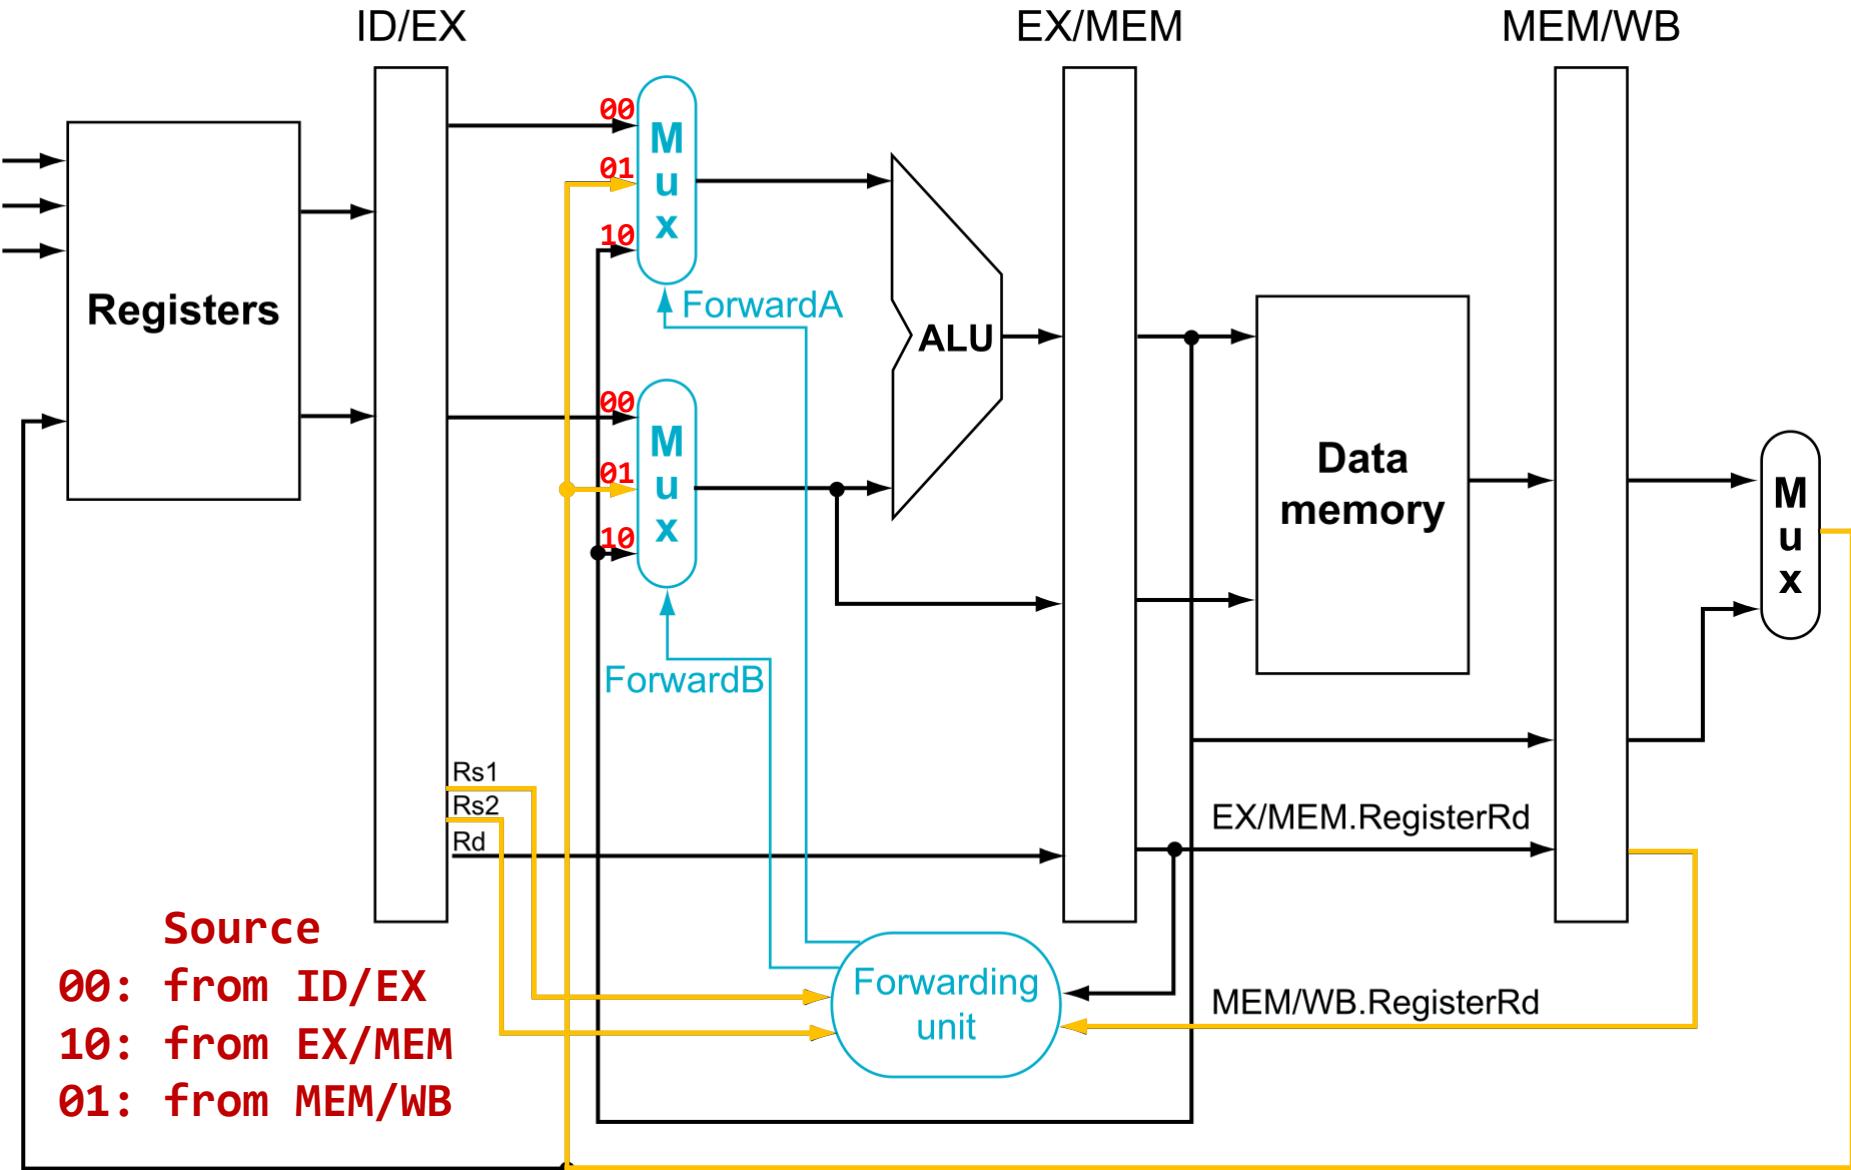
```
if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs1))  
    forwardA = 01  
if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs2))  
    forwardB = 01
```

Incomplete

Forwarding from EX/MEM



Forwarding from MEM/WB



Simultaneous Data Hazards

- Consider this sequence:



add	x1	x1, x2	①
add	x1	x1, x3	②
add	x1,	x1, x4	③

→ can refer to

ret of ① in MEM/WB Reg

or

ret of ② in EX/MEM Reg

⇒ have to use recent!

. : Using ② for ② operation

- Both hazards occur simultaneously

- Want to mimic serial execution

→ Use most recent result

- Revise MEM hazard condition

- Only forward if EX hazard condition isn't true

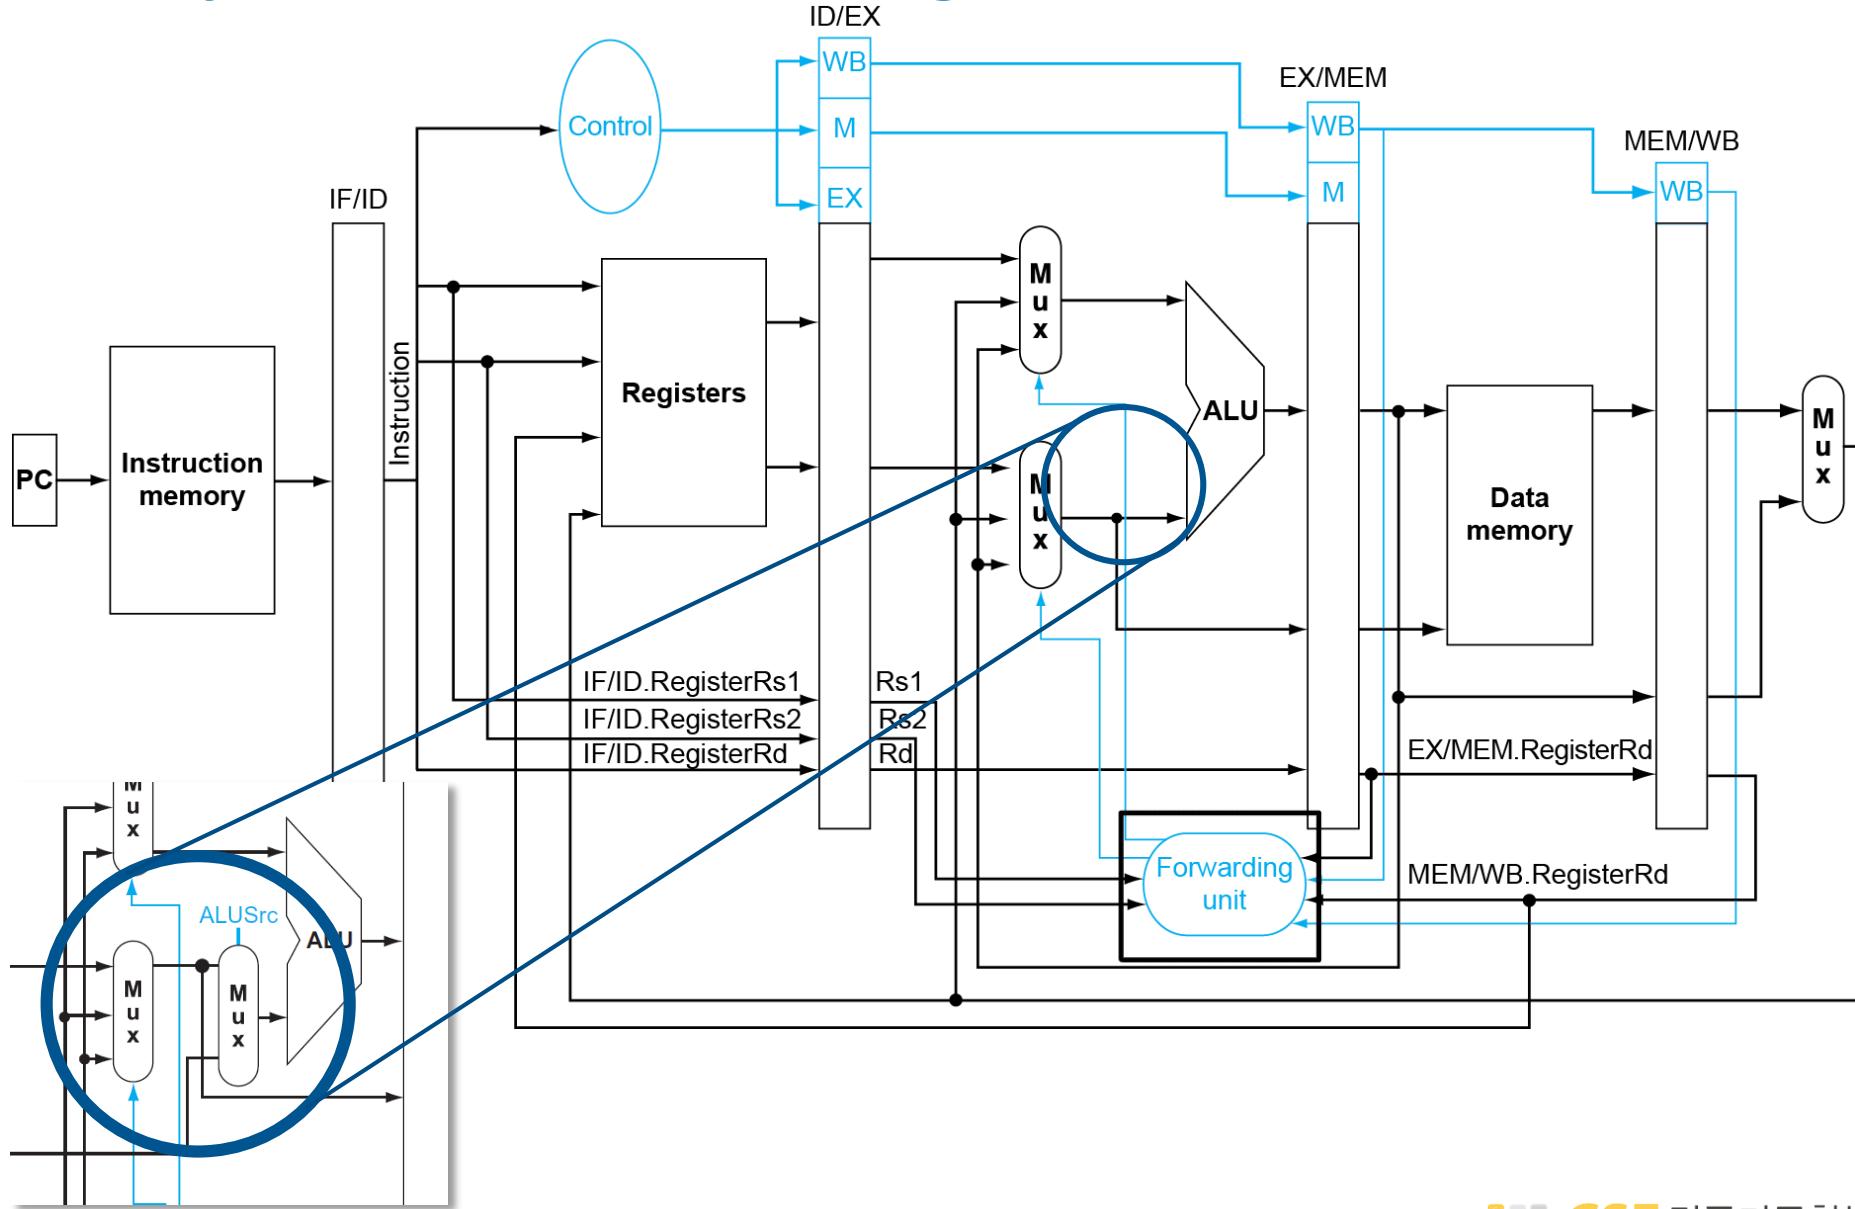
Revised Forwarding Conditions

■ MEM hazard

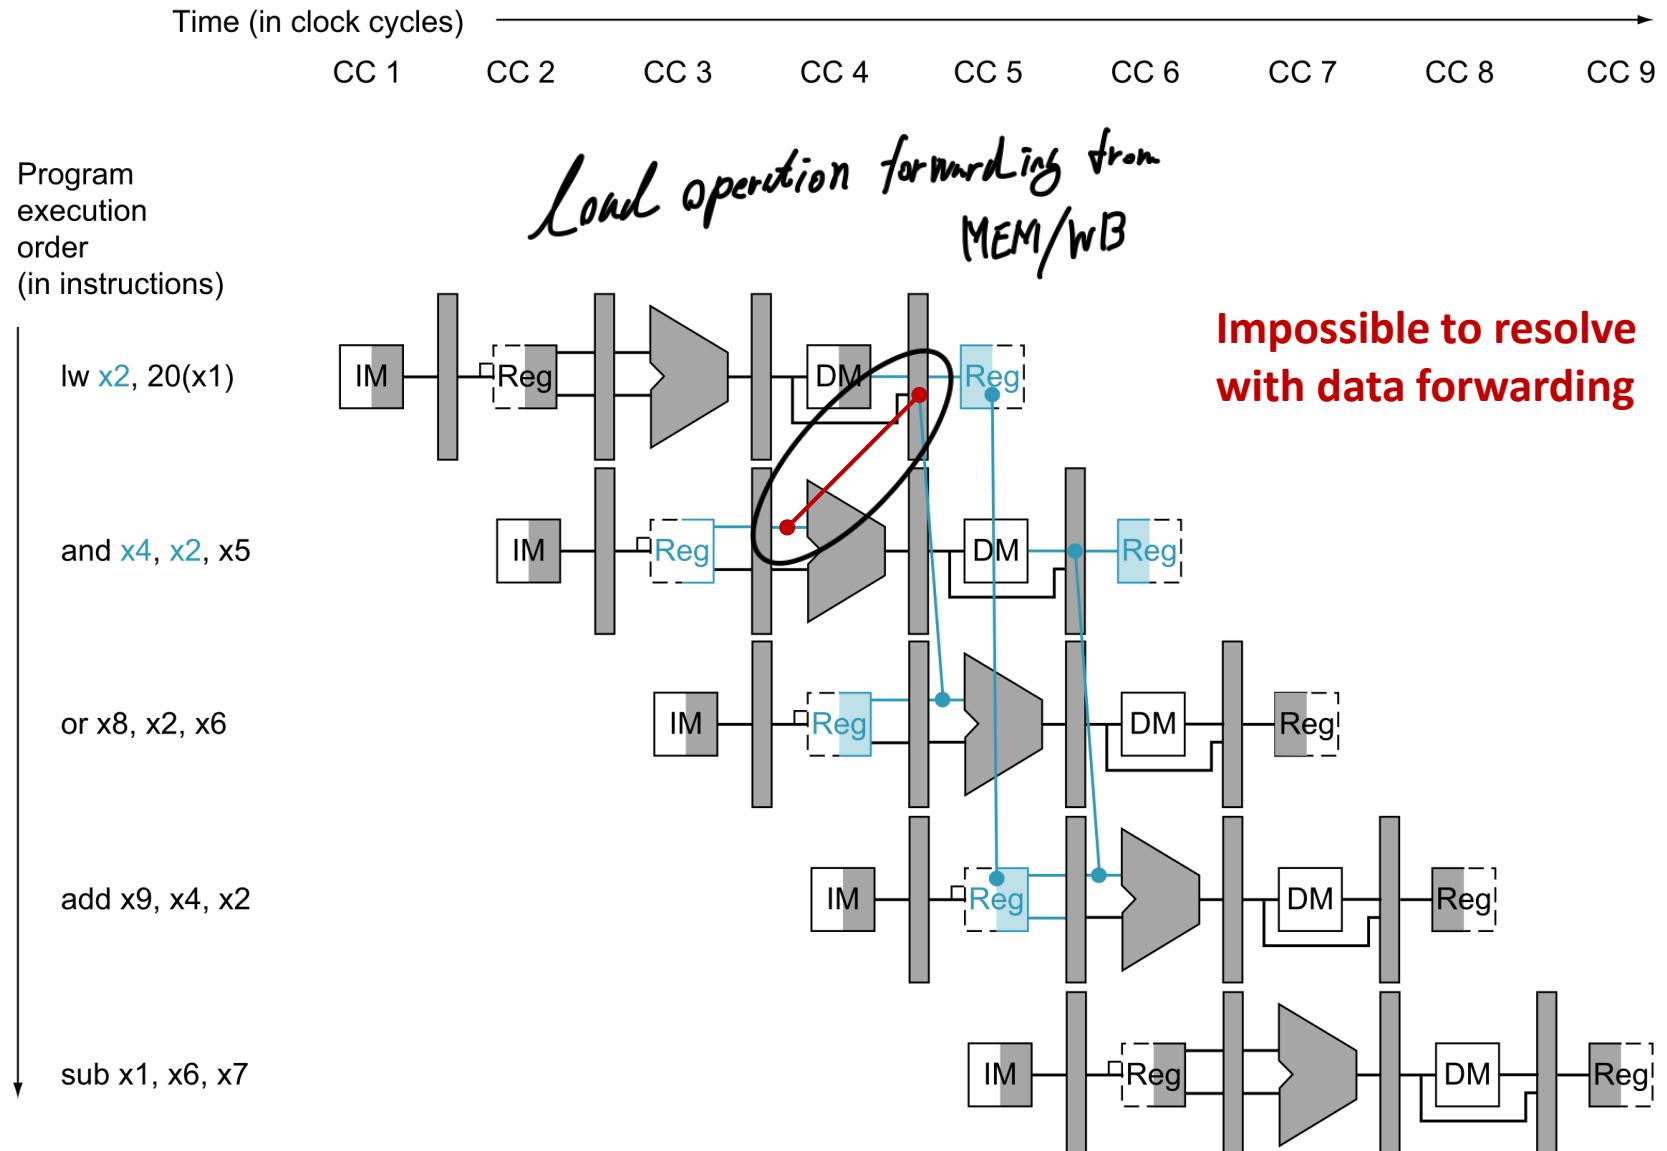
Check forwarding

```
if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd ≠ 0))  
    and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0))  
        and (EX/MEM.RegisterRd = ID/EX.RegisterRs1))  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs1))  
forwardA = 01  
Check MEM data hazard occur  
  
if (MEM/WB.RegWrite  
and (MEM/WB.RegisterRd != 0)  
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0))  
    and (EX/MEM.RegisterRd = ID/EX.RegisterRs2))  
    and (MEM/WB.RegisterRd = ID/EX.RegisterRs2))  
forwardB = 01
```

Datapath with Forwarding



Load-Use Hazard (Load Interlock)



Load-Use Hazard Resolution

■ Action

- Impossible to solve with forwarding
- Stall pipeline, insert bubble

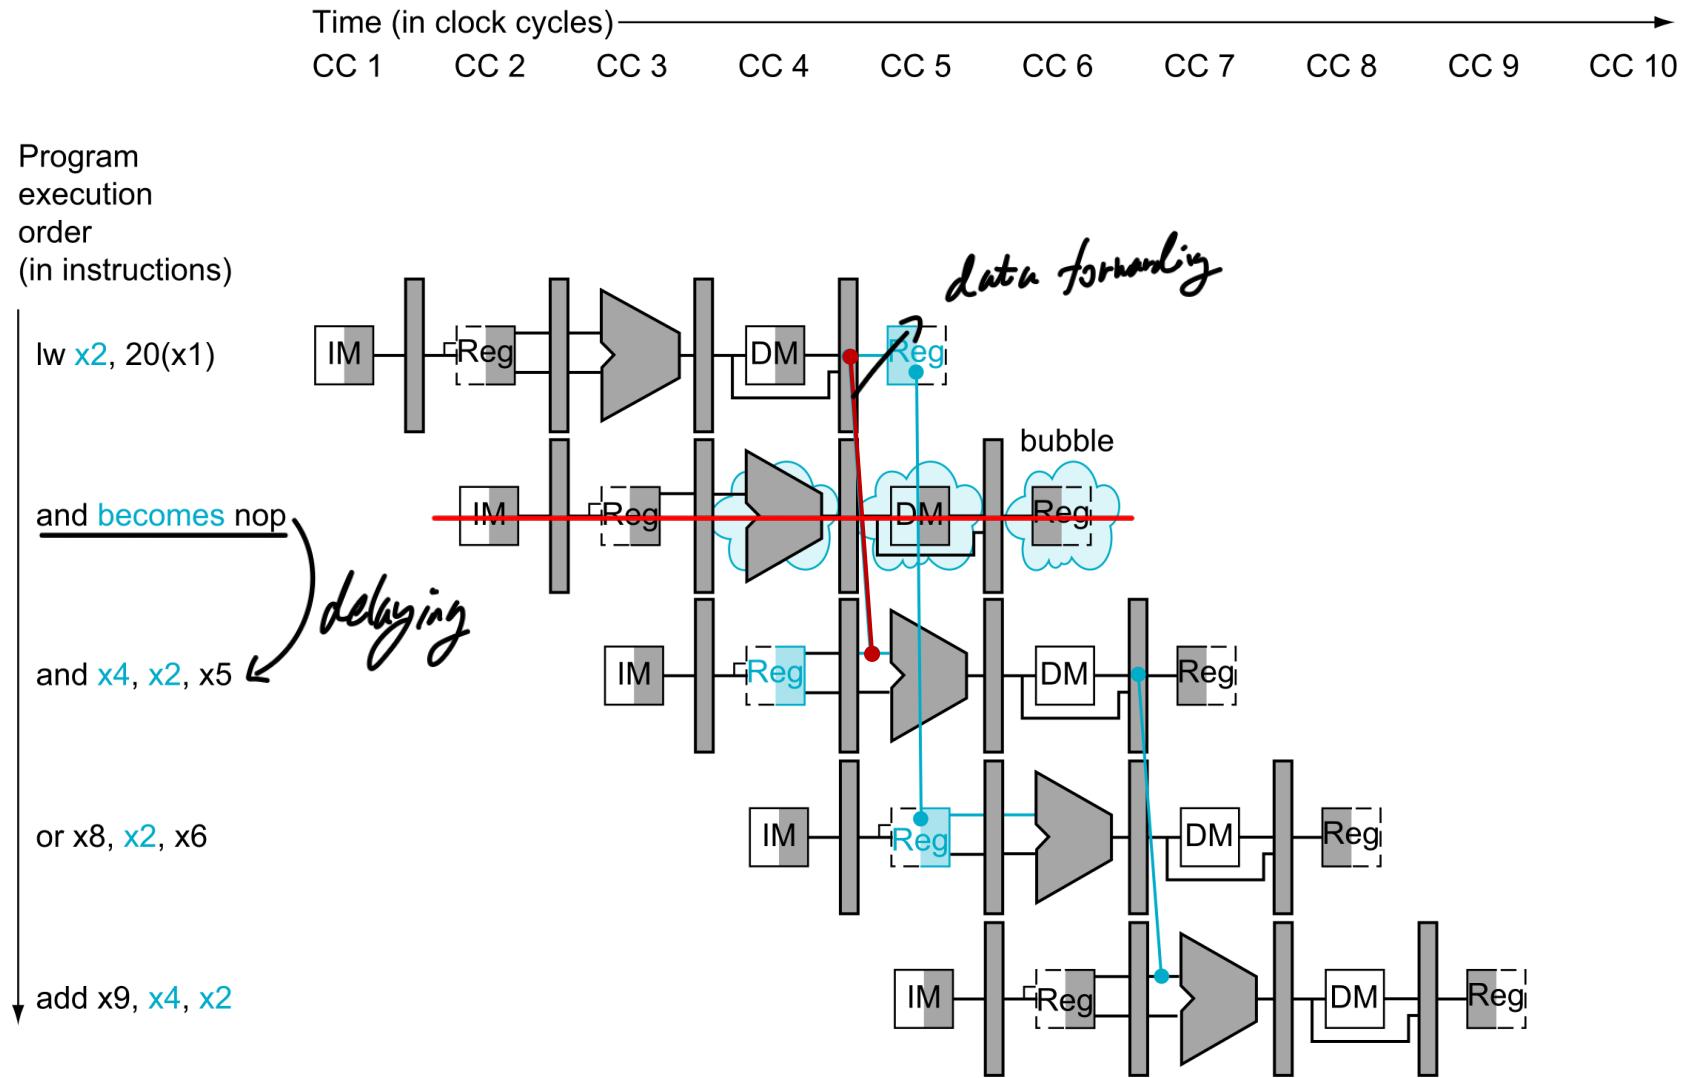
■ Detection

- Check for dependency when instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
 - ▶ IF/ID.RegisterRs1, IF/ID.RegisterRs2
- Load-use hazard when

Query type
ID/EX.MemRead and

$((\text{ID/EX.RegisterRd} = \text{IF/ID.RegisterRs1}) \text{ or } (\text{ID/EX.RegisterRd} = \text{IF/ID.RegisterRs2}))$ *Hazard condition*

Load-Use Hazard Resolution



How to Stall the Pipeline

■ Force control values in ID/EX register to 0

- Inserts a nop (no operation) into the ID/EX register
 - ▶ By setting the control signals to 0
- The nop bubble will propagate through EX, MEM and WB like a regular instruction

setup register

■ Prevent update of PC and IF/ID register (*delay*)

- Current instruction is decoded again
- Following instruction is fetched again
- 1-cycle stall allows MEM to read data for 1d that can subsequently be forwarded to EX stage

Datapath with Hazard Detection

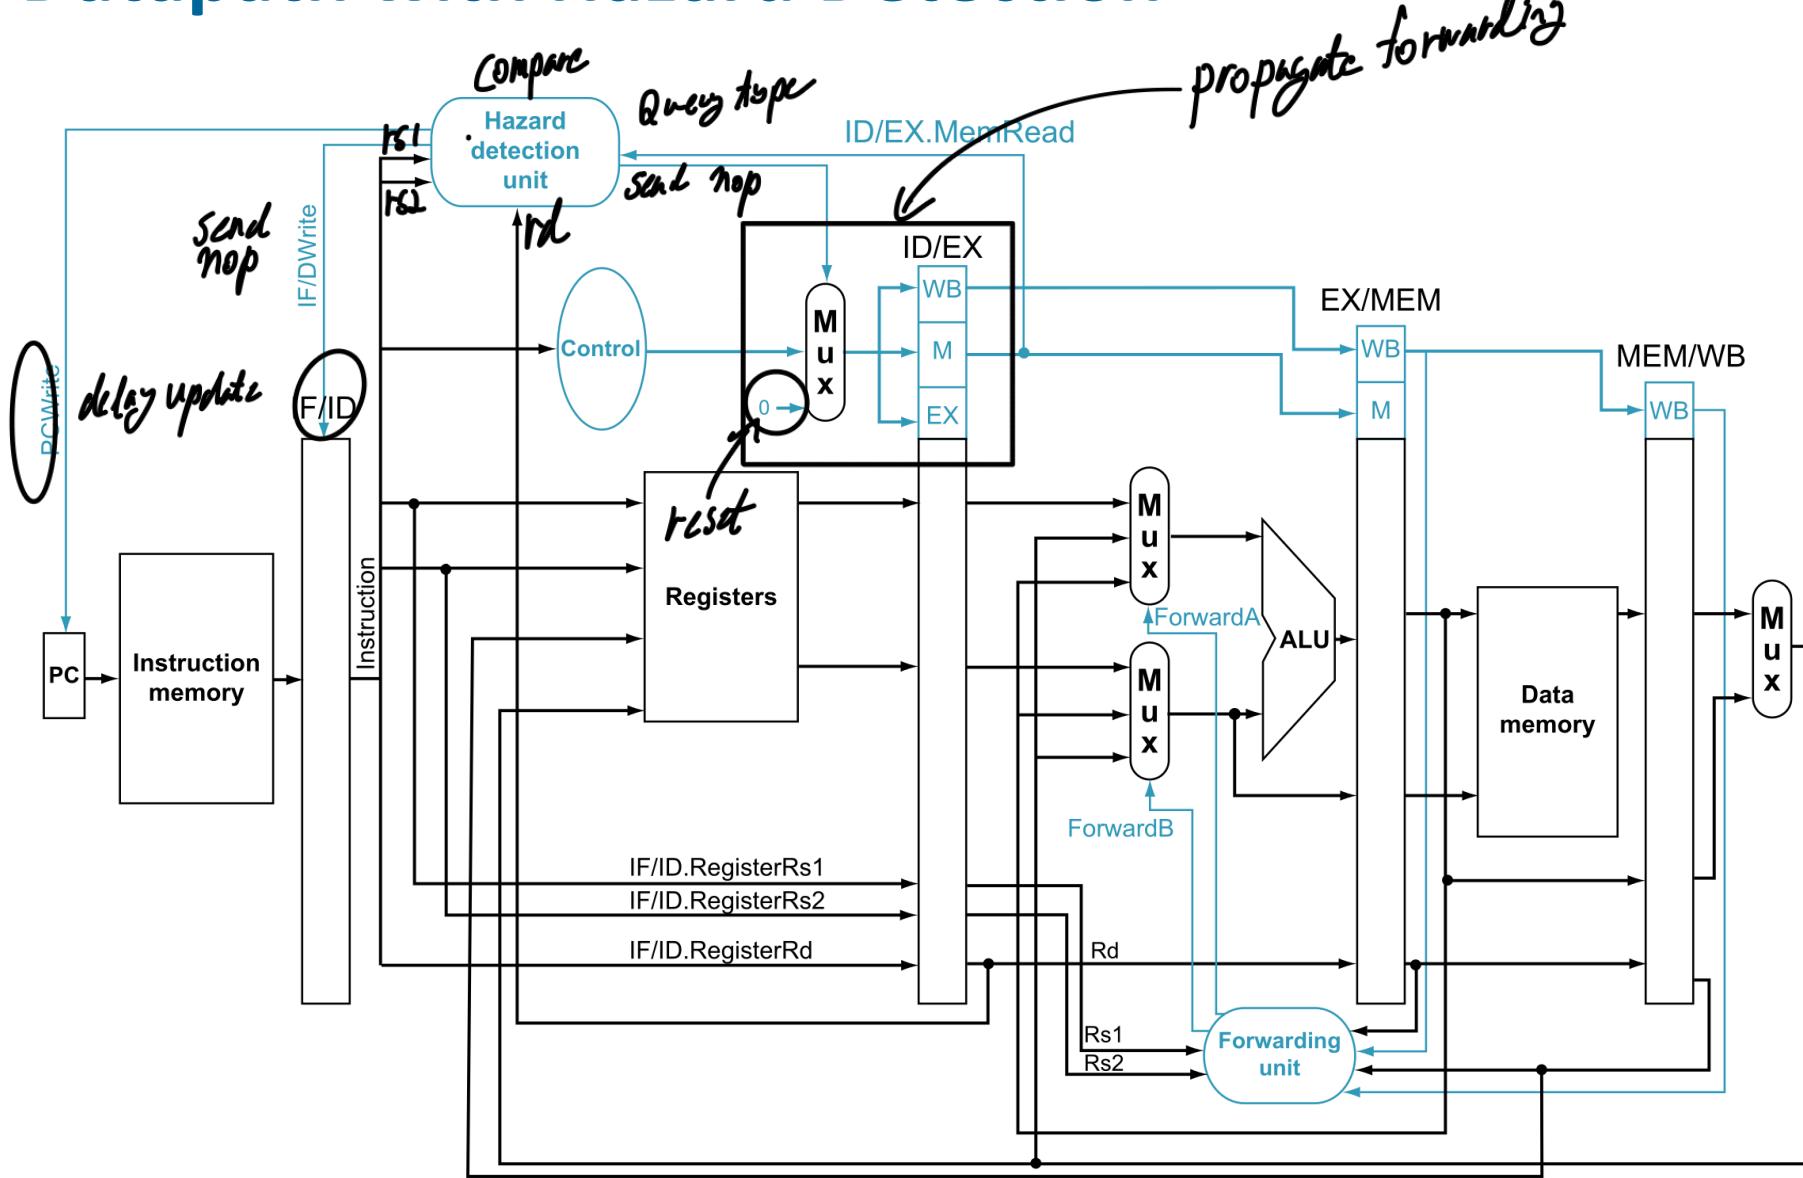


Illustration of Load Interlock Resolution

and x_4, x_2, x_5 1w $x_2, 40(x_1)$

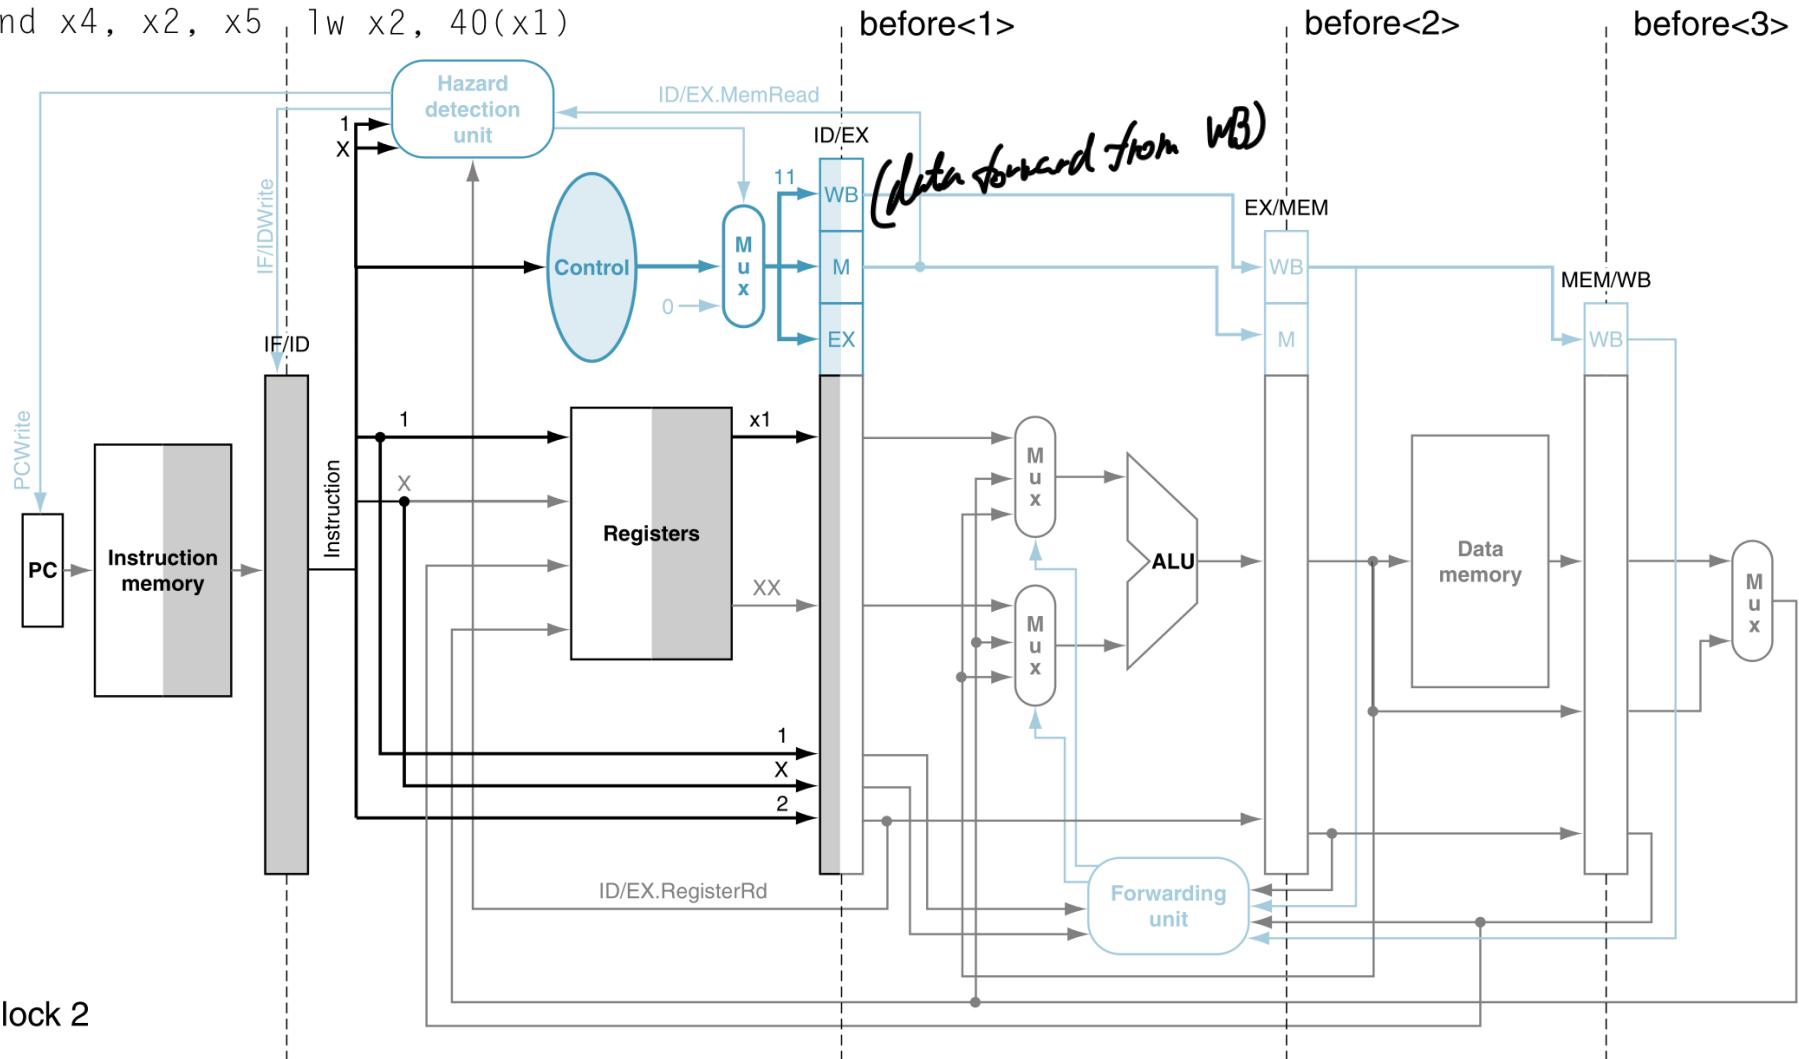


Illustration of Load Interlock Resolution

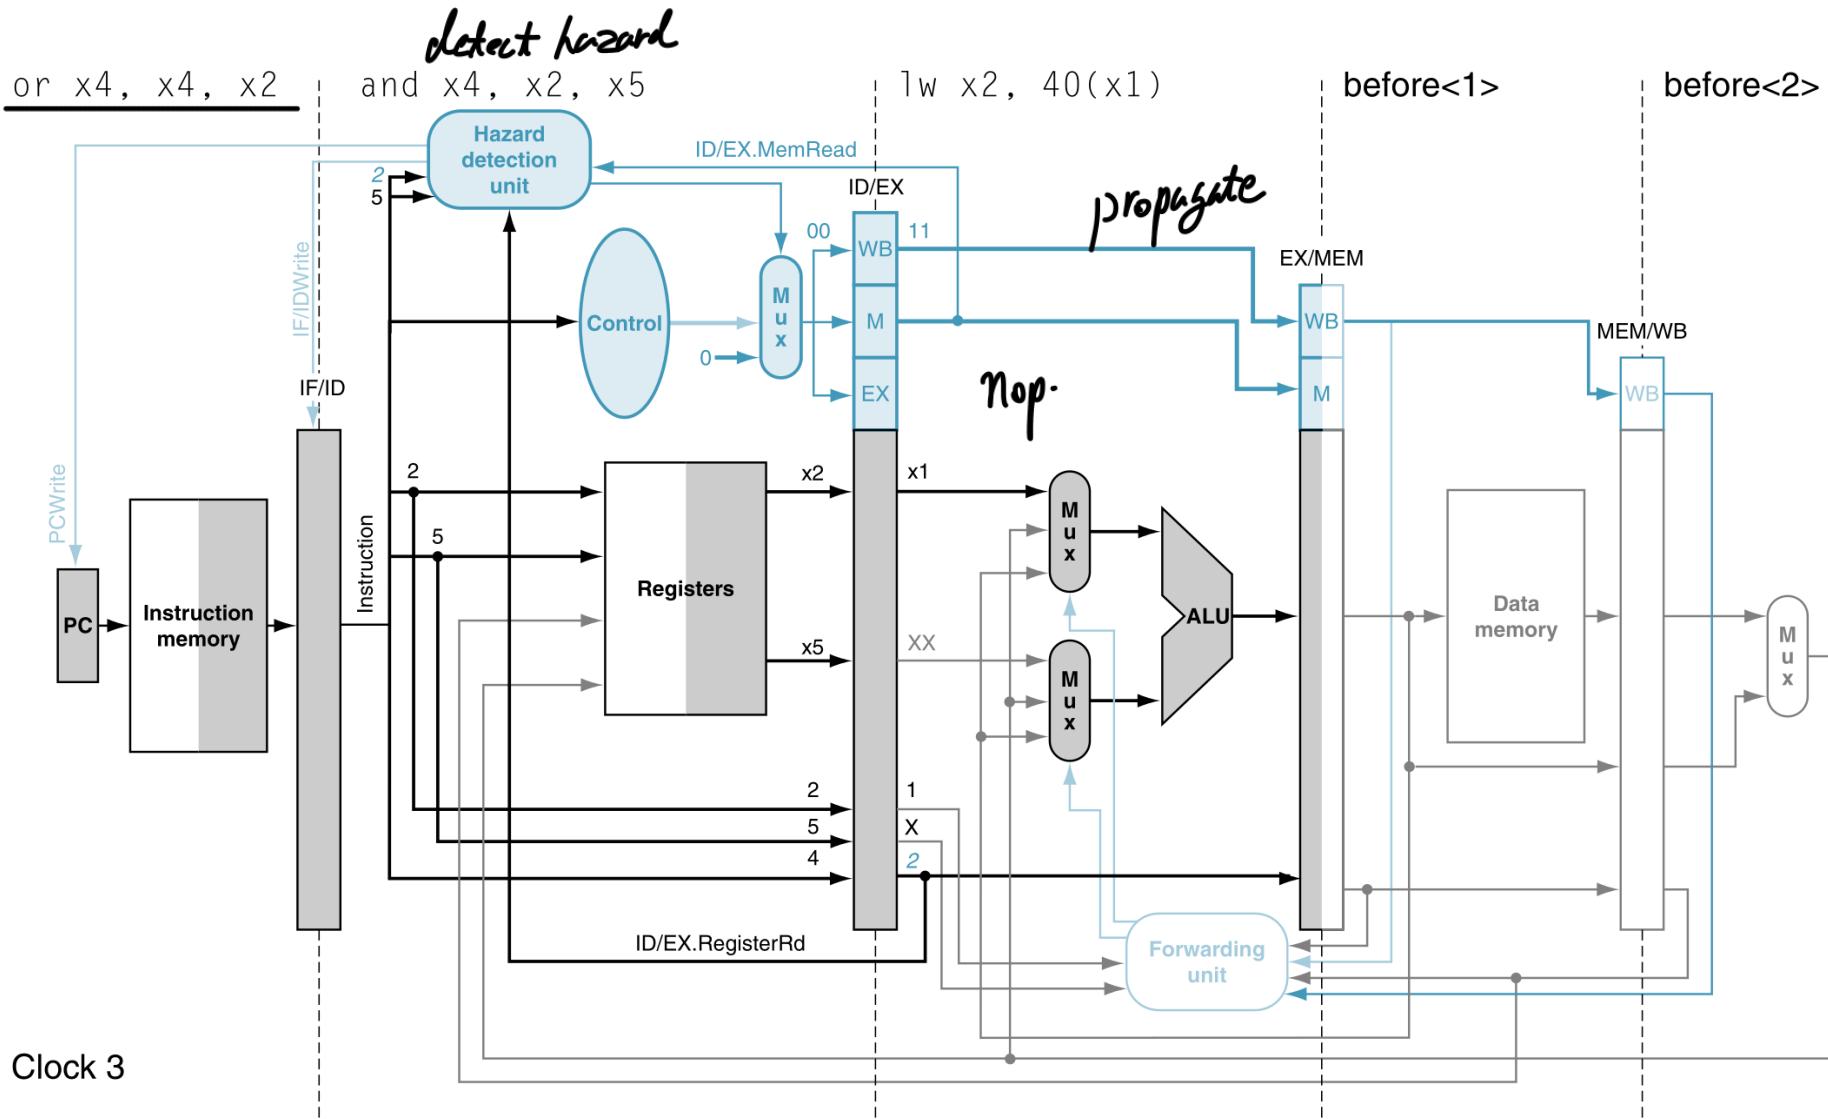


Illustration of Load Interlock Resolution

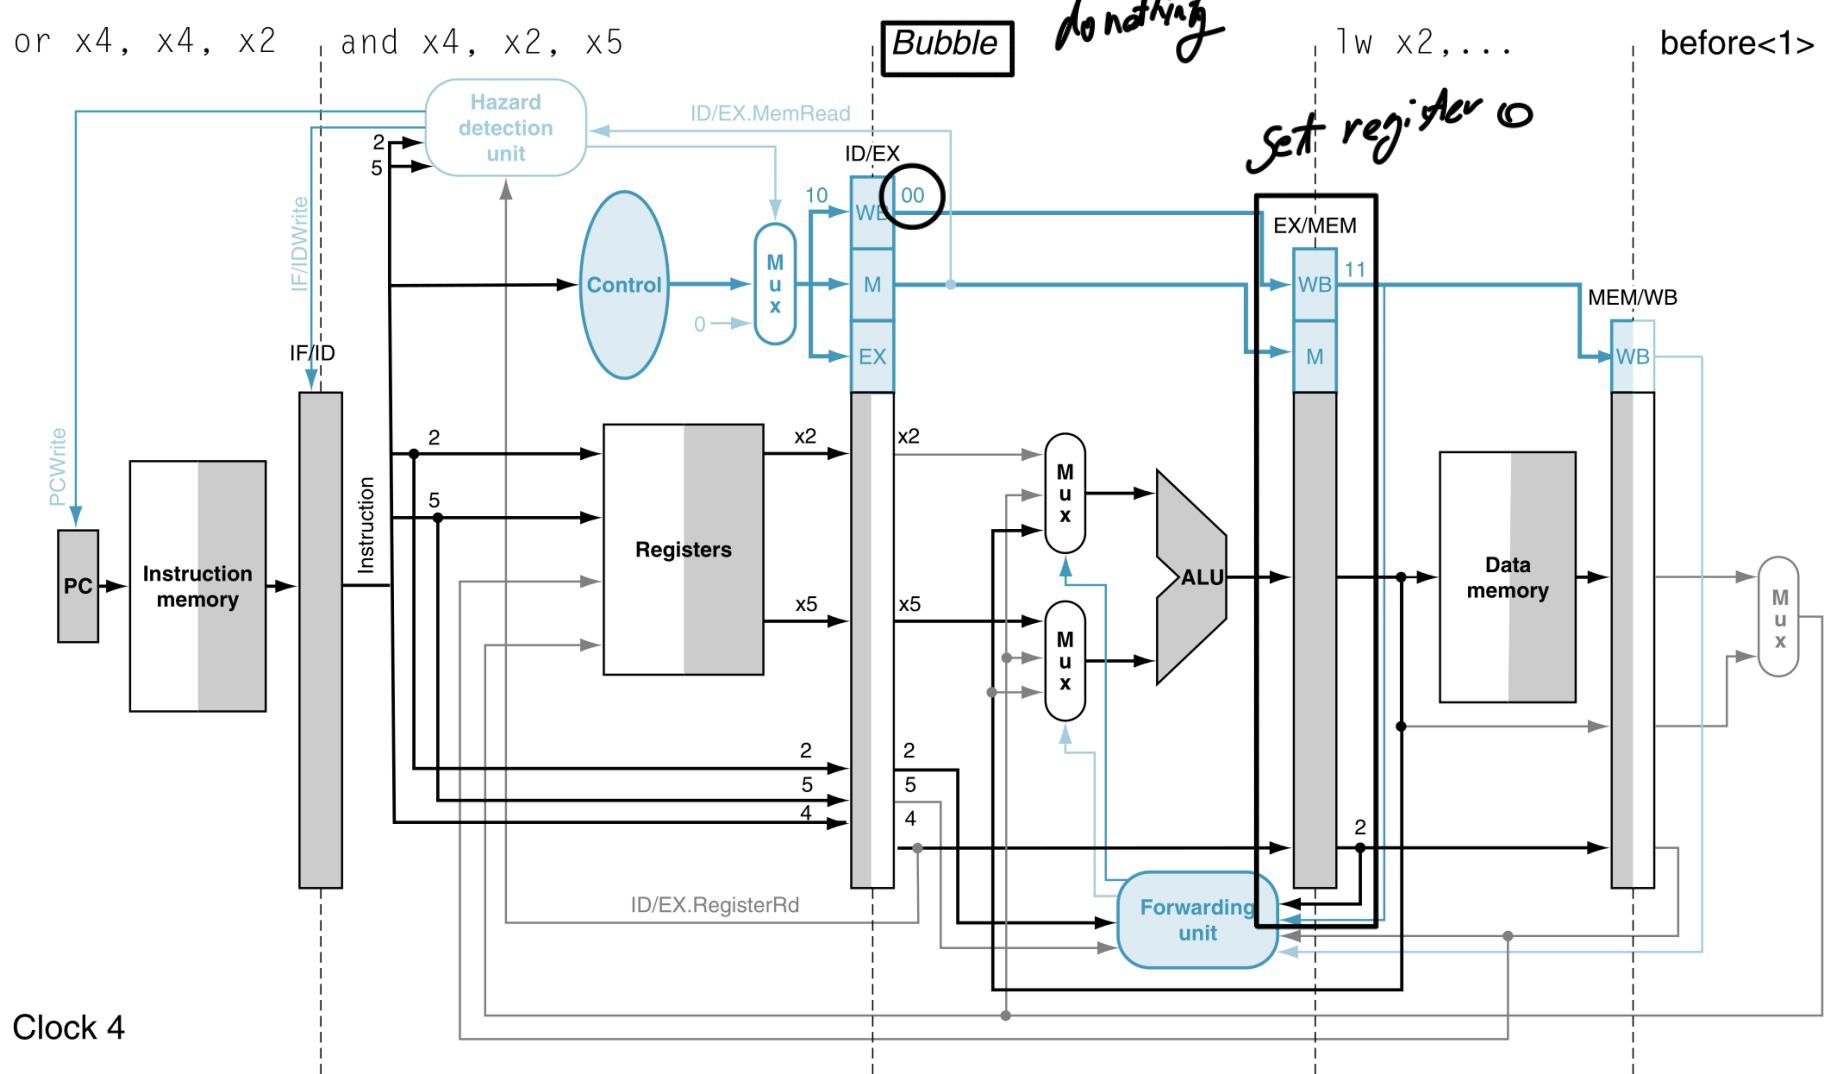


Illustration of Load Interlock Resolution

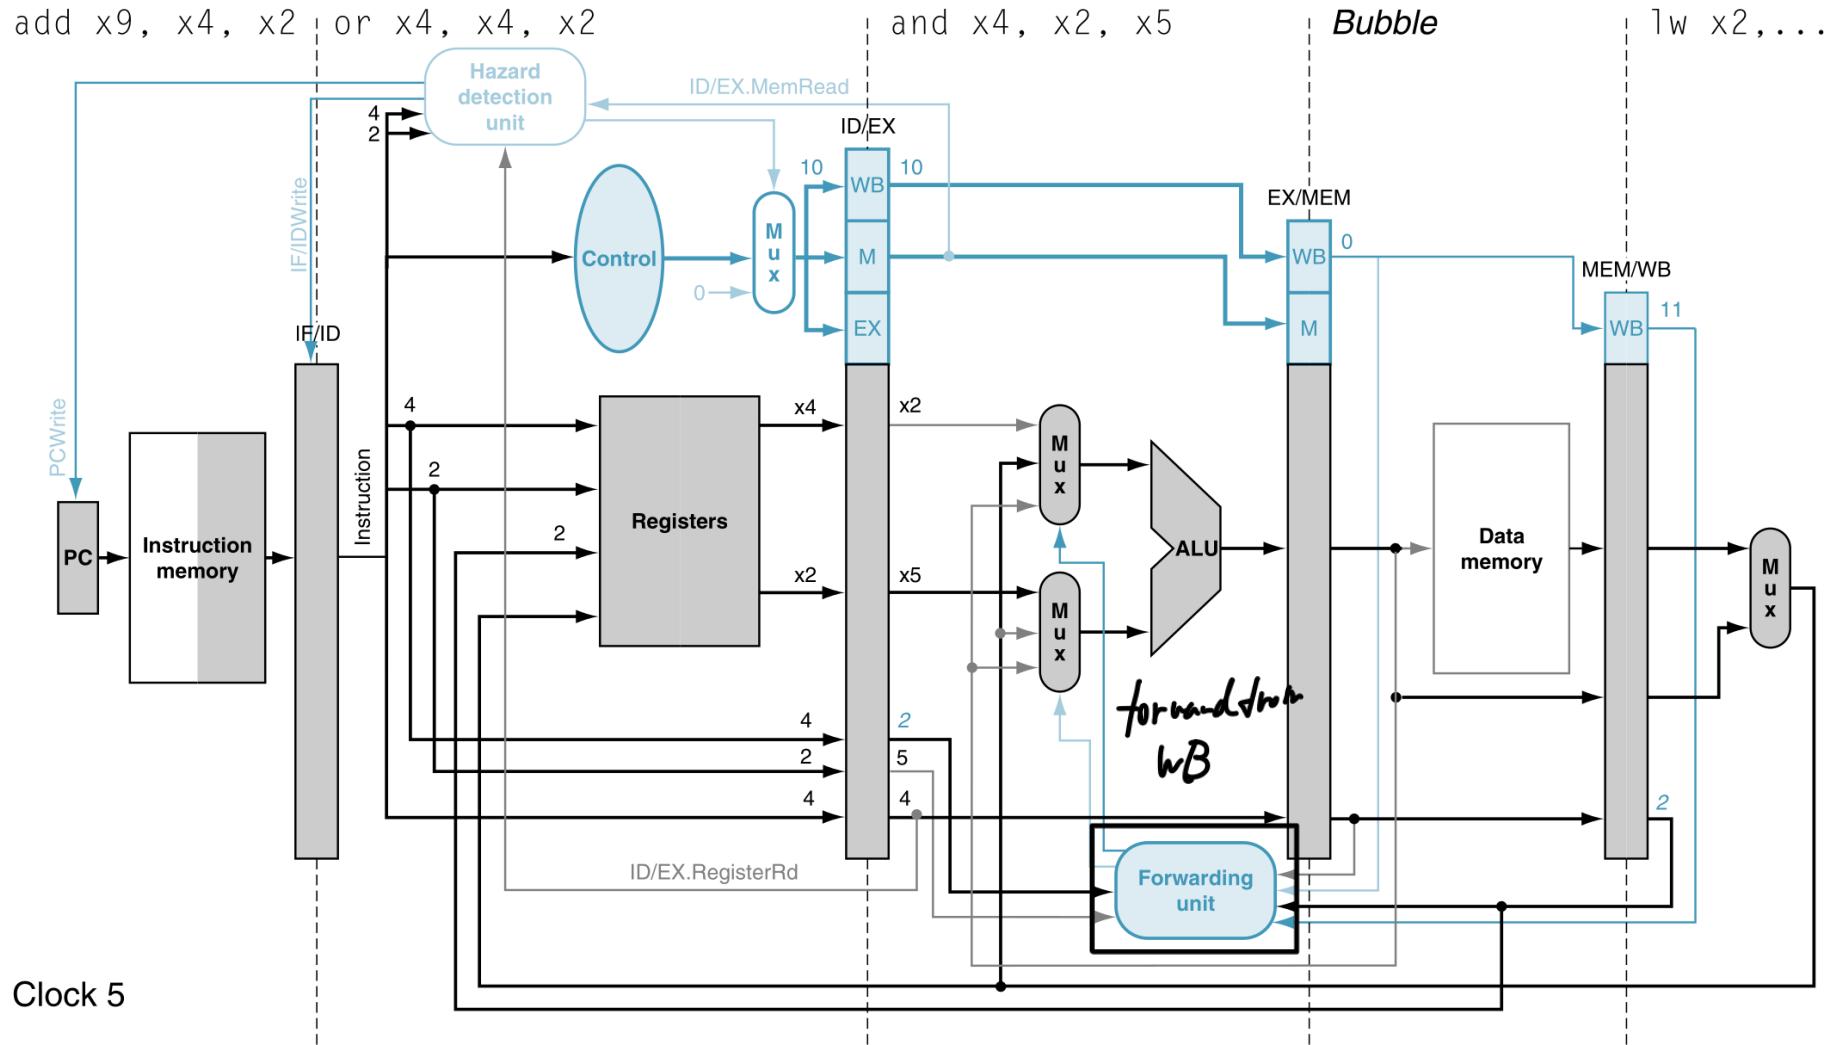
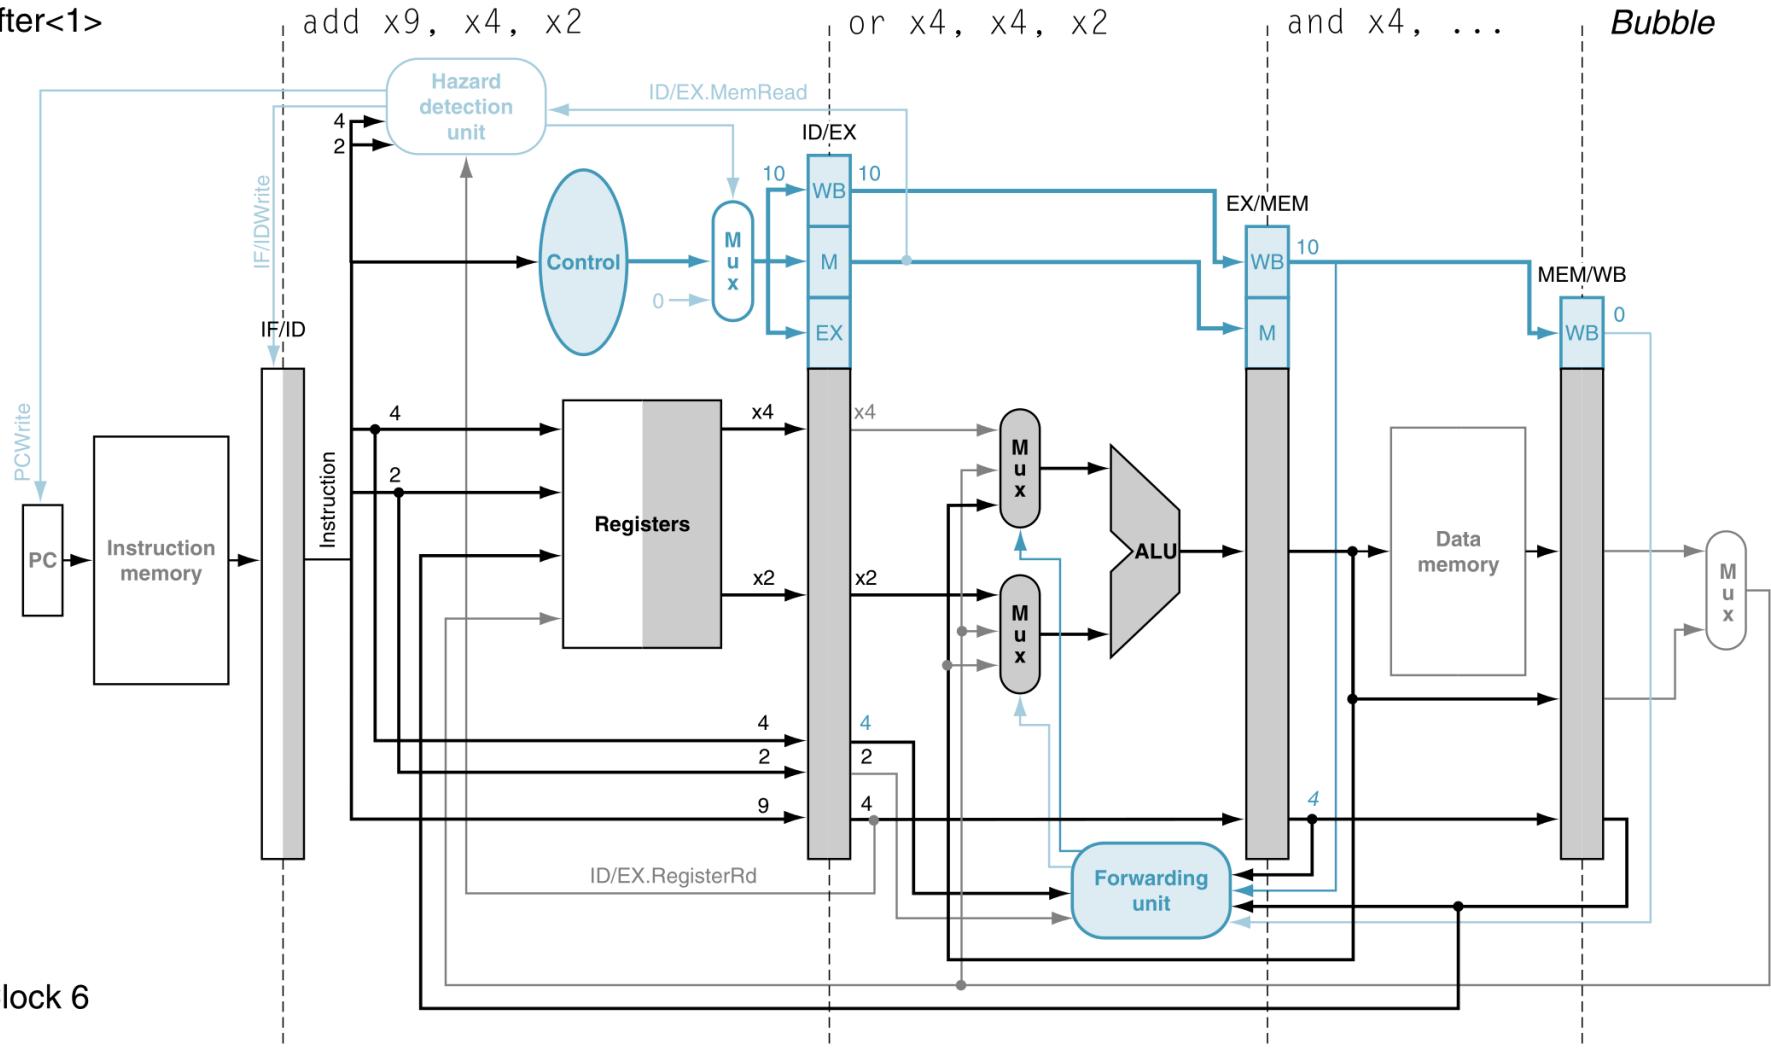


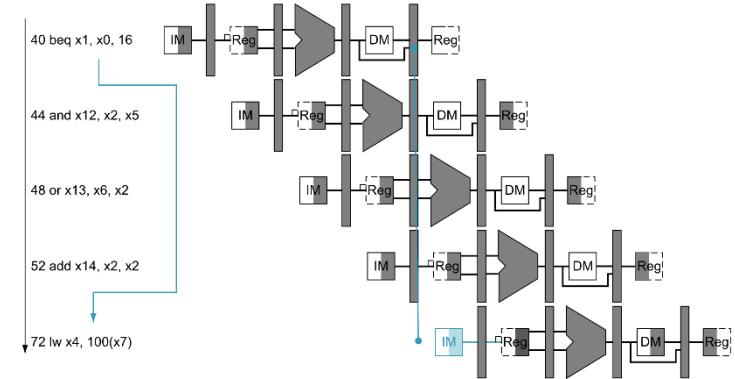
Illustration of Load Interlock Resolution

after<1>



Stalls and Performance

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure



Control Hazards

Recap: Control Hazards

Caused by PC-changing instructions

- conditional branches, function call/return
- may fetch wrong instructions

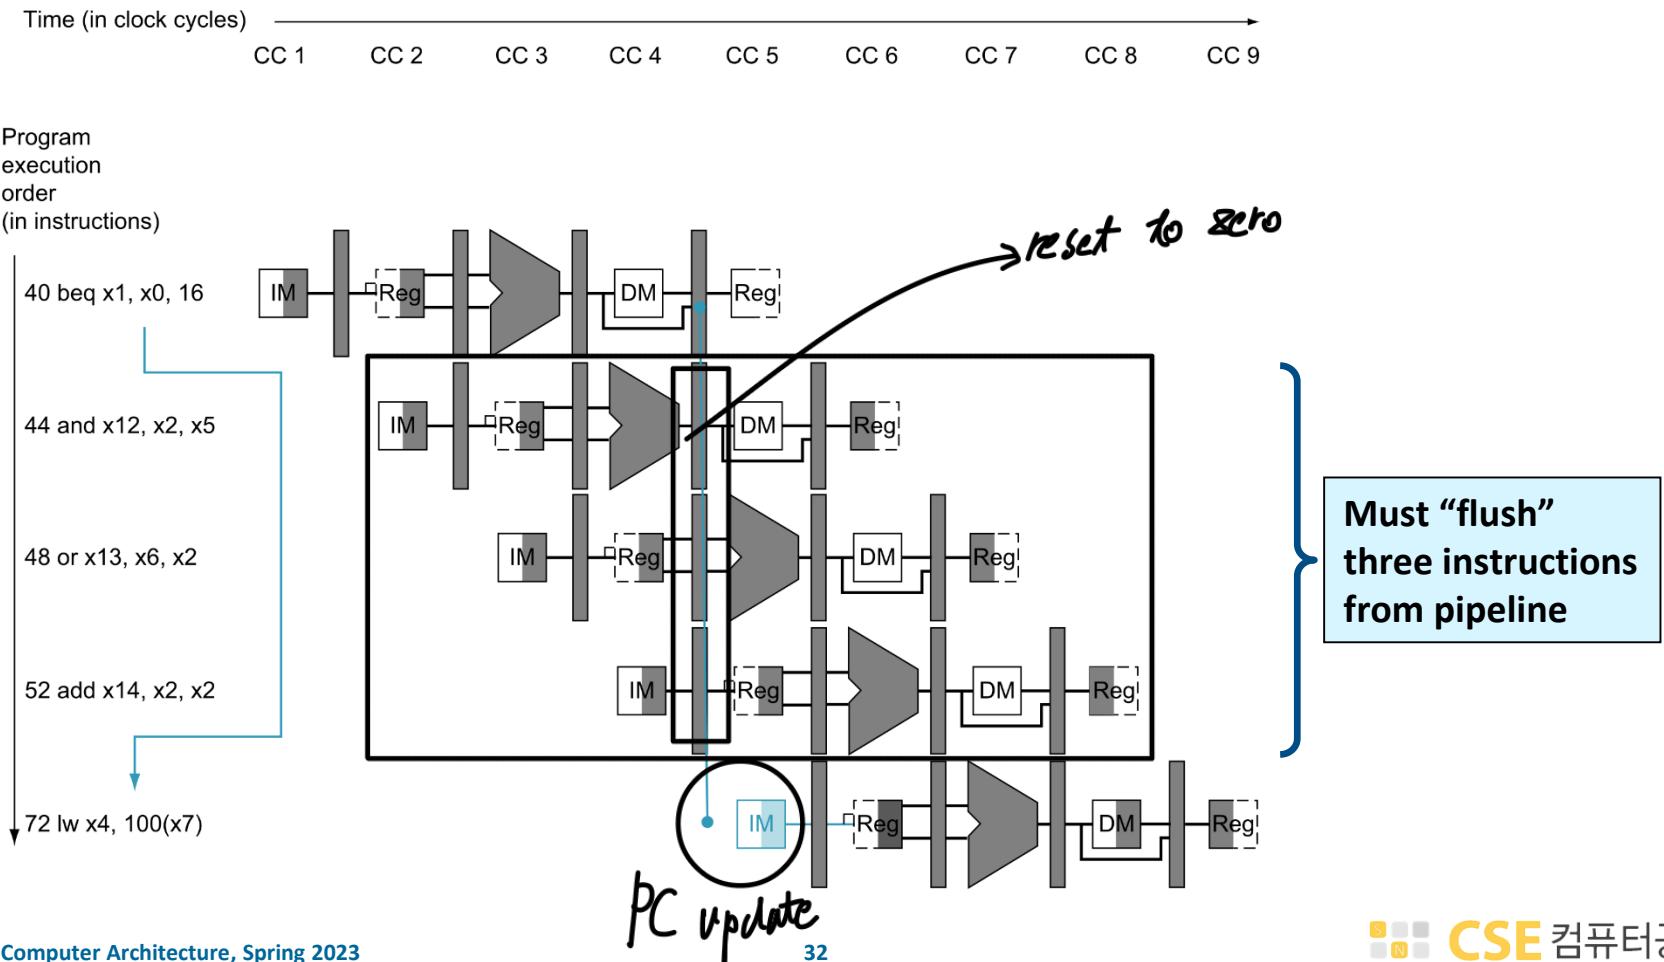
Branch Instruction	IF	ID	EX	MEM	WB			
Branch successor	— IF	stall	stall	—	IF	ID	EX	MEM WB
Branch successor + 1	Cancel process				IF	ID	EX	MEM WB
Branch successor + 2						IF	ID	EX MEM
Branch successor + 3						IF	ID	EX
Branch successor + 4						IF	ID	
Branch successor + 5						IF		

run after PC updated

Cancel process

Control Hazards

- Branch outcome determined in MEM
- Always assume branch not taken *static prediction*



Reducing the Branch Delay

■ Move hardware to determine outcome to ID stage

- Easy

$EX \rightarrow ID$

- ▶ Target address adder
- ▶ Register comparator

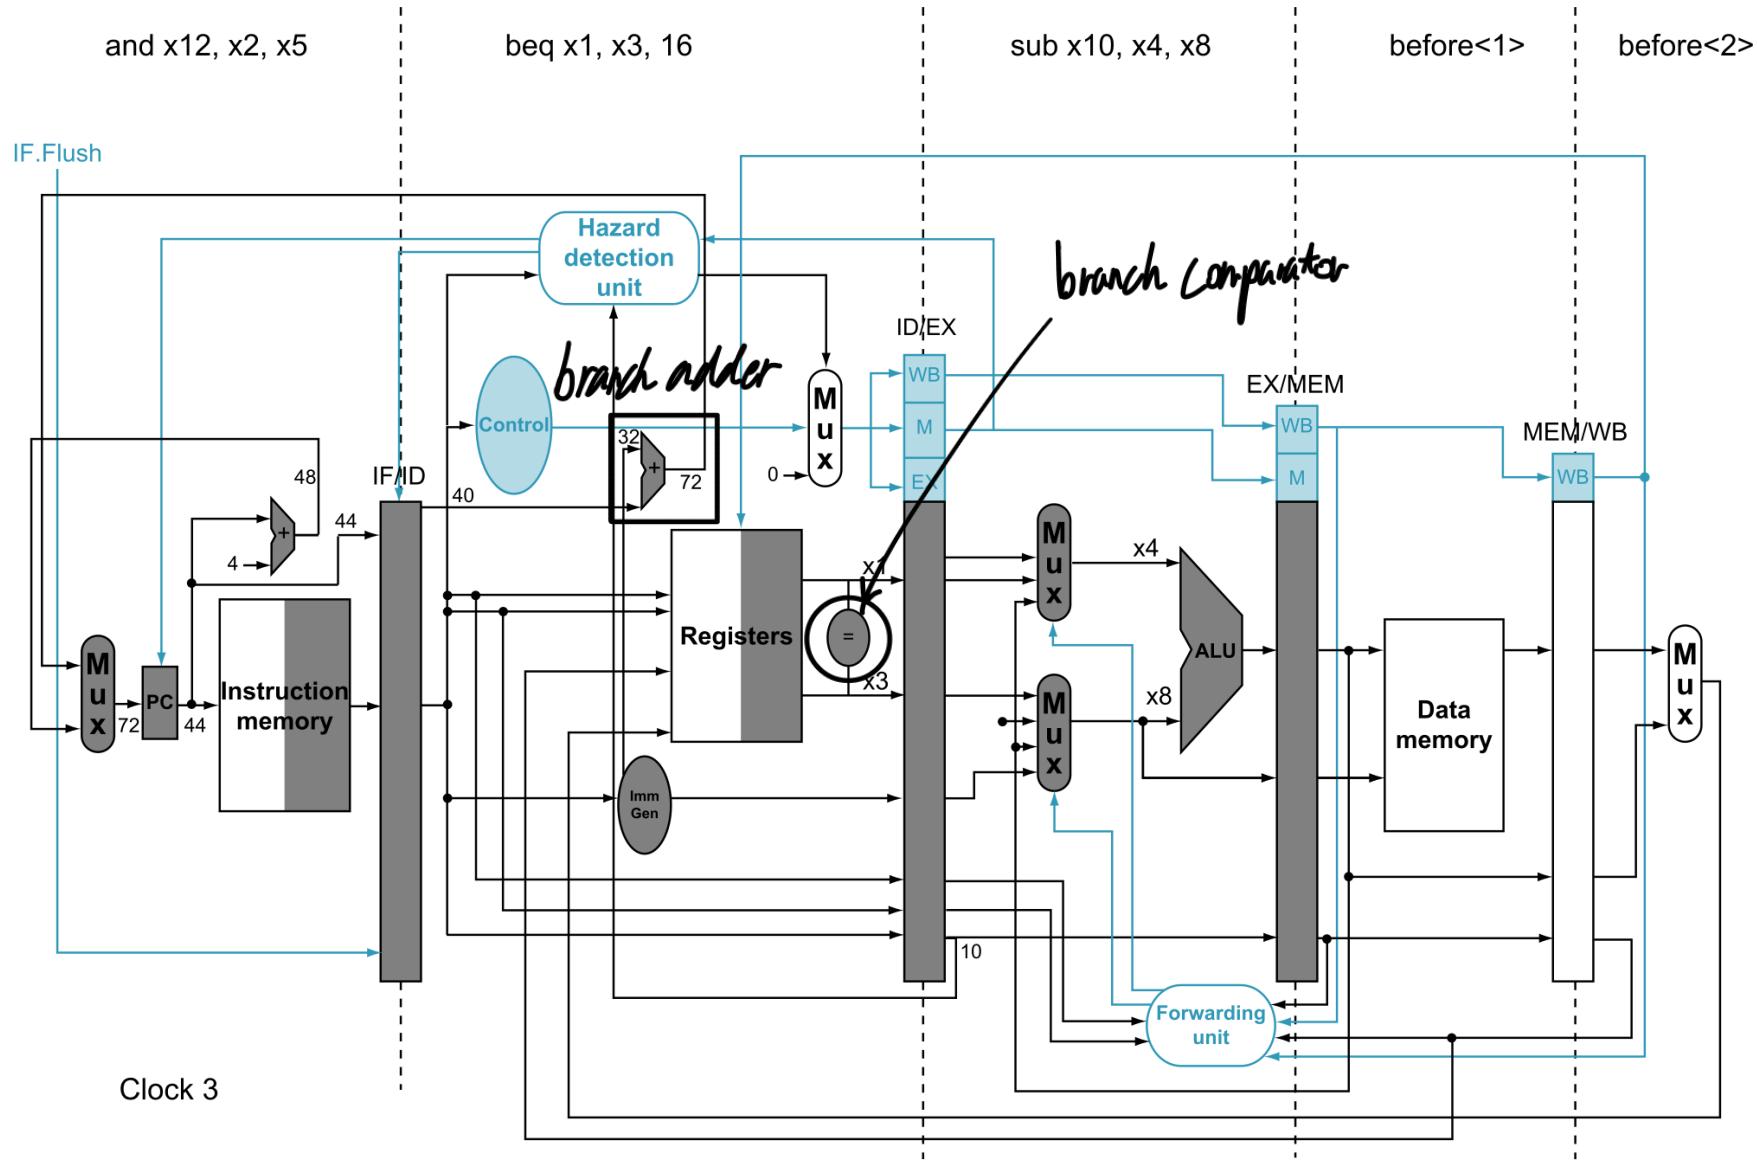
- Difficult

- ▶ Branch decision: comparing two registers in ID stage requires additional forwarding and hazard detection logic

■ Example: branch taken

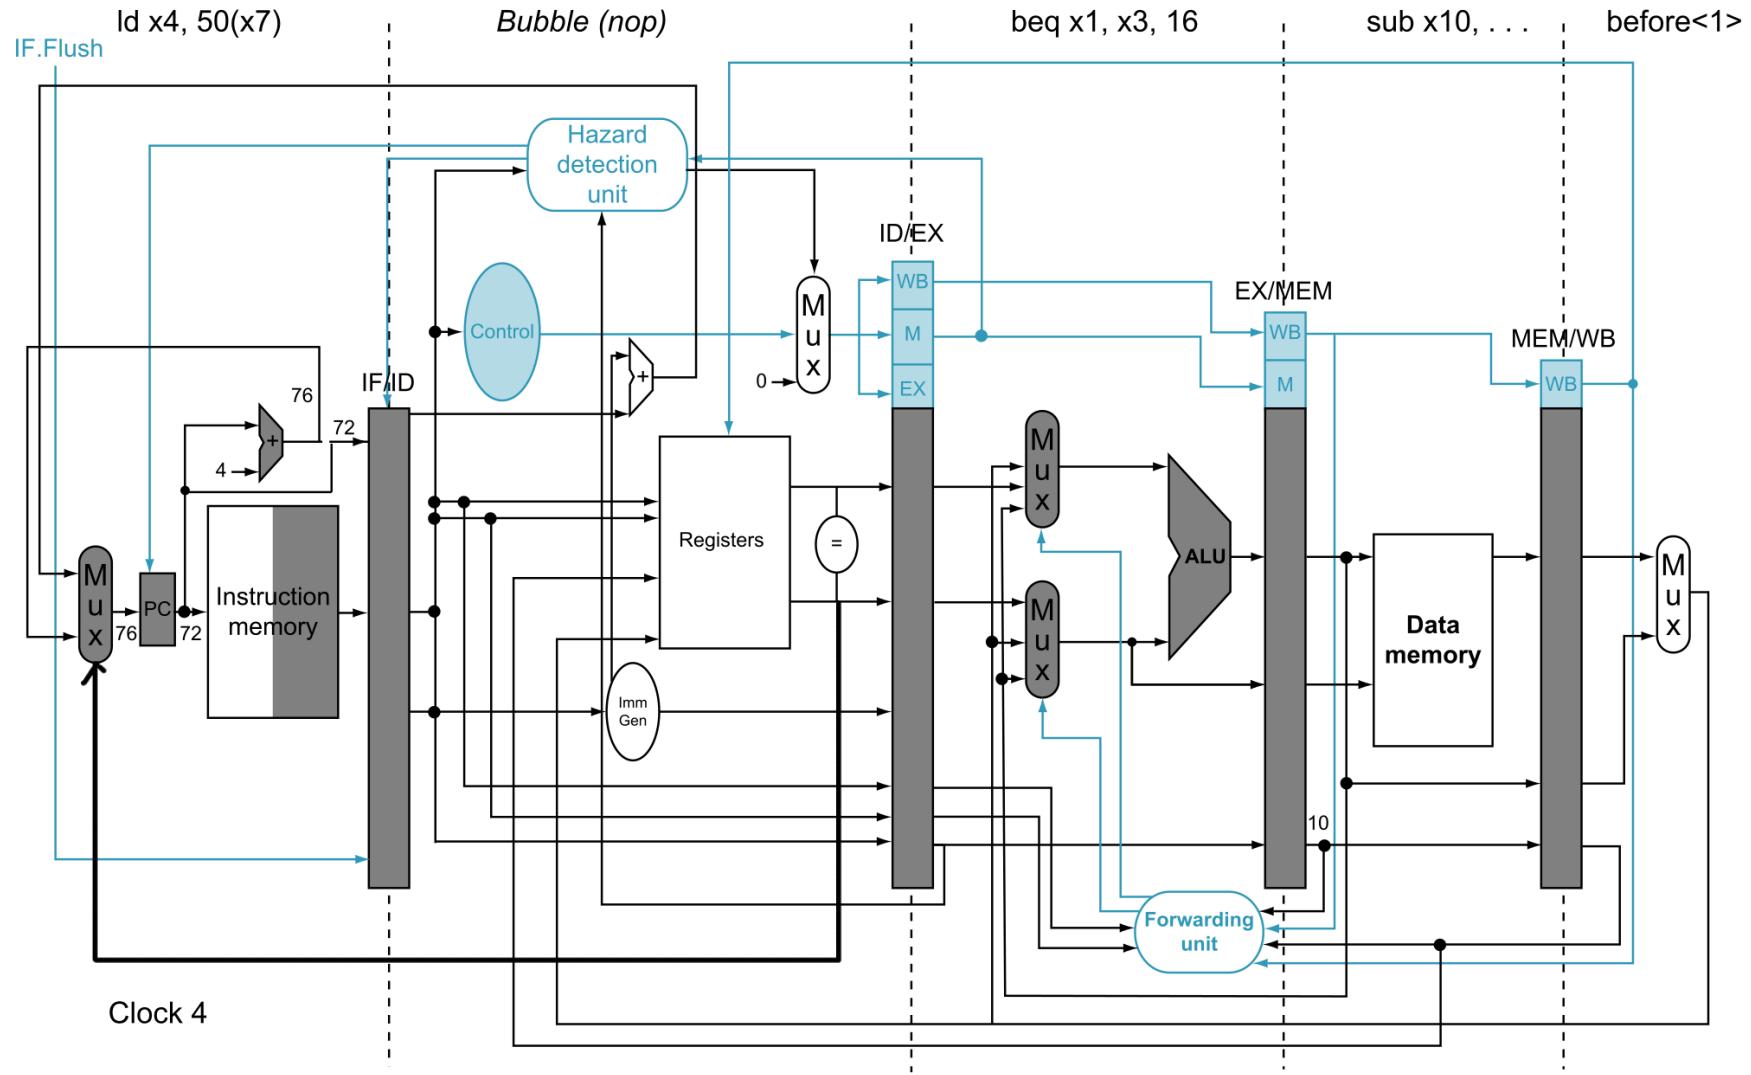
```
36:    sub     x10, x4, x8
40:    beq     x1, x3, 16      // PC-relative branch to 40+16*2 = 72
44:    and     x12, x2, x5
48:    or      x13, x2, x6
52:    add     x14, x4, x2
56:    sub     x15, x6, x7
...
72:    ld      x4, 50(x7)
```

Example: Branch Taken



Example: Branch Taken

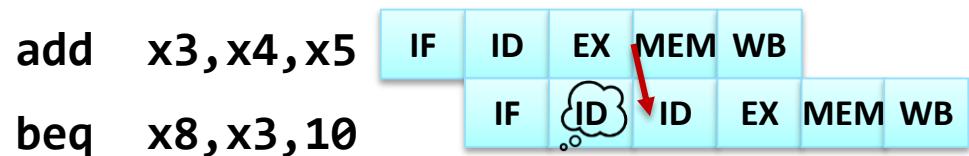
before, and



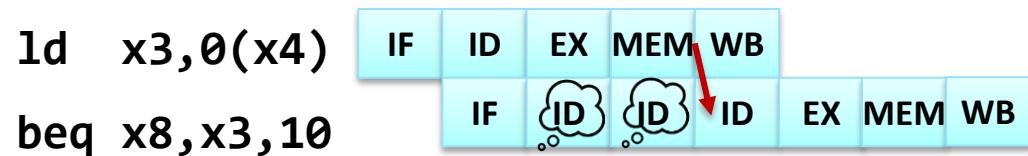
Cost of Moving Branch Test to ID

⇒ Must forward to 'ID' phase

- Register operands may require forwarding
 - New forwarding logic from EX/MEM or MEM/WB pipeline registers to ID needed
- Stalls due to data hazard *(if data from forward, can't determine)
⇒ need stall*
 - 1-cycle stall if the preceding instruction is an ALU instruction



- 2-cycle stall if the preceding instruction is a load instruction



Dynamic Branch Prediction

Static Branch Prediction

- always taken 50%
- always not taken

- In deeper and superscalar pipelines branch penalty is more significant

+d : Based on previous decision

- Use dynamic prediction

- Branch prediction buffer (or branch history table) : history based predict (Super expensive)
- Indexed by recent branch instruction addresses (lower part of address)
- Stores outcome (taken / not taken)

- To execute a branch

- Check table, expect the same outcome
- Start fetching from fall-through or target
- If wrong, flush pipeline and flip prediction

0x104

0x108

0x112

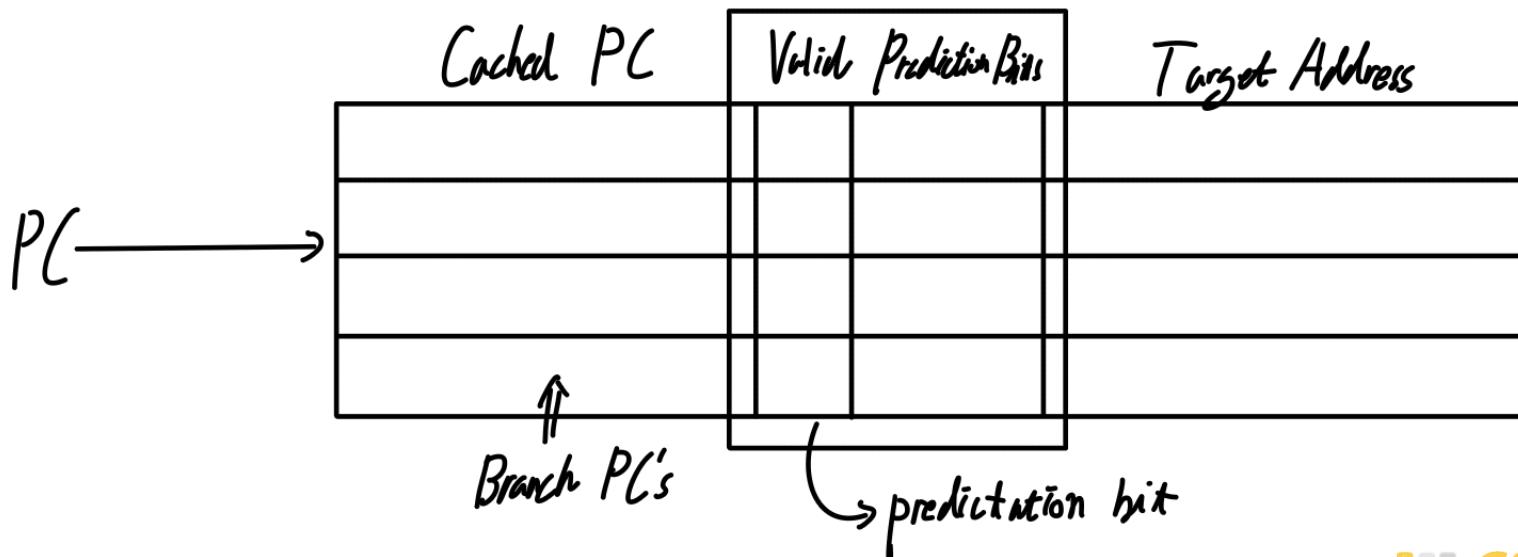
0x116

0x120

↓ 4 (\Rightarrow 0 or 1, 2 bits)

Calculating Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer (BTB)
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately



1-Bit Predictor Example

BTB

Code

0x14380: L1: add x3,x3,x5
0x14384: beq x8,x3,L1
...
...

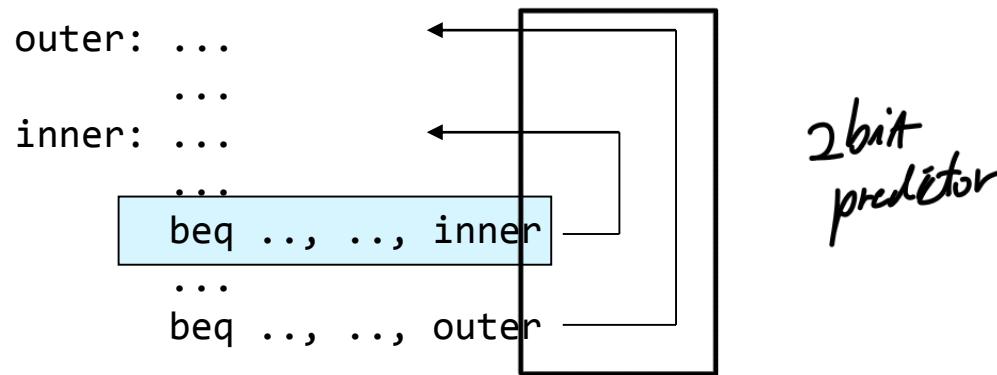
Address tag	Valid	Branch Target Buffer
...0	0	
...4	0	
...8	0	
...c	0	

Execution

CC	Address	Instruction	x8	x3	Adr tag	Valid	BTB
10	14384	beq x8,x3,L1	10	7	4	0	first meet
11	14380	add x3,x3,x5	8		4	1	14380
12	14384	beq x8,x3,L1	10	8	4	1	14380
13	14380	add x3,x3,x5	9				
14	14384	beq x8,x3,L1	10	9	4	1	14380
15	14380	add x3,x3,x5	10				
16	14384	beq x8,x3,L1	10	10	4	1	14380
17	14388	...				0	invalid now!

1-Bit Predictor: Shortcoming

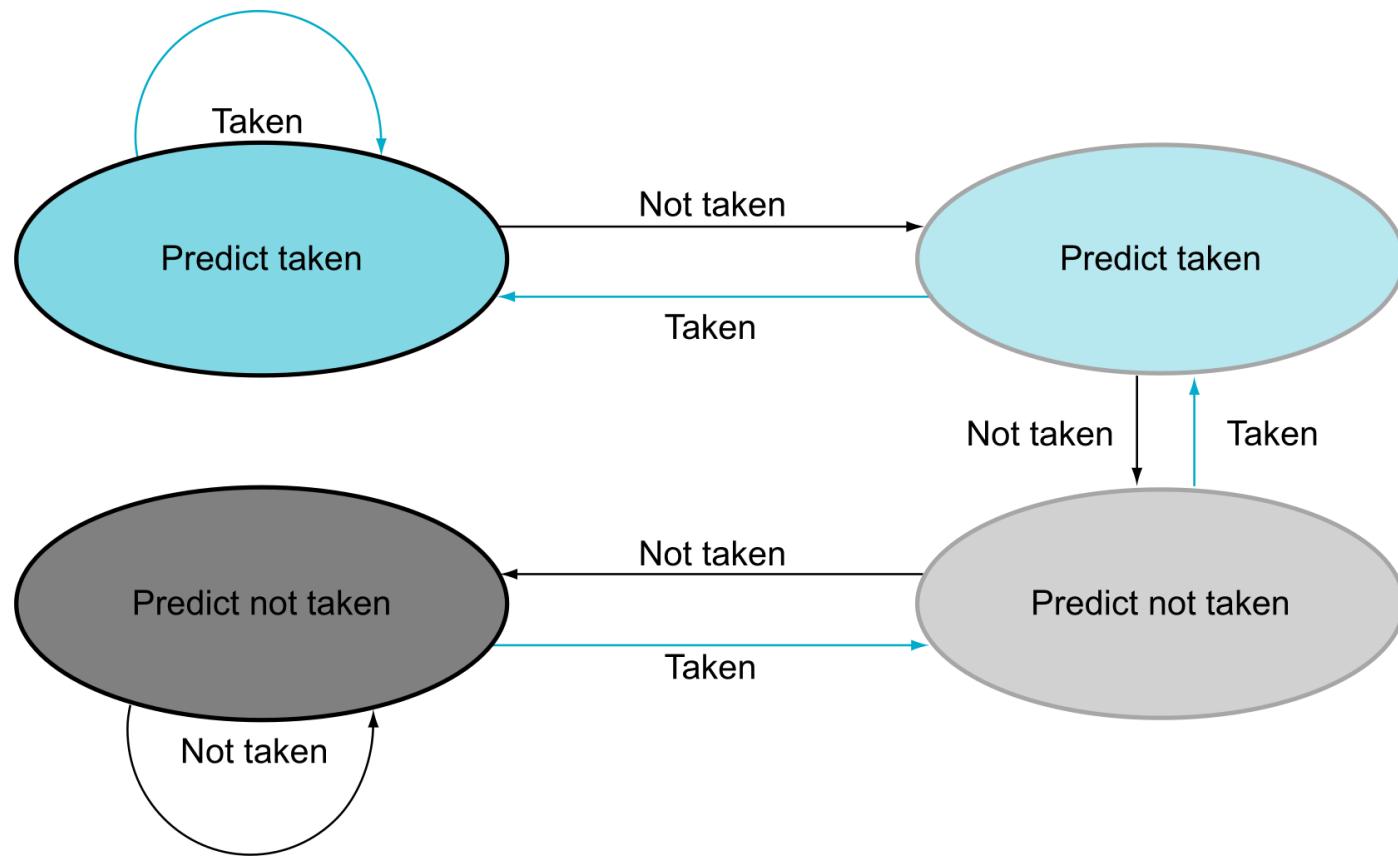
- Inner loop branches always mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

- Only change prediction after two successive mispredictions



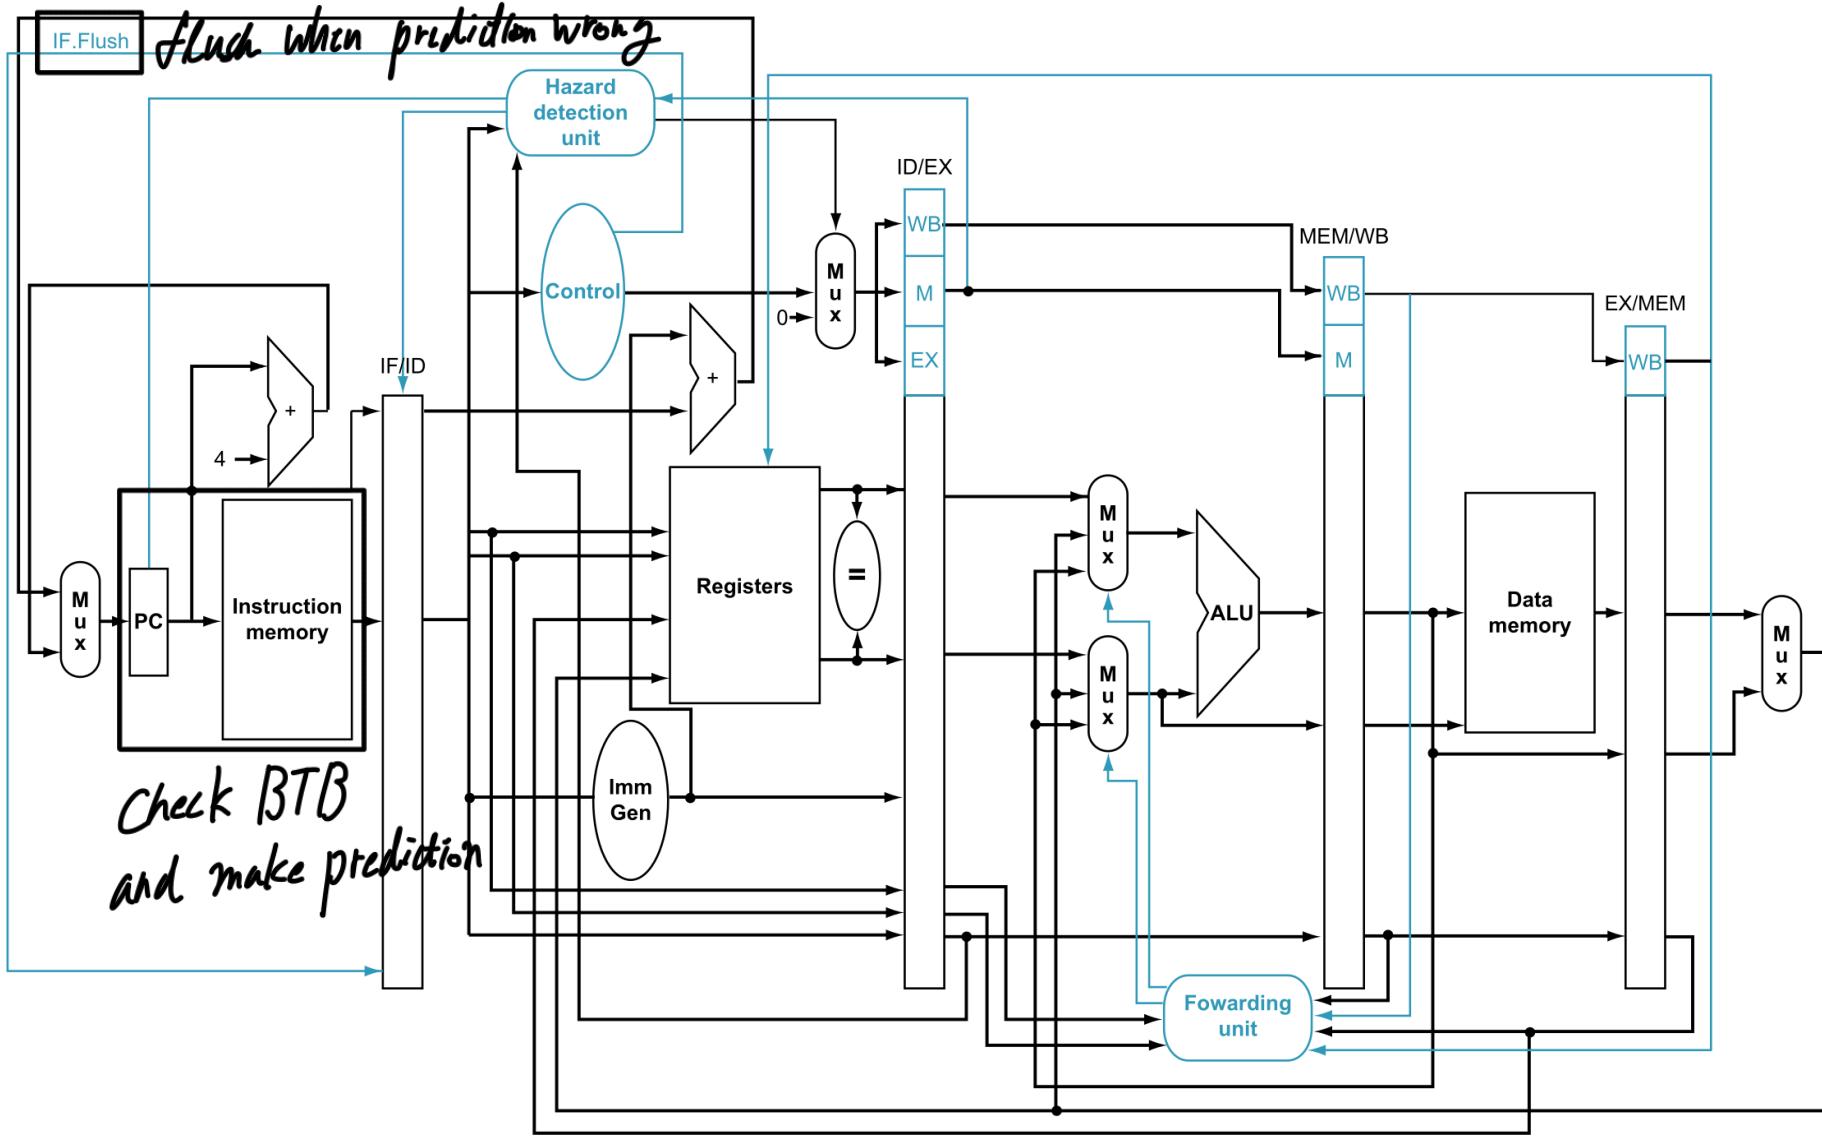
2-Bit Predictor Example

11: Predict taken
 10: Predict taken
 01: Predict not taken
 00: Predict not taken

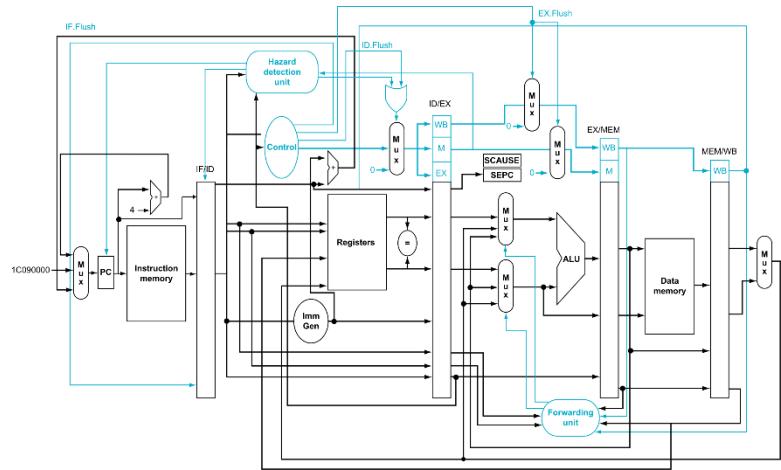
Execution

CC	Address	Instruction	x8	x3	Adr tag	Pred	BTB
inner Loop1	10	14384 beq x8,x3,L1	10	7	4	00	
	11	14380 add x3,x3,x5		8	4	01	taken +1
	12	14384 beq x8,x3,L1	10	8	4	01	recording
	13	14380 add x3,x3,x5		9	4	10	taken +1
	14	14384 beq x8,x3,L1	10	9	4	10	14380
	15	14380 add x3,x3,x5		10	4	11	taken +1 14380
	16	14384 beq x8,x3,L1	10	10	4	11	14380
	17	14388 ...			4	10	not taken -1 14380
...							
inner Loop2	20	14384 beq x8,x3,L1	10	7	4	10	14380
	21	14380 add x3,x3,x5		8	4	11	14380
	22	14384 beq x8,x3,L1	10	8	4	11	14380
	23	14380 add x3,x3,x5		9	4	11	14380
	24	14384 beq x8,x3,L1	10	9	4	11	14380
	25	14380 add x3,x3,x5		10	4	11	14380
	26	14384 beq x8,x3,L1	10	10	4	11	14380
	27	14388 ...			4	10	14380

Datapath and Control



Exceptions



Accident type	Source	RISC-V
System reset	External	Exception
I/O device request	External	Interrupt
OS system call	Internal	Exception
Undefined Operation	Internal	Exception
Hardware Error	In/External	Exception / Interrupt

Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
 - Different ISAs use the terms differently
 - Exceptions
 - RISC-V: any unexpected change in control flow (either internal or external)
 - Internal exceptions arise within the CPU (e.g., undefined opcode, syscall, ...)
 - Interrupt
 - RISC-V: event from an external I/O controller
 - e.g., hard disks, network adapters, keyboard, ...
 - Dealing with them without sacrificing performance is hard
 - But also not necessary if exceptions do not happen frequently
 - Make the common case fast
- (Non-error phase)*
- divide by zero*
- 정상적인 계산이나 예외 처리*
- handle in OS*

Handling Exceptions in RISC-V

- Save PC of offending (or interrupted) instruction *(Hardware error)*
 - Supervisor Exception Program Counter (SEPC)
Save address for decision or return
- Save indication of the problem
 - Supervisor Exception Cause Register (SCAUSE); *reason*
 - 64-bits, but most bits unused
 - Exception code field: 2 for undefined opcode, 12 for hardware malfunction, ...
detect reason of error by looking for SEPC + state register
- Jump to handler
 - Assume at 0x0000 0000 1C09 0000 (*+ error's handler address*)

An Alternate Mechanism

- Vectored interrupts
 - Handler address determined by the cause
- Exception vector address to be added to a vector table base register: *(base address + Vectorize address)*
 - Undefined opcode: 00 0100 0000
 - Hardware malfunction: 01 1000 0000
- Instructions either
 - Deal with the interrupt, or
 - Jump to real handler

Interrupt Type	Handler Address
	.
	.
	.
	.

Handler Actions

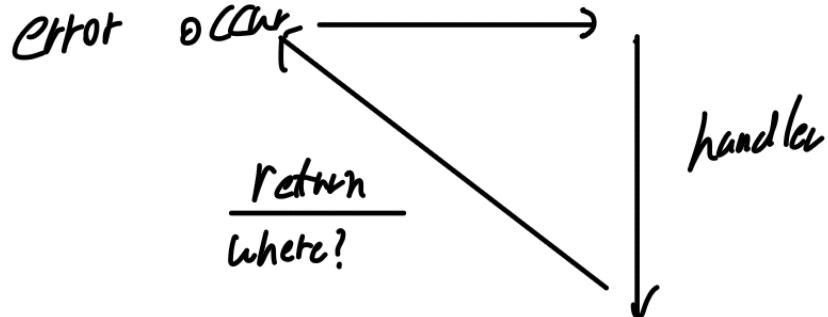
from inflation and compare it with SCAUSE

- Read cause, and transfer to relevant handler

- Determine action required

- If restartable
 - Take corrective action
 - Use SEPC to return to program

- Otherwise
 - Terminate program
 - Report error using SEPC, SCAUSE, ...



① after 'sys call': next operation

② after 'load': same address

③ general: stop program

Exception Properties

- Restartable exceptions
 - Pipeline can flush the instruction
 - Handler executes, then returns to the instruction:
Refetched and executed from scratch
- PC saved in SEPC register
 - Identifies causing instruction

Exceptions in a Pipeline

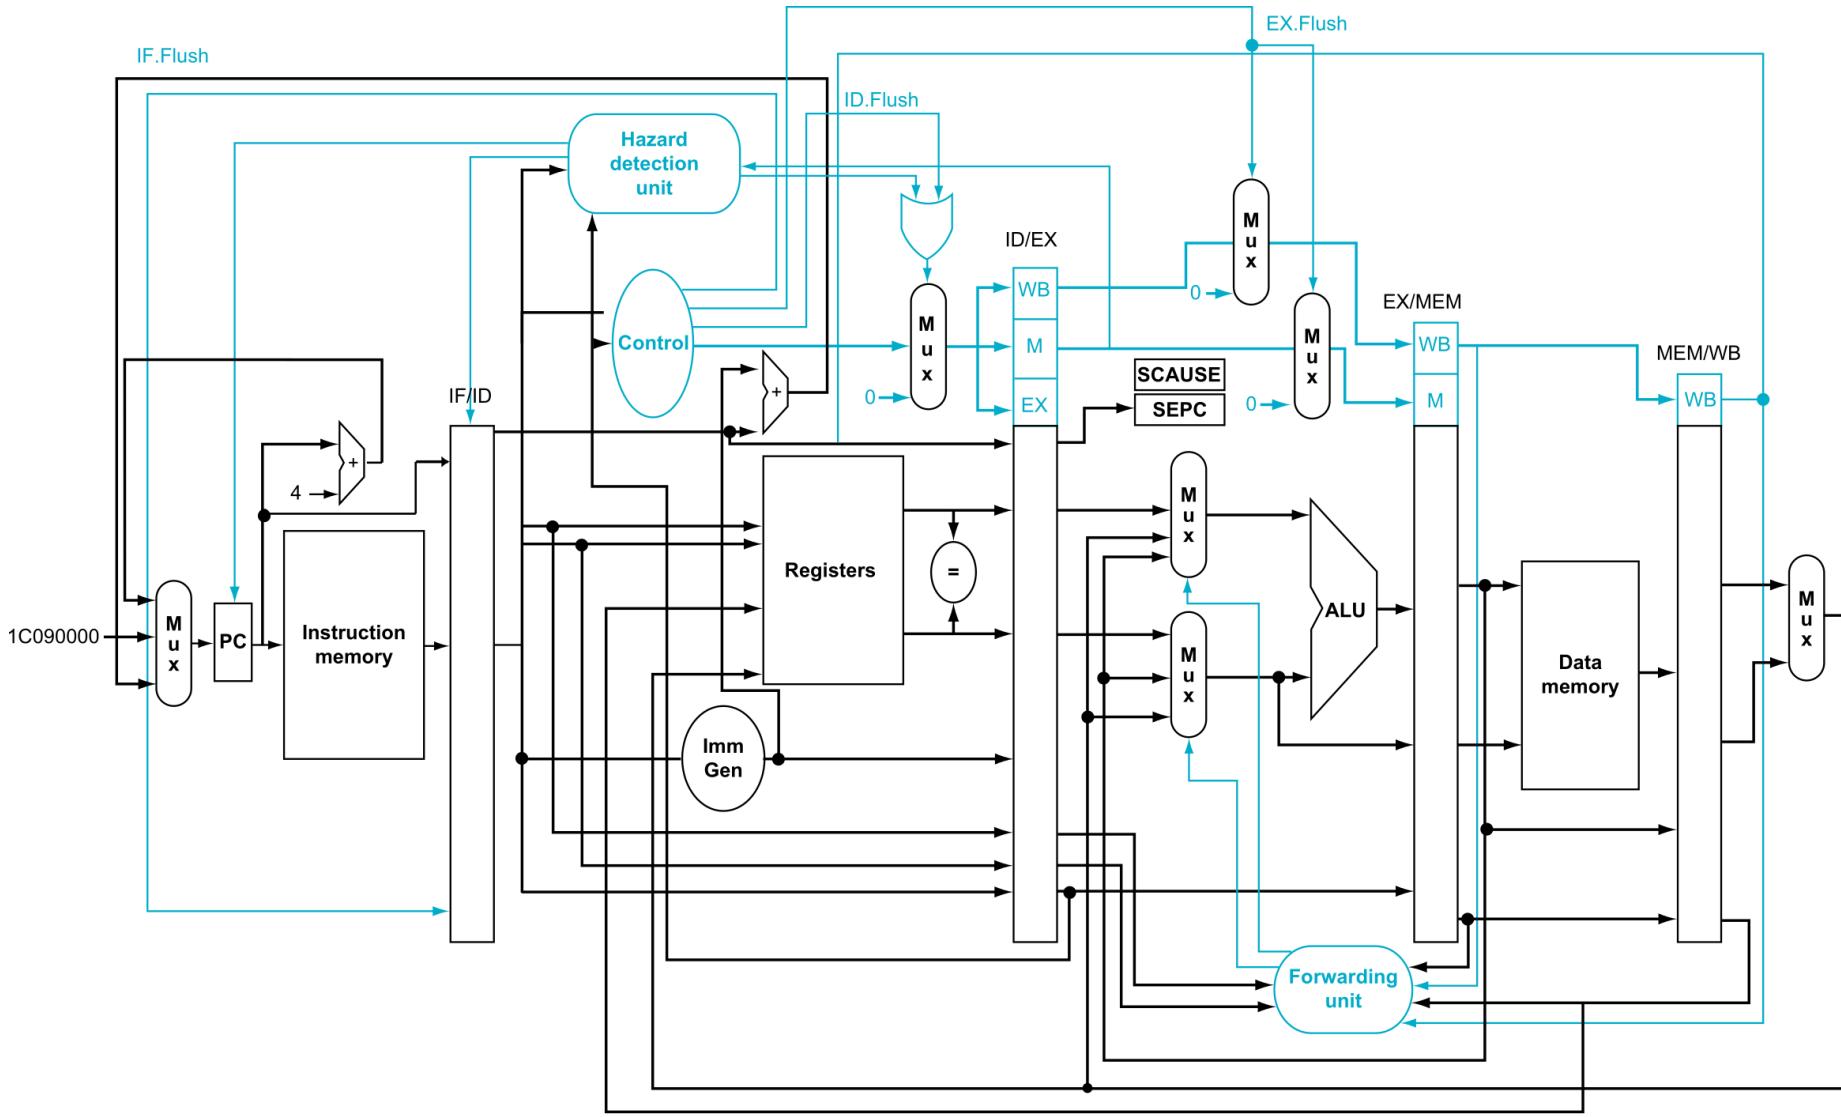
- Another form of control hazard
- Consider malfunction on add in EX stage

add x1, x2, x1

- Prevent x1 from being clobbered : *operation fail, so must not override*
- Complete previous instructions
- Flush add and subsequent instructions (*flush register*)
- Set SEPC and SCAUSE register values (*address and reason*)
- Transfer control to handler

- Similar to mispredicted branch: use much of the same hardware

Pipeline with Exception Control



Exception Example

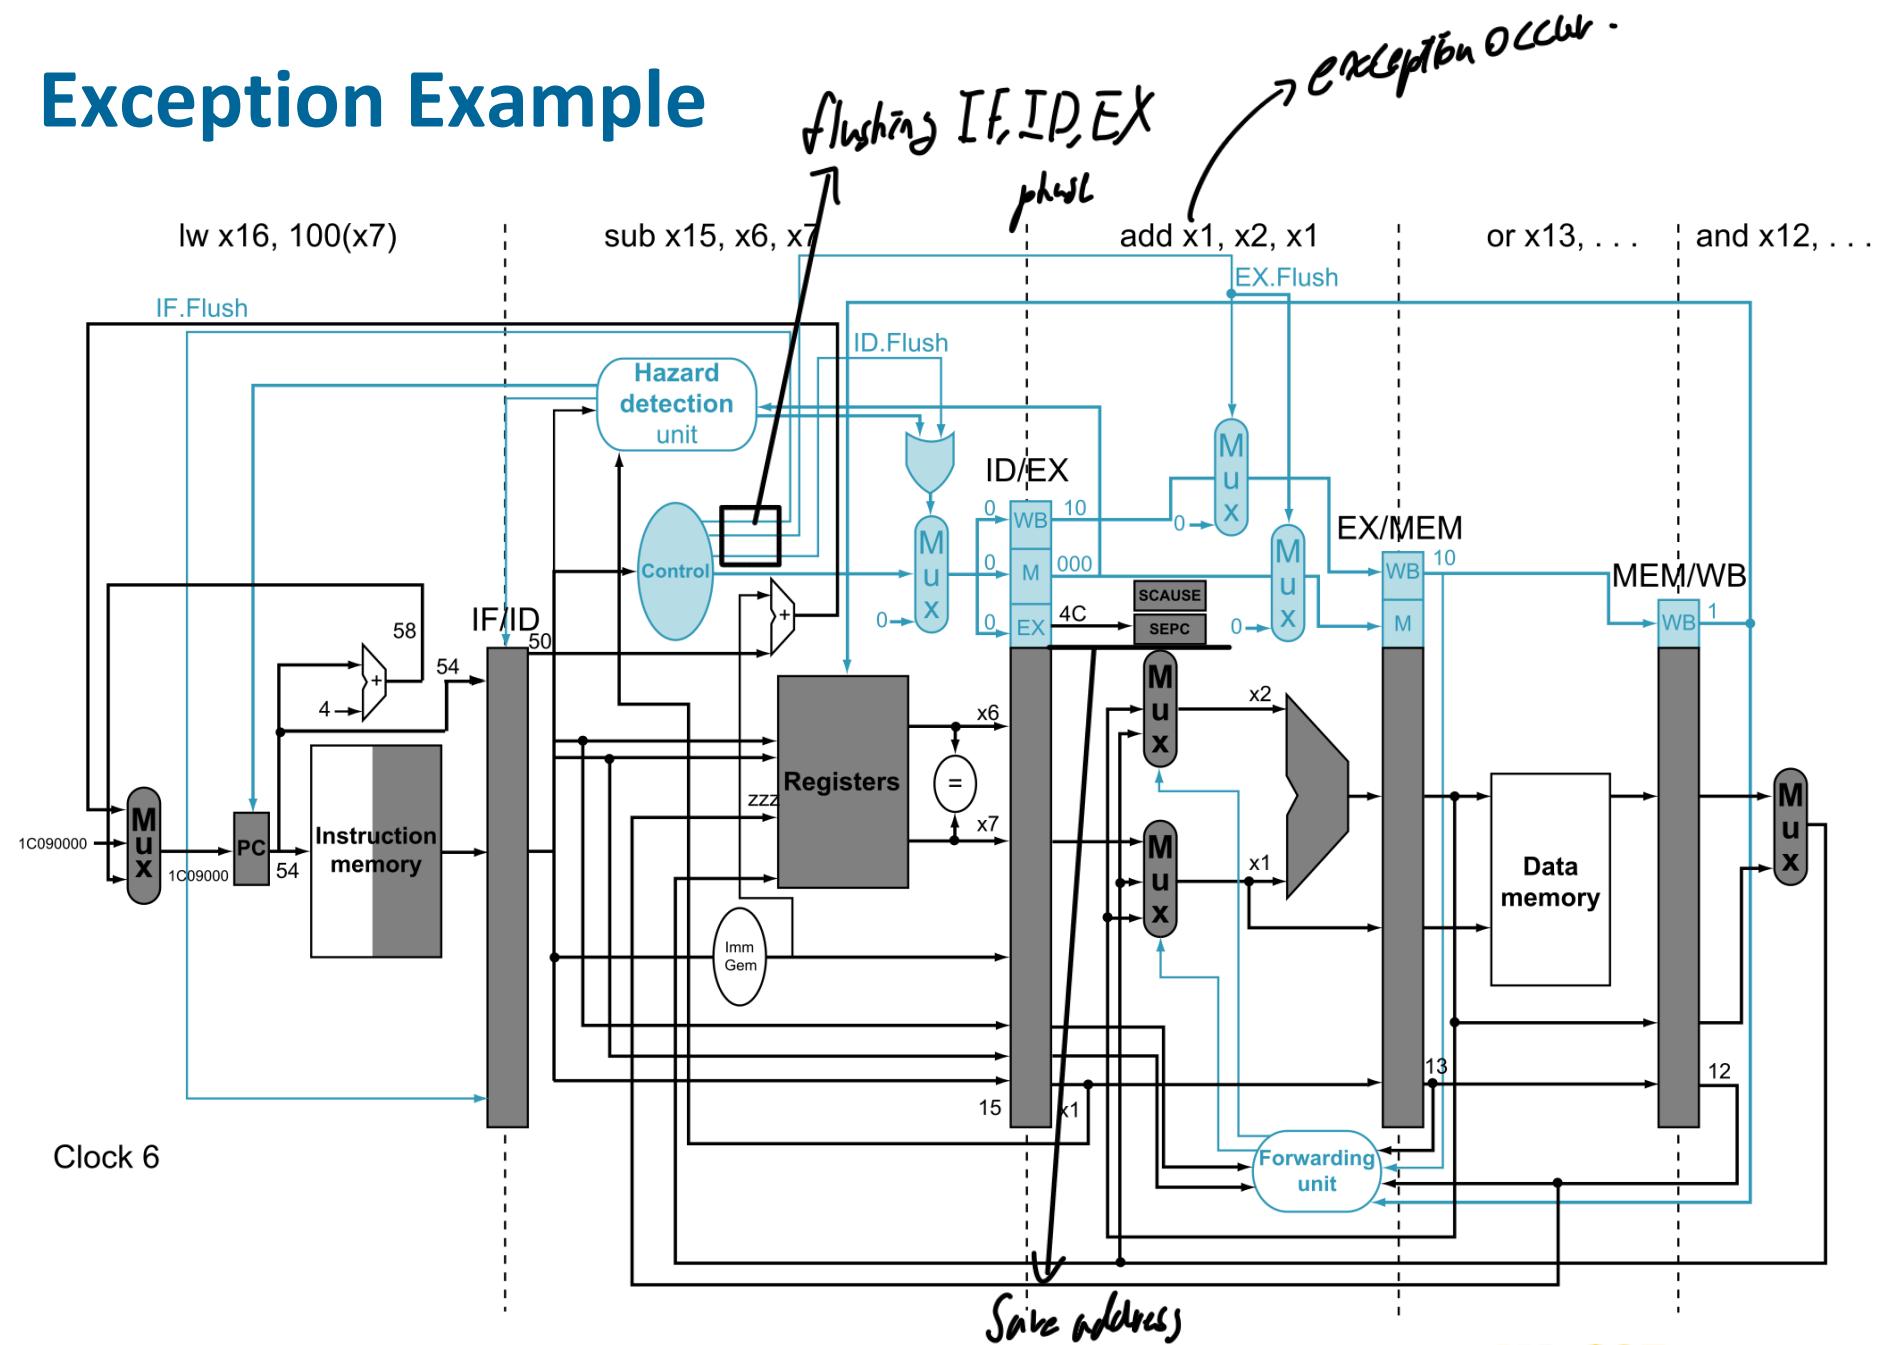
- Exception (hardware fault) on **add** in EX stage

```
40:    sub    x11, x2, x4
44:    and    x12, x2, x5
48:    or     x13, x2, x6
→4c:    add    x1, x2, x1 Error!)
50:    sub    x15, x6, x7
54:    lw     x16, 100(x7)
```

- Handler

```
1c090000      sw    x26, 1000(x10)
1c090004      sw    x27, 1008(x10)
...
```

Exception Example



Exception Example

handler

flush instruction input after exception

sw x26, 1000(x0)

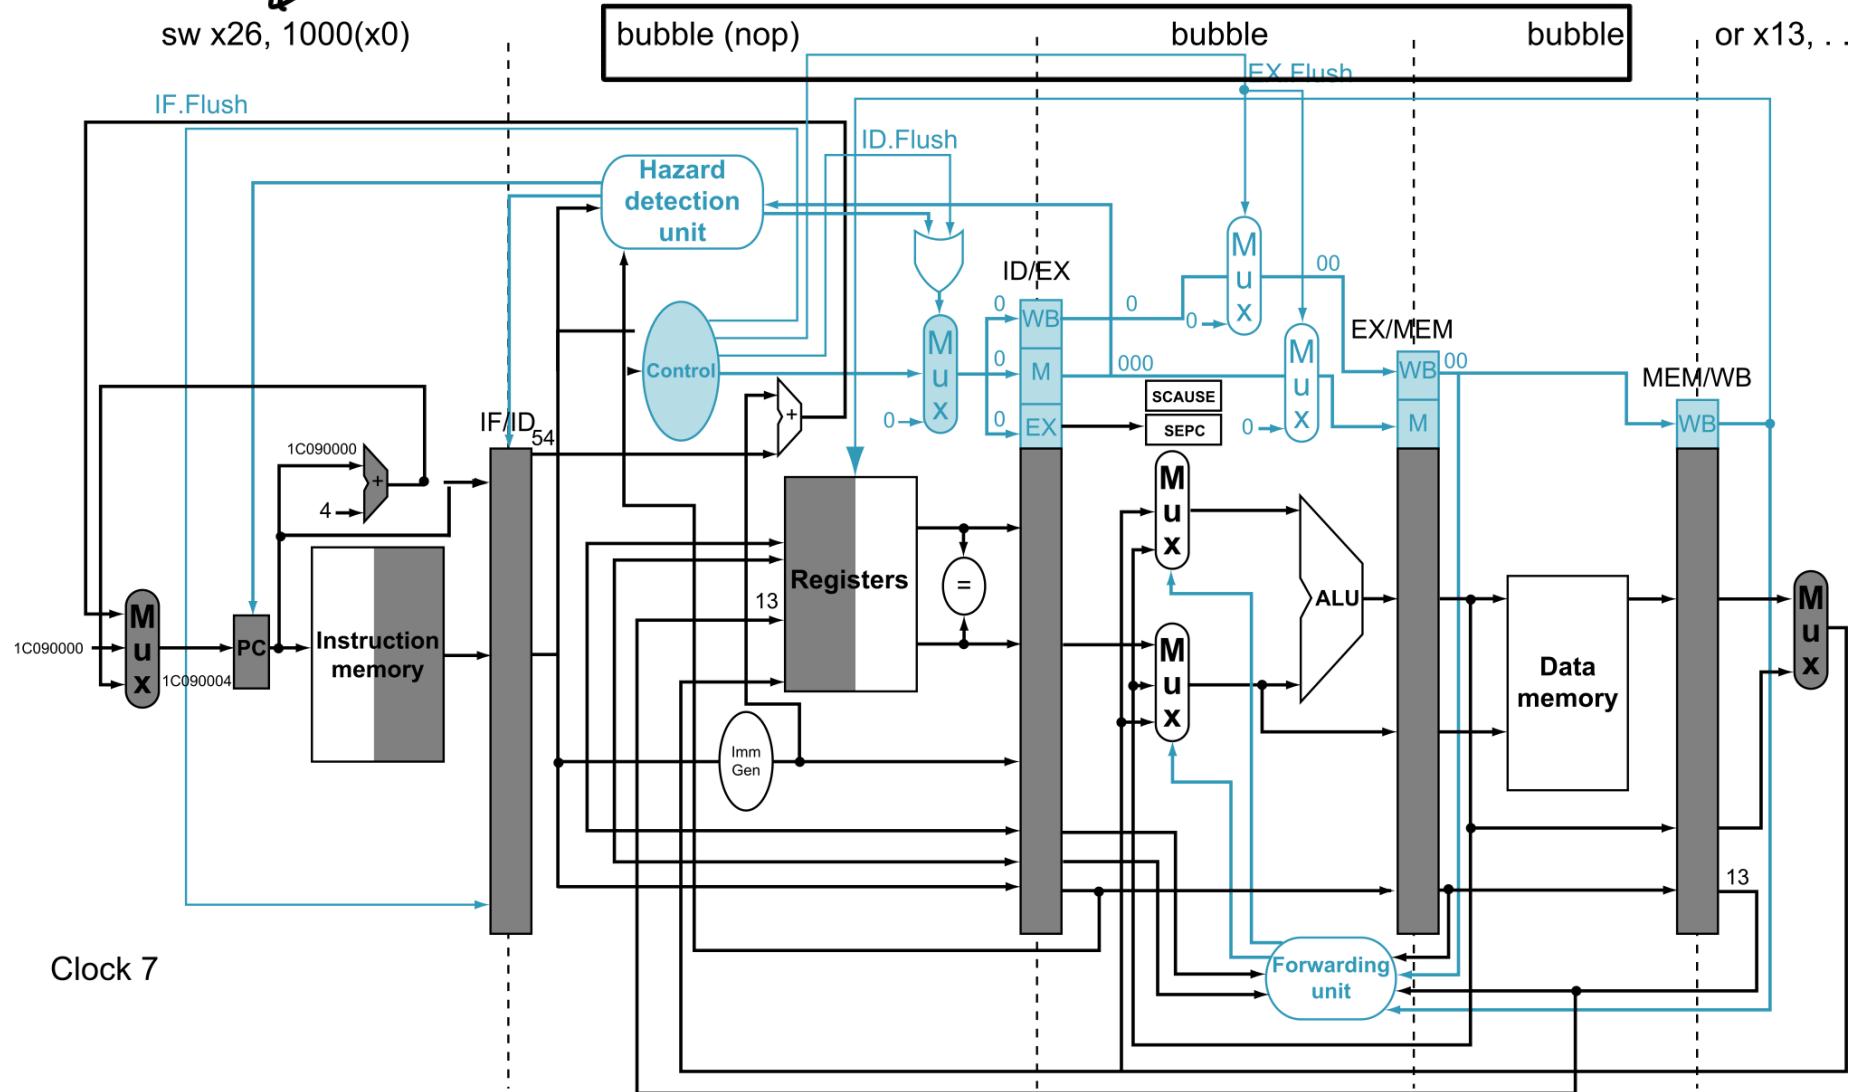
bubble (nop)

bubble

bubble

or x13, . . .

IF.Flush

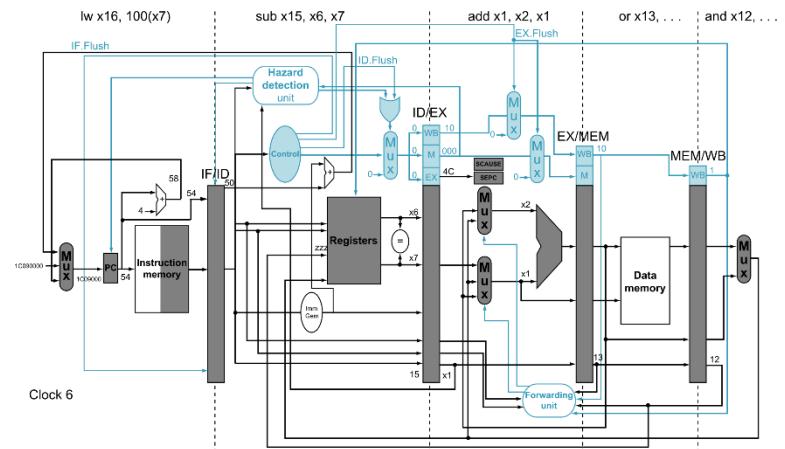


More on Exceptions

■ Multiple exceptions

- Pipelining overlaps multiple instructions - could have multiple exceptions at once!
- Simple approach: deal with exception from earliest instruction
 - ▶ Flush subsequent instructions
 - ▶ “Precise” exceptions
- In complex pipelines
 - ▶ Multiple instructions issued per cycle
 - ▶ Out-of-order completion
 - ▶ Maintaining precise exceptions is difficult!
- “Imprecise” exceptions
 - ▶ Exception not associated with exact instruction that caused the interrupt
- Majority of modern processors, incl. RISC-V, support precise interrupts

Module Summary



Module Summary

■ Data hazards

- Solved with data forwarding and hazard detection
- Can eliminate stalls completely except for load interlock
- Relatively simple

■ Control hazards

- Reduce branch penalty by moving branch decision to the ID stage
 - ▶ At the expense of additional forwarding paths and hazard detection logic
- Speculative execution
 - ▶ Speculate, flush pipeline if incorrect
 - ▶ Branch prediction with 1 and 2 bit buffers

■ Exceptions

- Unforeseen internal or external events trigger a jump to an interrupt handler
- Address of affected instruction and cause of exception recorded in registers
- Exceptions in pipelines are passed on, similar to other data, and triggered when the instruction retires

Storage Technology and Trend

SATA, M2, NVMe, PCIe

Multi-Level Cells